

# Convolutional Neural Networks

Theo Heimel

May 2025

CP3, UCLouvain



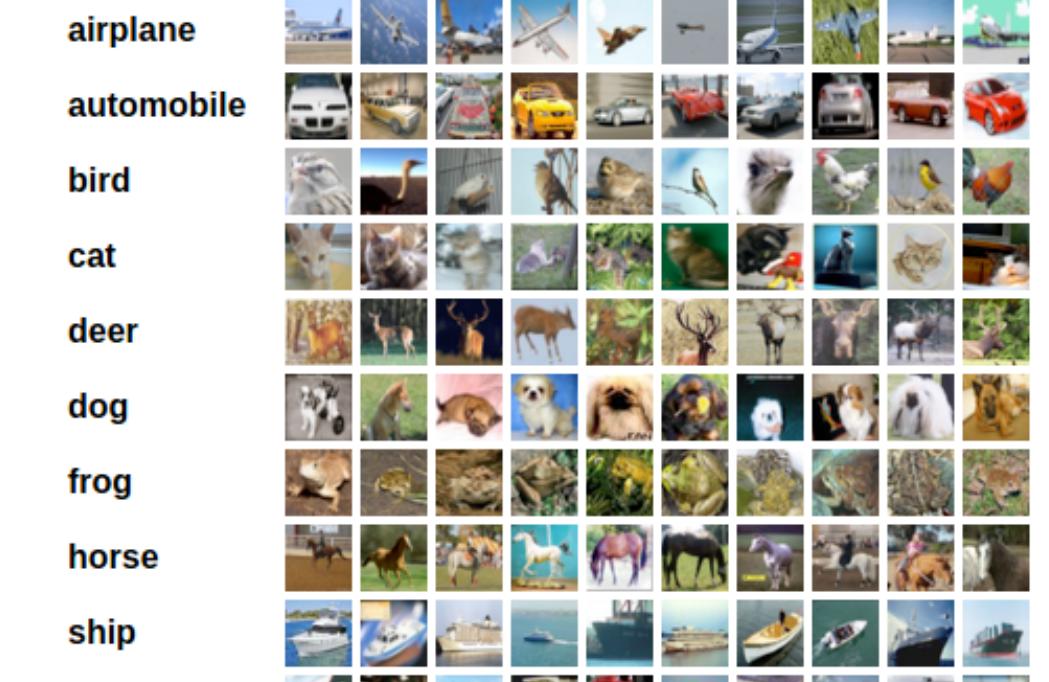
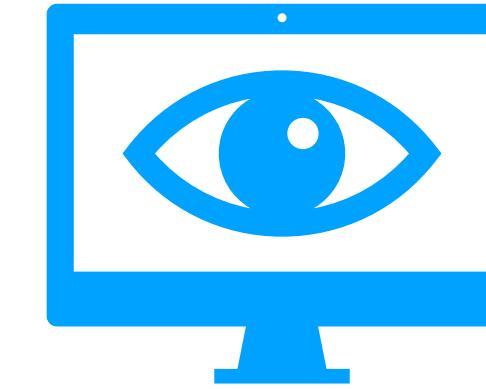
# Introduction



Tagging



Segmentation



Classification

Computer vision tasks appear in many areas of science, engineering and technology:  
self-driving cars, medicine, astronomy, and many more...

# Introduction



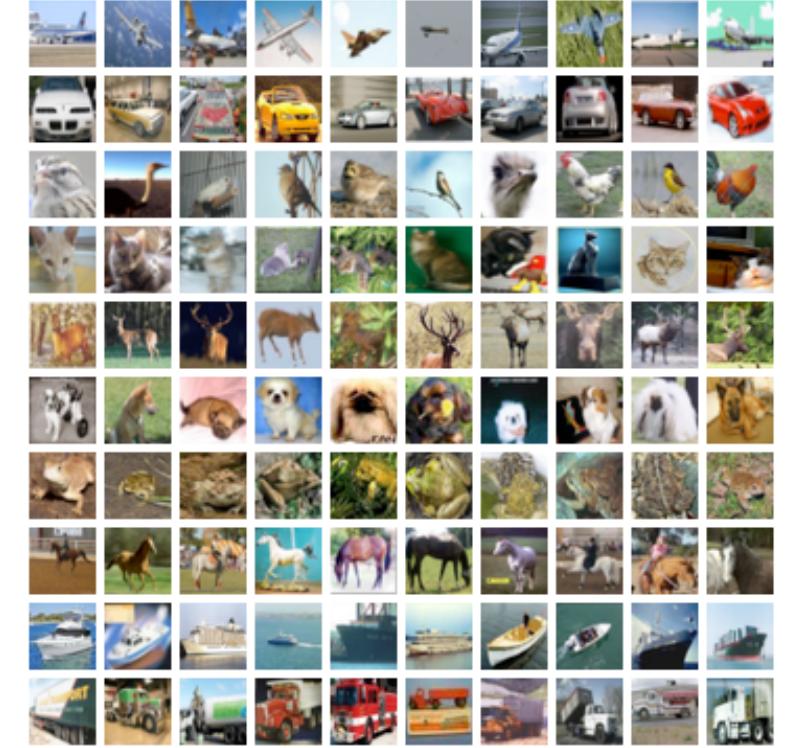
Tagging



Segmentation



airplane  
automobile  
bird  
cat  
deer  
dog  
frog  
horse  
ship  
truck



Classification

Underlying most neural networks for these tasks:  
**convolutional neural networks**

# Outline

- ① Challenges of image processing
  - ② Introduction to convolutions
  - ③ Convolutional layers
  - ④ CNN building blocks
  - ⑤ A short history of CNNs
  - ⑥ Outlook
- Slides based on previous CNN lectures at ERuM schools, especially those by Sascha Diefenbacher.
- More lectures on CNNs:  
[ErUM DataHub Material Collection](#)

# Outline

- ① **Challenges of image processing**
- ② Introduction to convolutions
- ③ Convolutional layers
- ④ CNN building blocks
- ⑤ A short history of CNNs
- ⑥ Outlook

# What is an image?



color image  
↓  
red, green,  
blue channels  
↓  
matrices of  
intensities

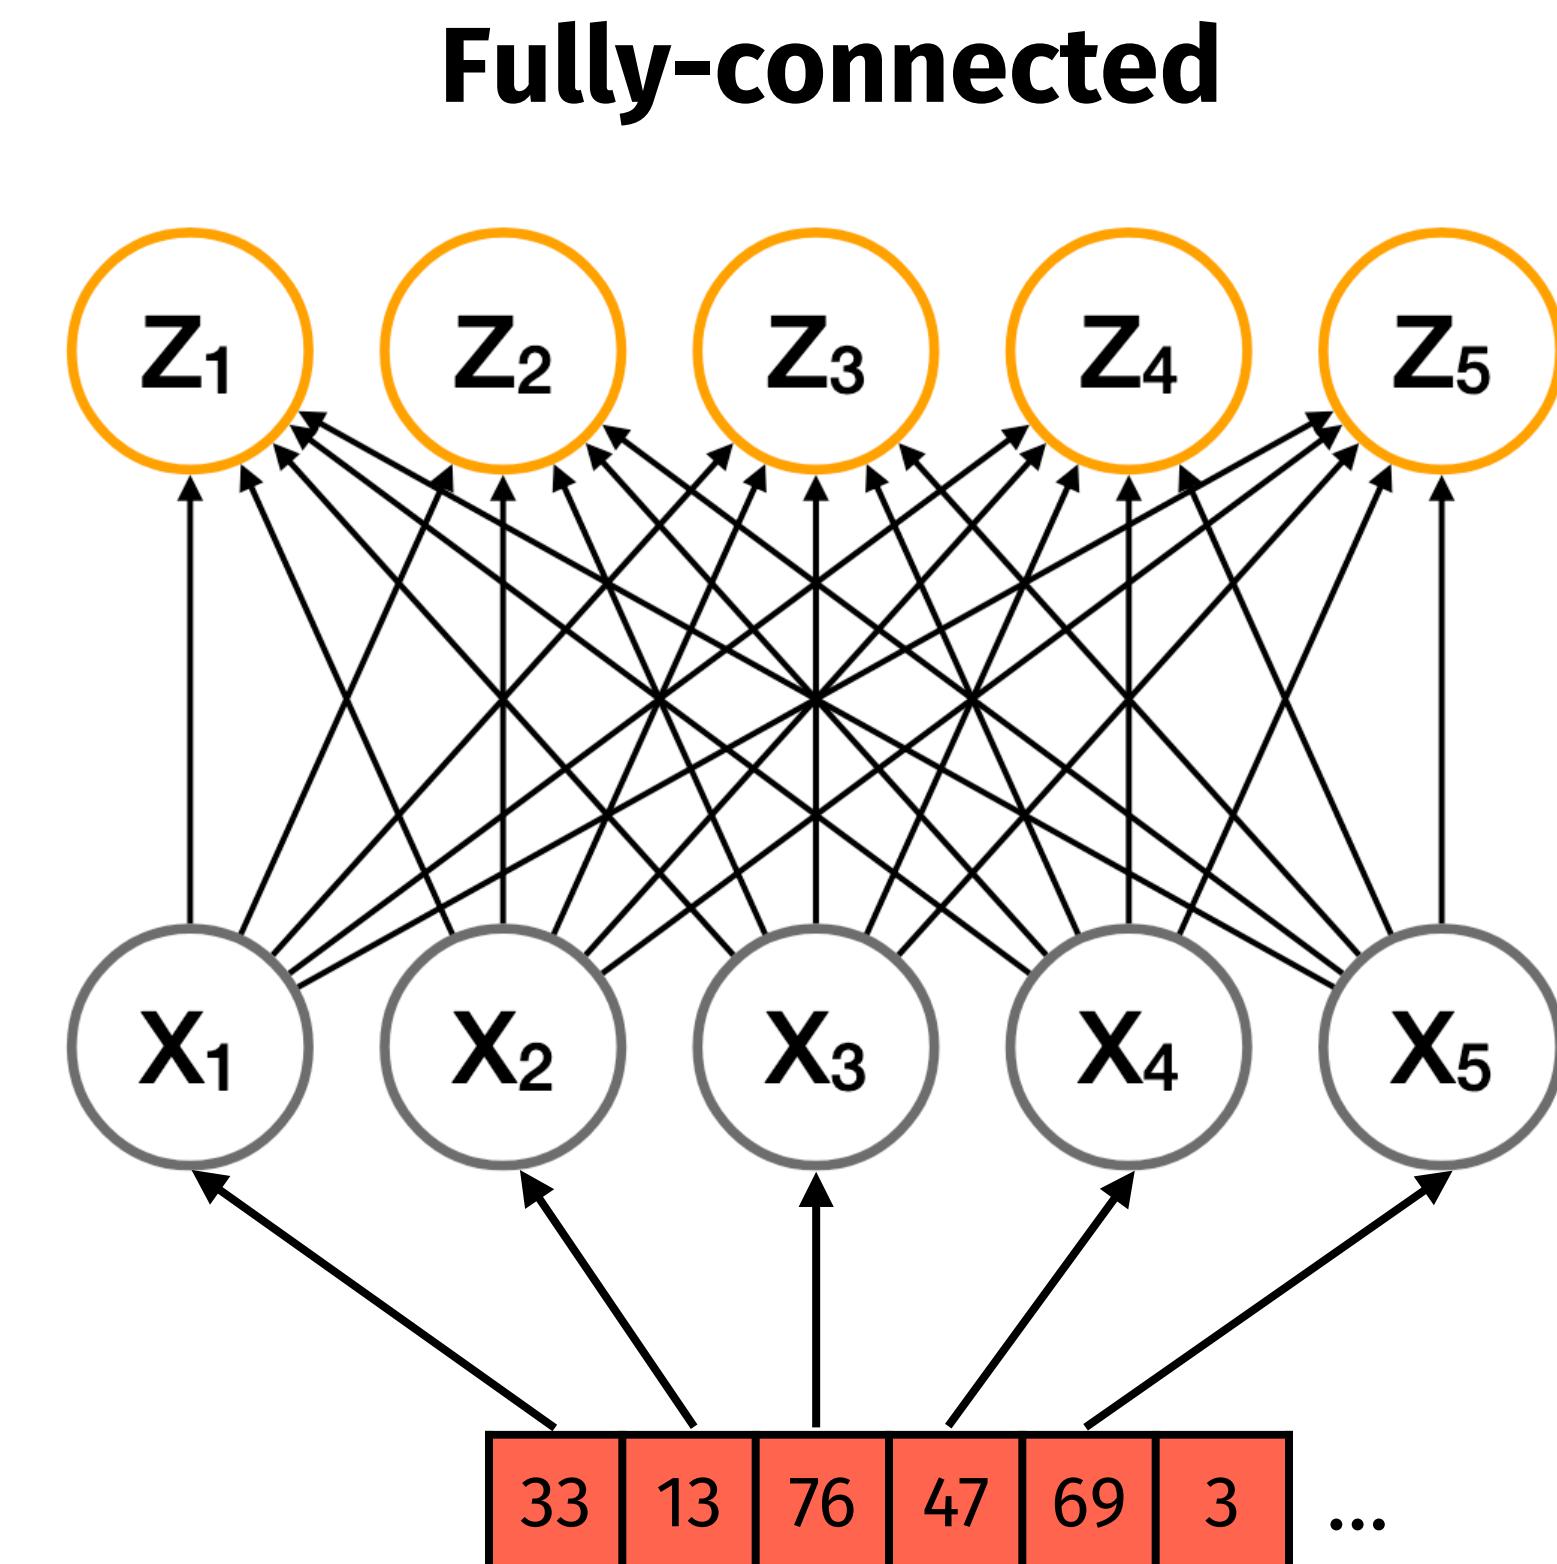
33	13	76	47	69	3
68	99	37	2	16	25
3	1	48	52	52	16
12	87	24	21	44	86
19	48	0	27	64	69
66	88	31	15	16	47

52	15	8	82	37	53
32	20	85	40	74	77
62	30	74	78	36	90
58	95	58	7	69	55
65	8	66	69	94	8
0	65	80	41	26	60

10	21	45	48	15	19
33	93	77	57	75	71
24	9	1	45	75	1
41	71	68	95	71	38
43	33	74	34	86	56
61	88	21	79	95	30

# Training neural nets on images

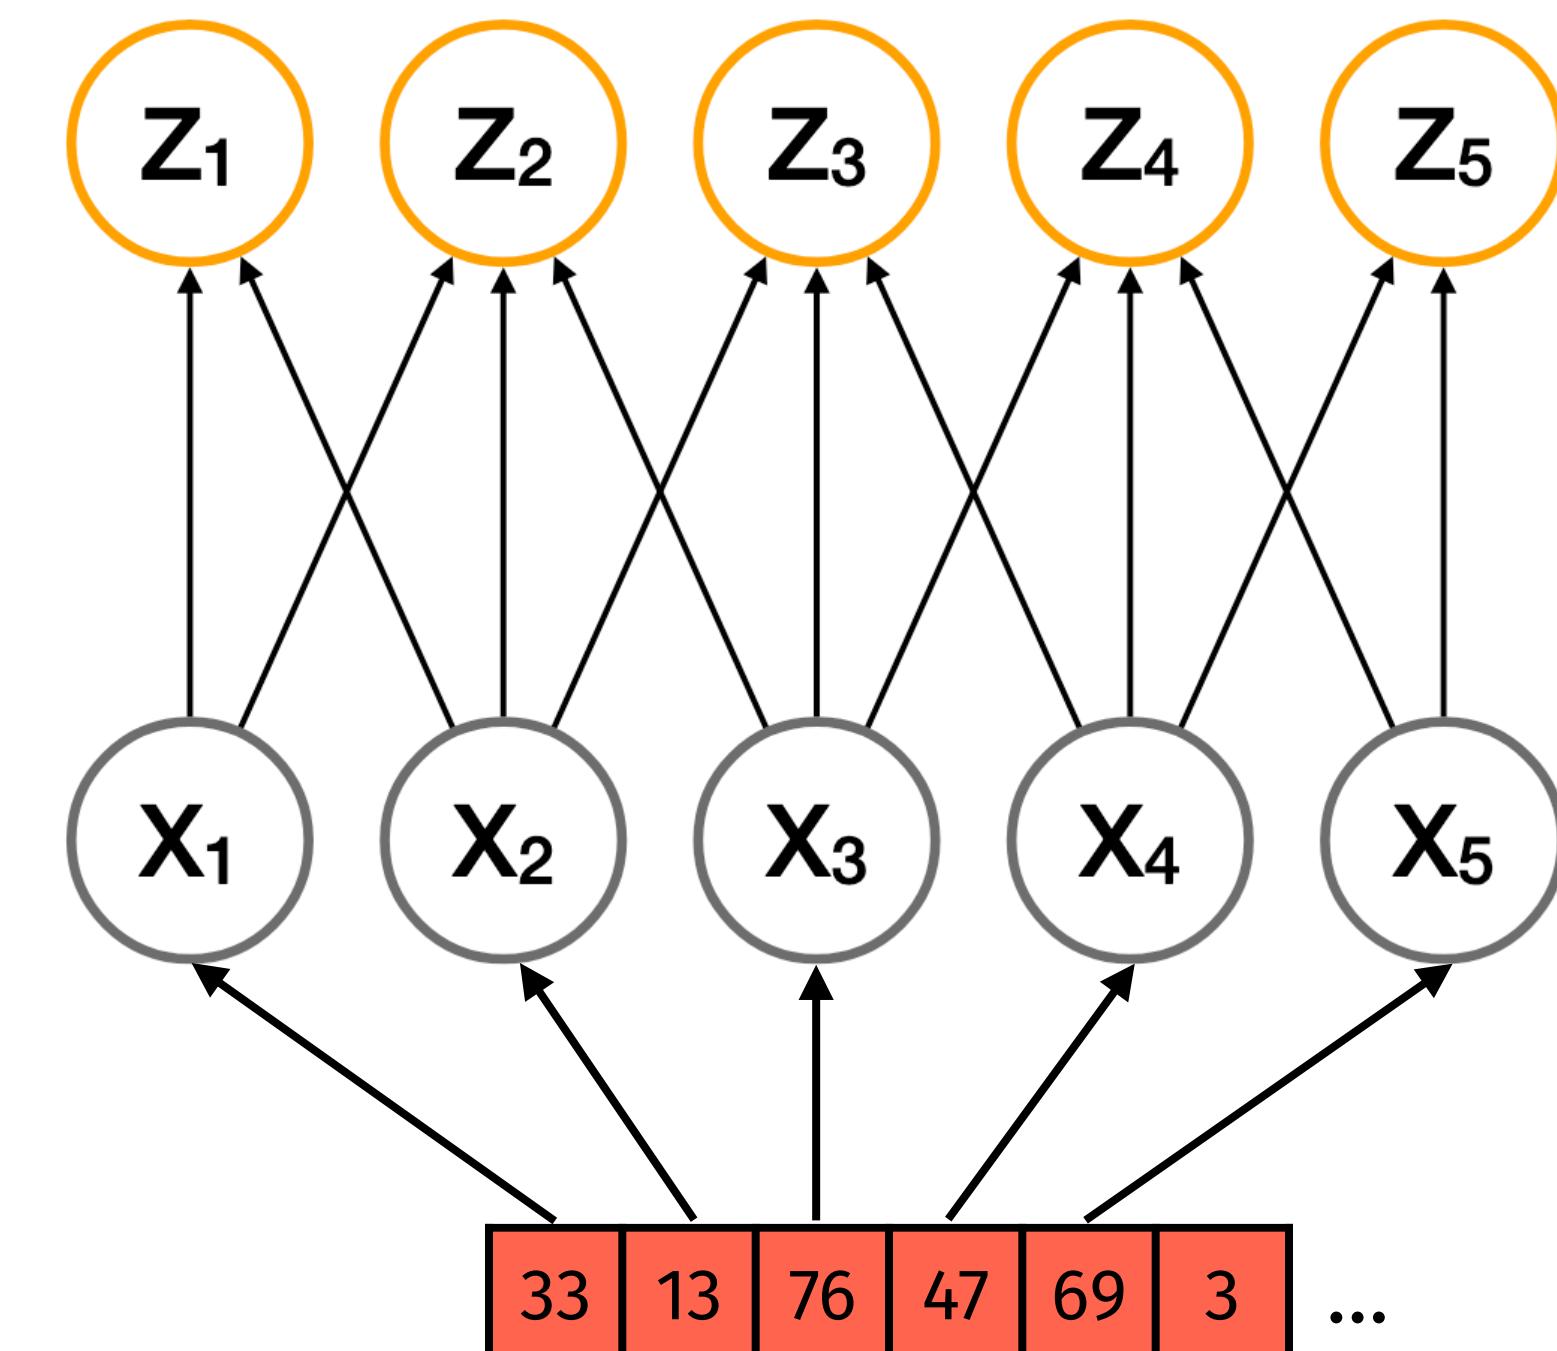
- Naive idea: flatten all values, build fully-connected network
- Problem: network would too large, even at low resolution
- Example: CIFAR-10 RGB image  
→  $32 \times 32 \times 3 = 3072$  inputs  
→  $3072 \times 3072 = 9440256$  parameters/layer



# Training neural nets on images

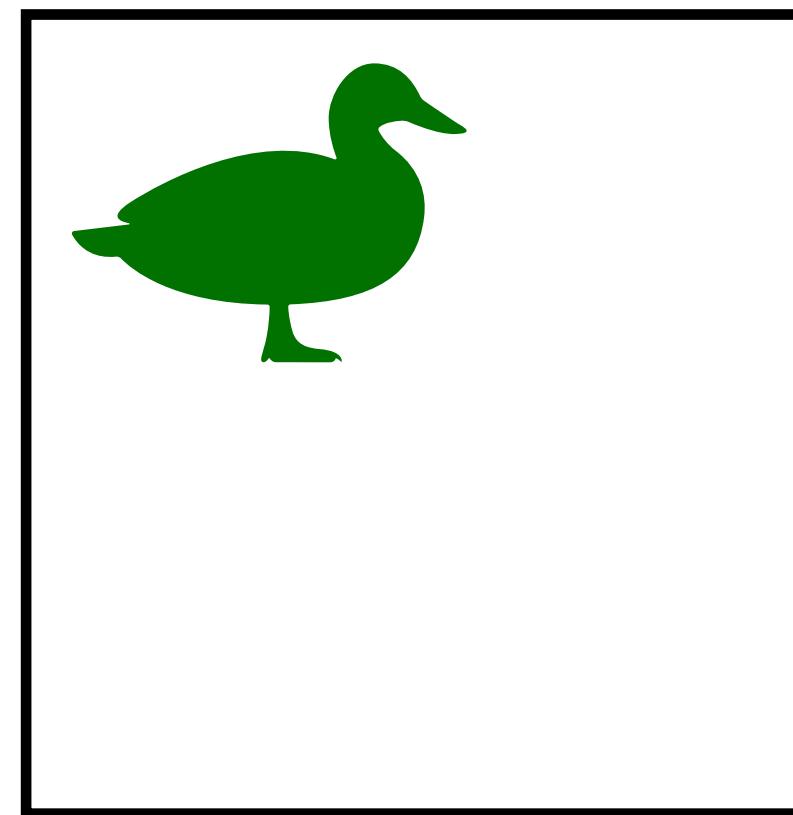
- Problem: network would too large, even at low resolution
- Solution: image features have limited size  
→ don't need to see whole image at once  
→ use **locally connected network**
- Much lower number of parameters  
→ can we reduce them even more?

**Locally connected**

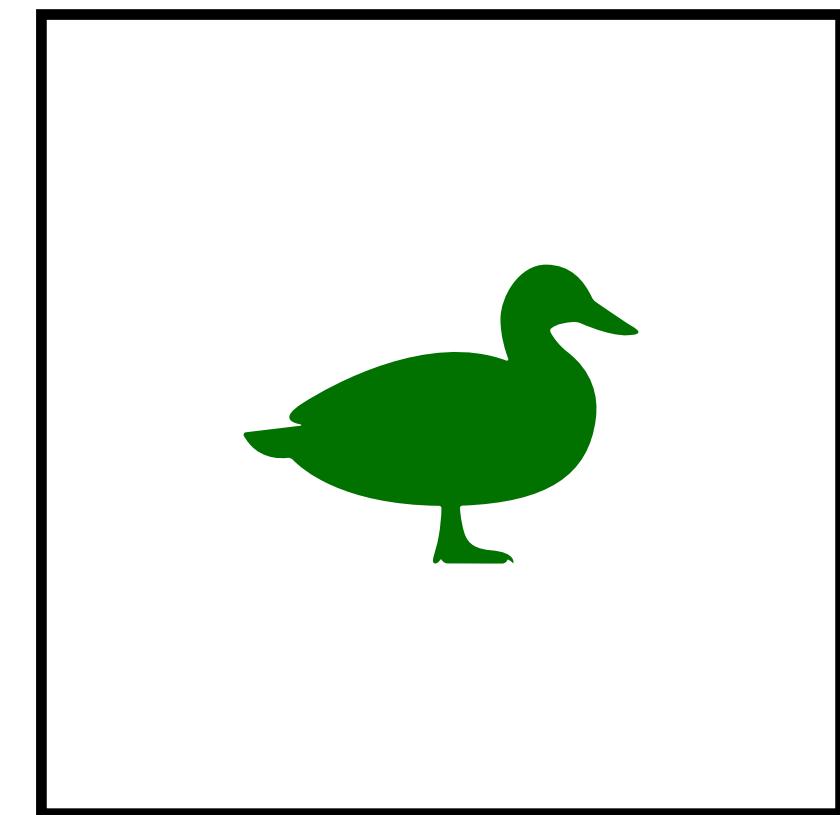


# Translation invariance

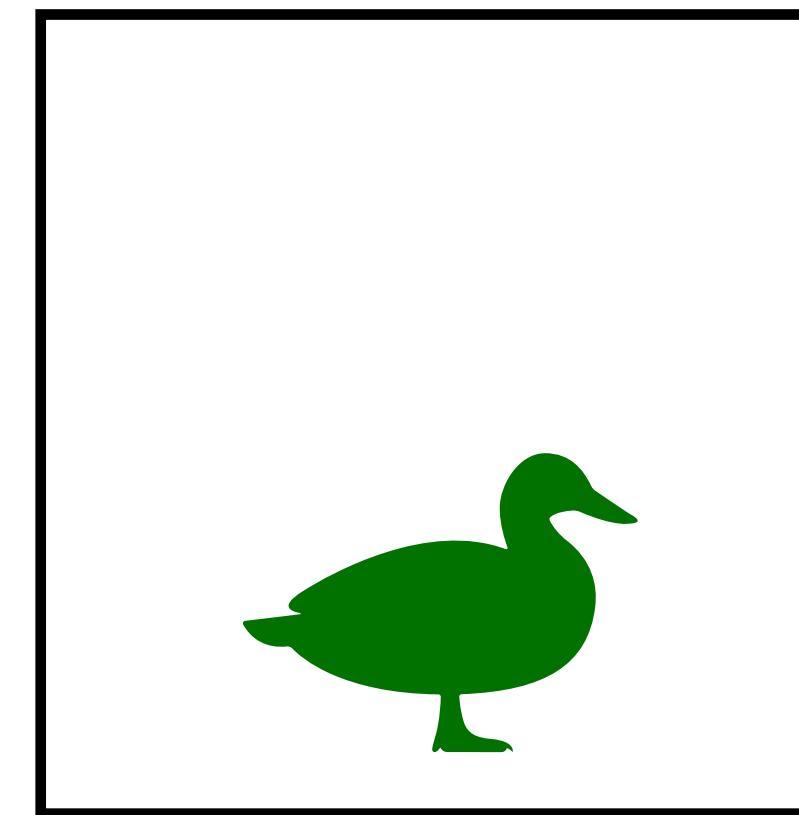
A duck



Also a duck



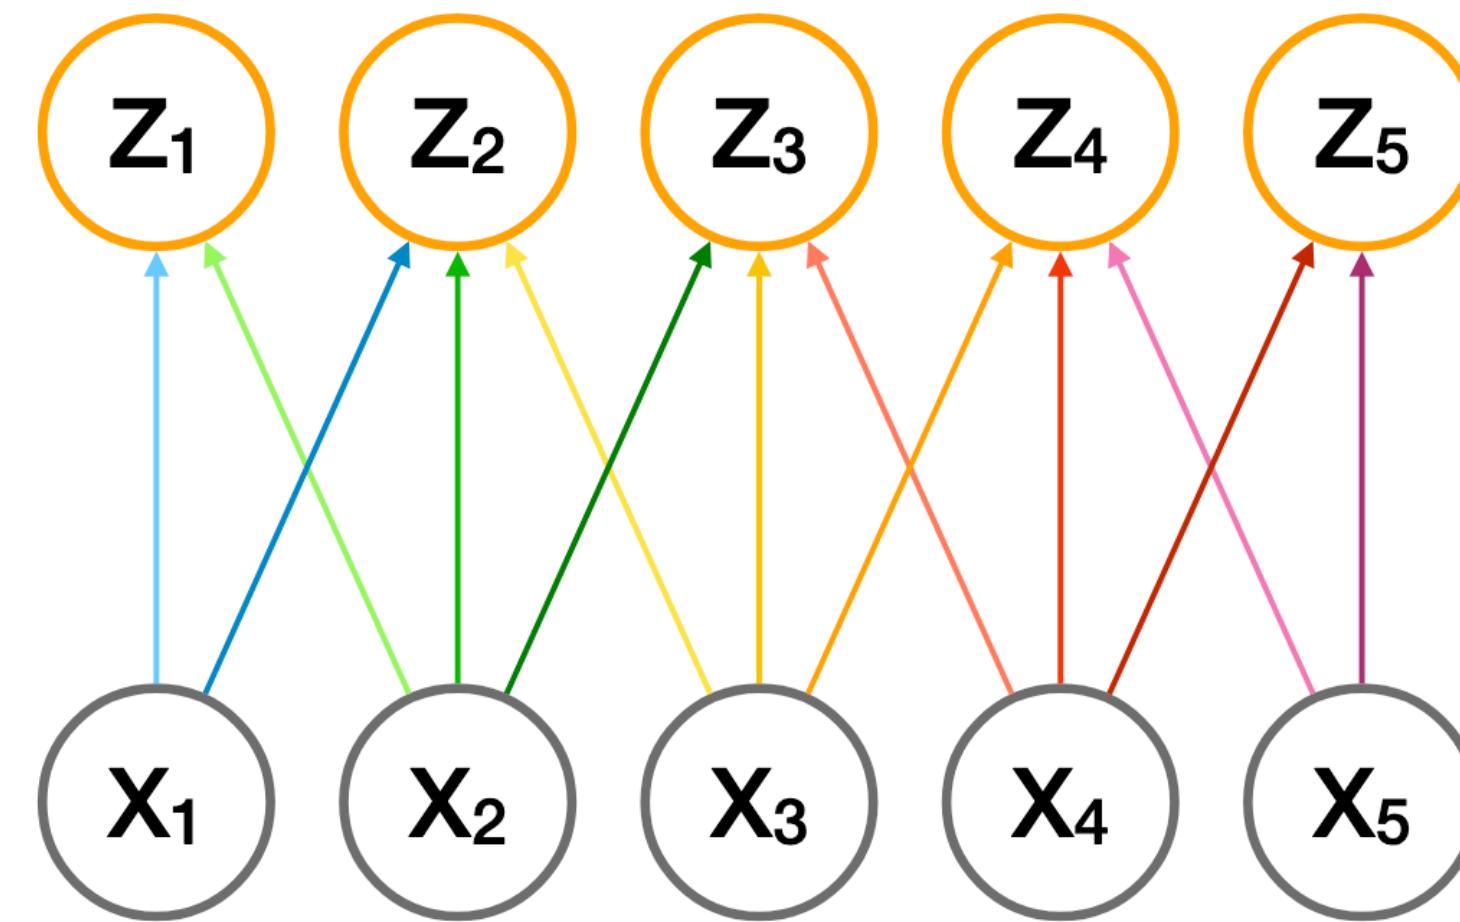
Another duck



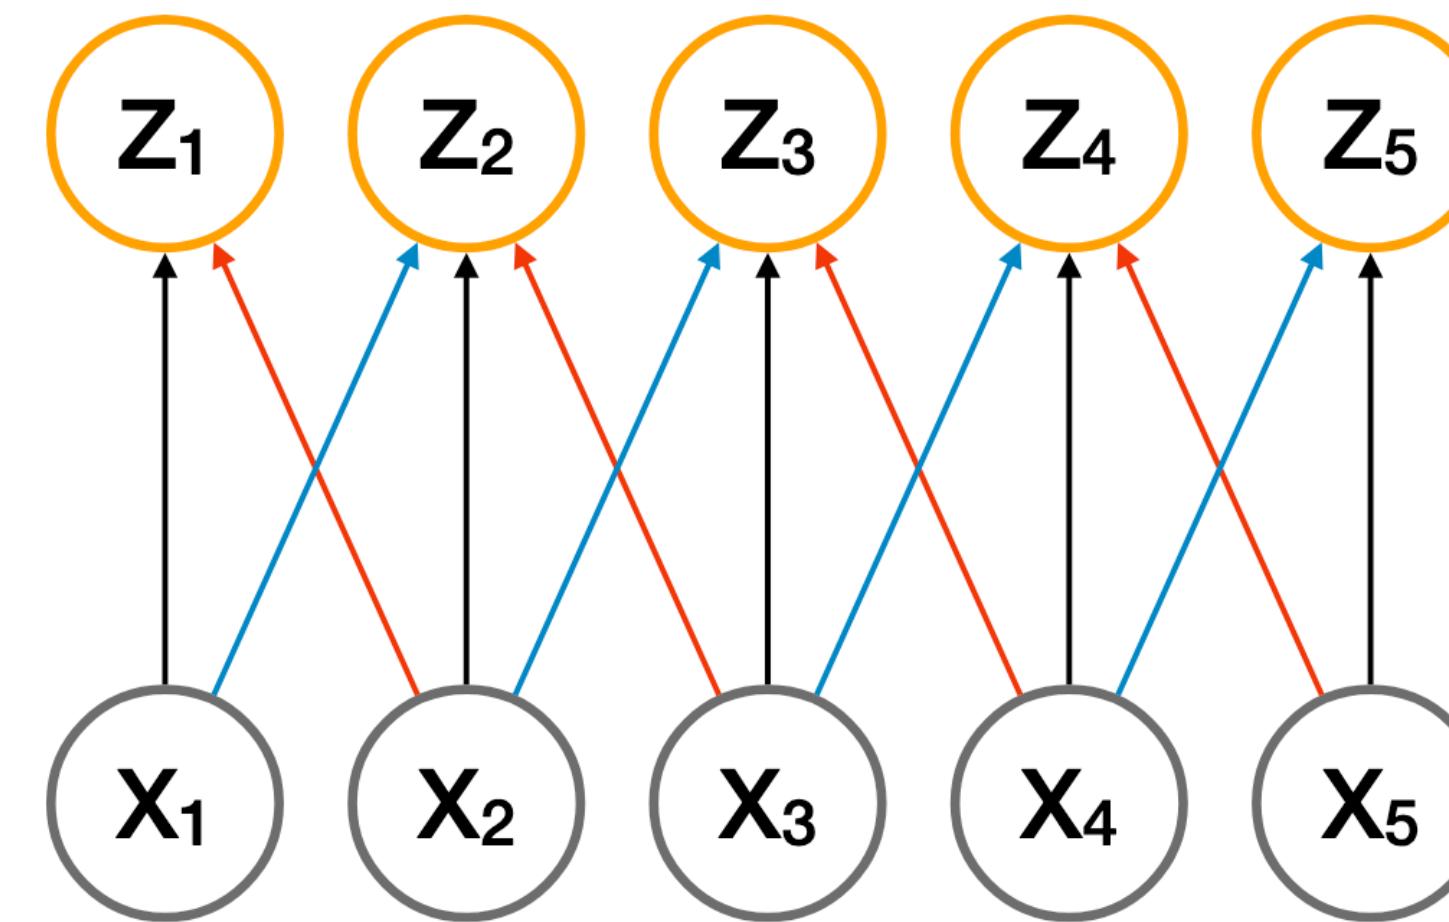
- Classification of object is independent of position in image  
→ translation invariance
- With locally connected network:  
have to relearn task for every position

# Parameter sharing

**Locally connected**



**Convolution**

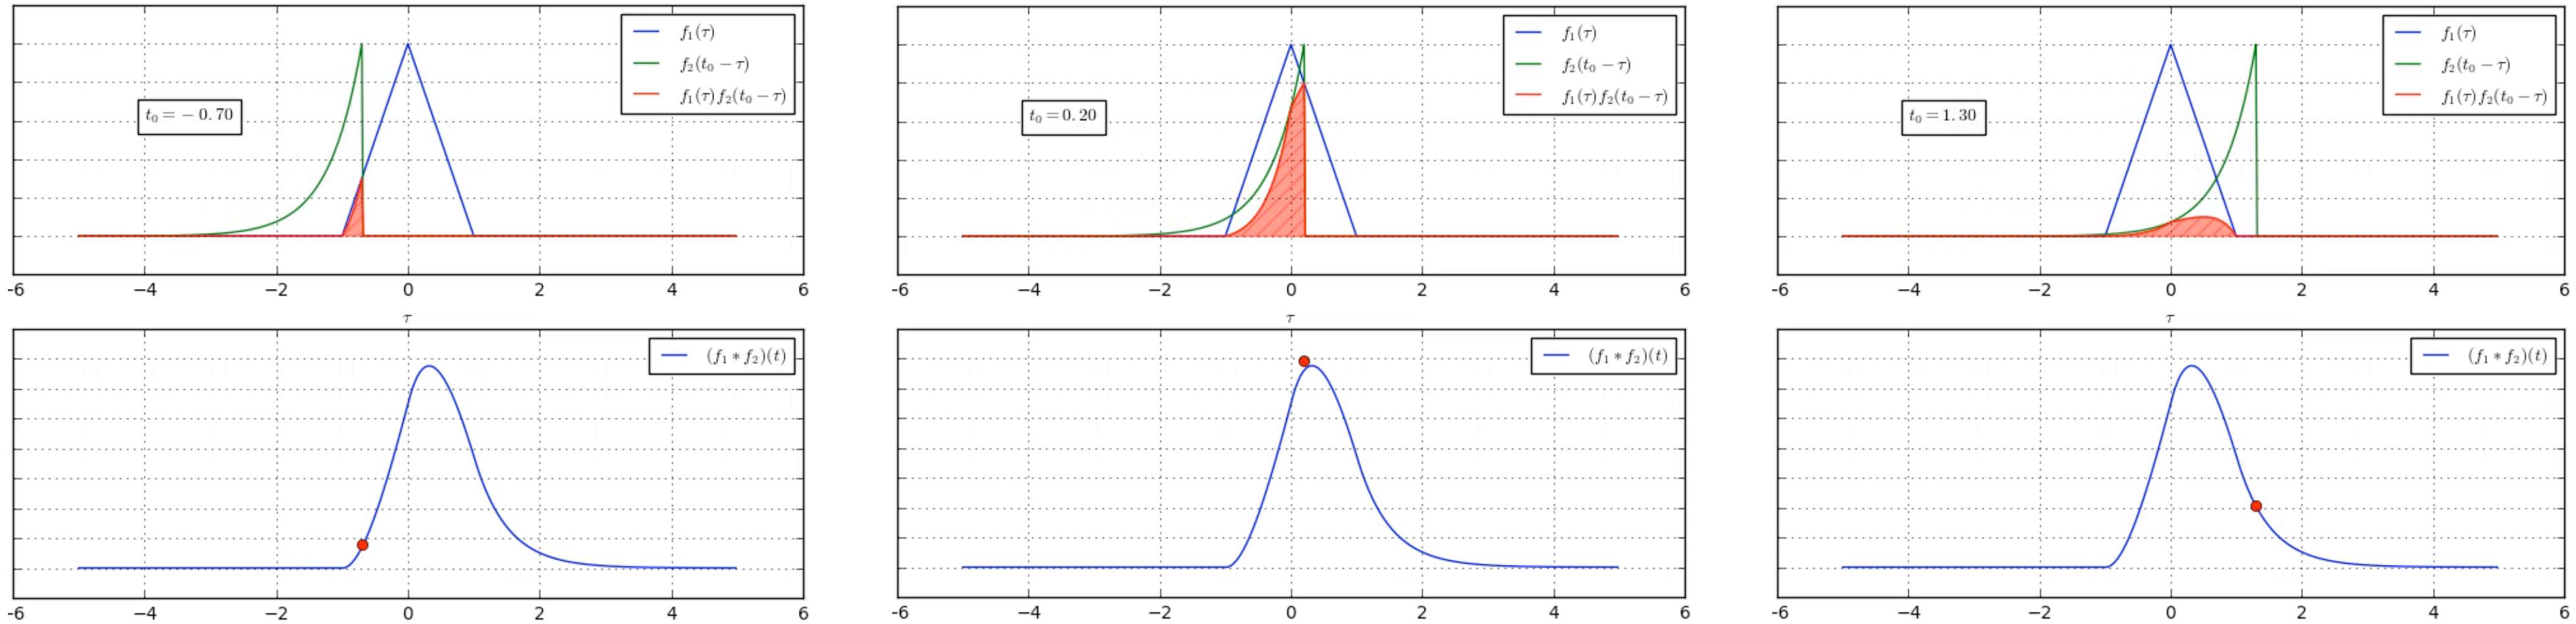


- Classification of object is independent of position in image  
→ translation invariance
- With locally connected network:  
have to relearn task for every position
- Solution: convolution  
→ **shared parameters for different positions**

# Outline

- ① Challenges of image processing
- ② **Introduction to convolutions**
- ③ Convolutional layers
- ④ CNN building blocks
- ⑤ A short history of CNNs
- ⑥ Outlook

# Mathematical convolutions

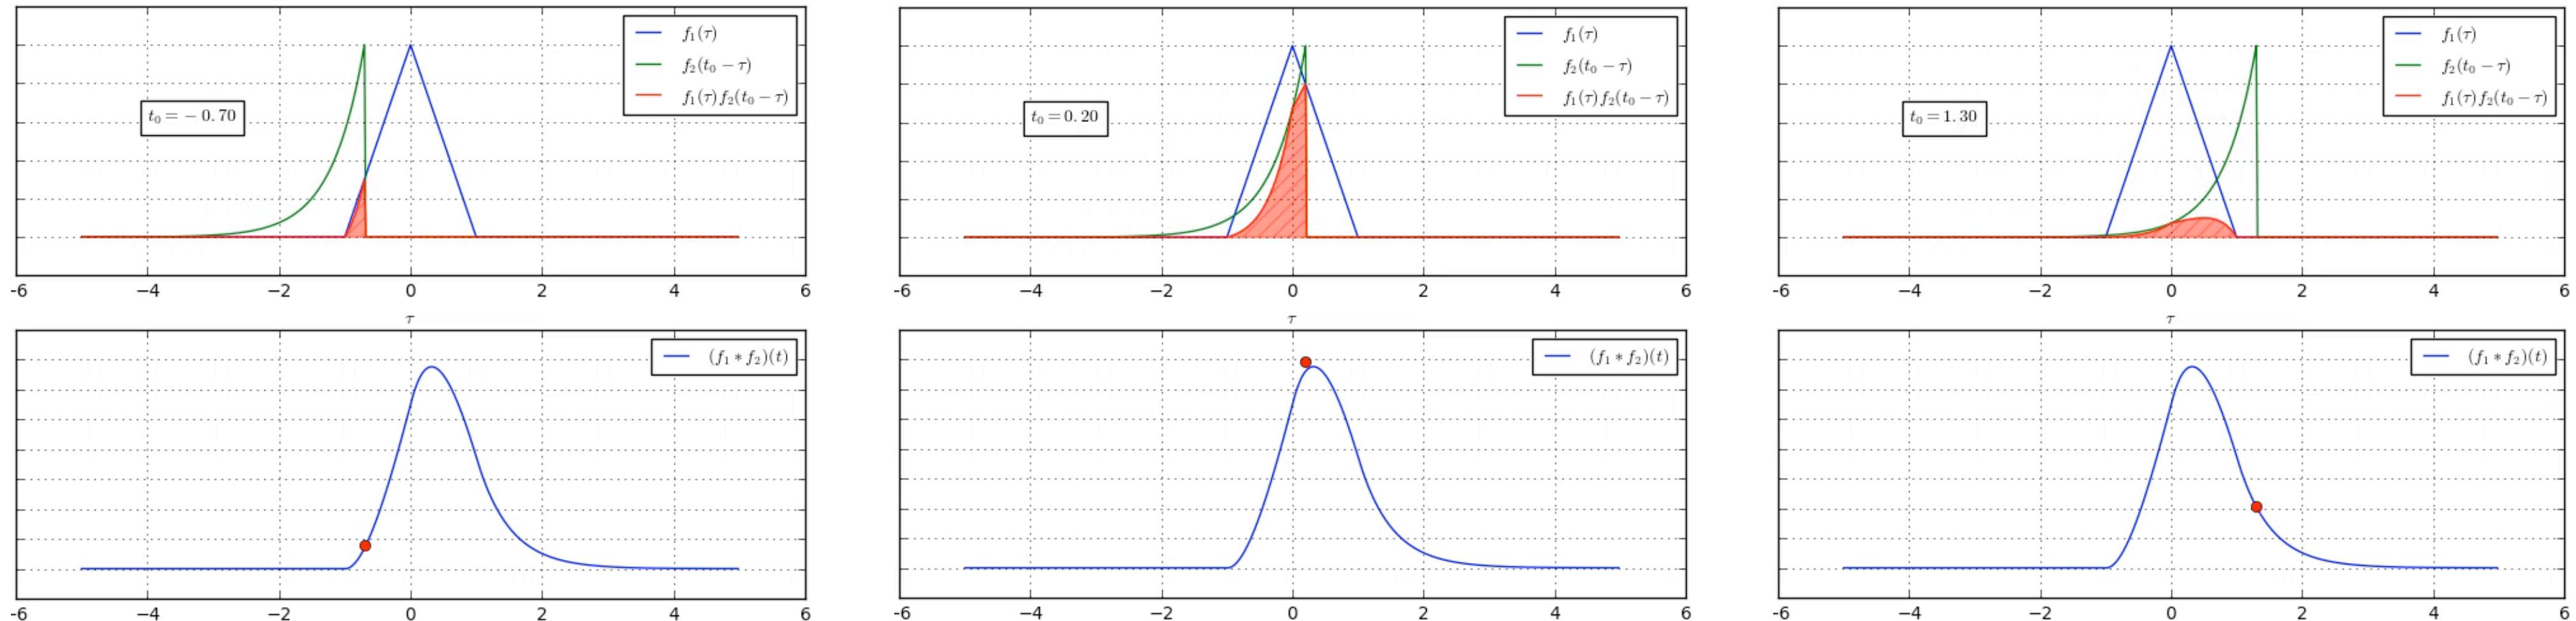


Animated illustration:

<https://dspillustrations.com/pages/posts/misc/convolution-examples-and-the-convolution-integral.html>

- Convolution in mathematics: defined for real-valued functions  $f, g$   
$$(f * g)(t) = \int_{-\infty}^{\infty} d\tau f(\tau) g(t - \tau)$$
- Applications in statistics, signal processing, ... and many more

# Mathematical convolutions



Animated illustration:

<https://dspillustrations.com/pages/posts/misc/convolution-examples-and-the-convolution-integral.html>

- To apply to images, first discretize:  $(I * K)(i) = \sum_{m=-\infty}^{\infty} I(m)K(i - m)$
- Then add a second dimension:

$$(I * K)(i, j) = \sum_{m=-\infty}^{\infty} \sum_{n=-\infty}^{\infty} I(m, n)K(i - m, j - n)$$

# Image convolutions

<b>Image <math>I</math></b>				
1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

<b>Kernel <math>K</math></b>		
0	1	0
1	2	1
0	1	0

*	=

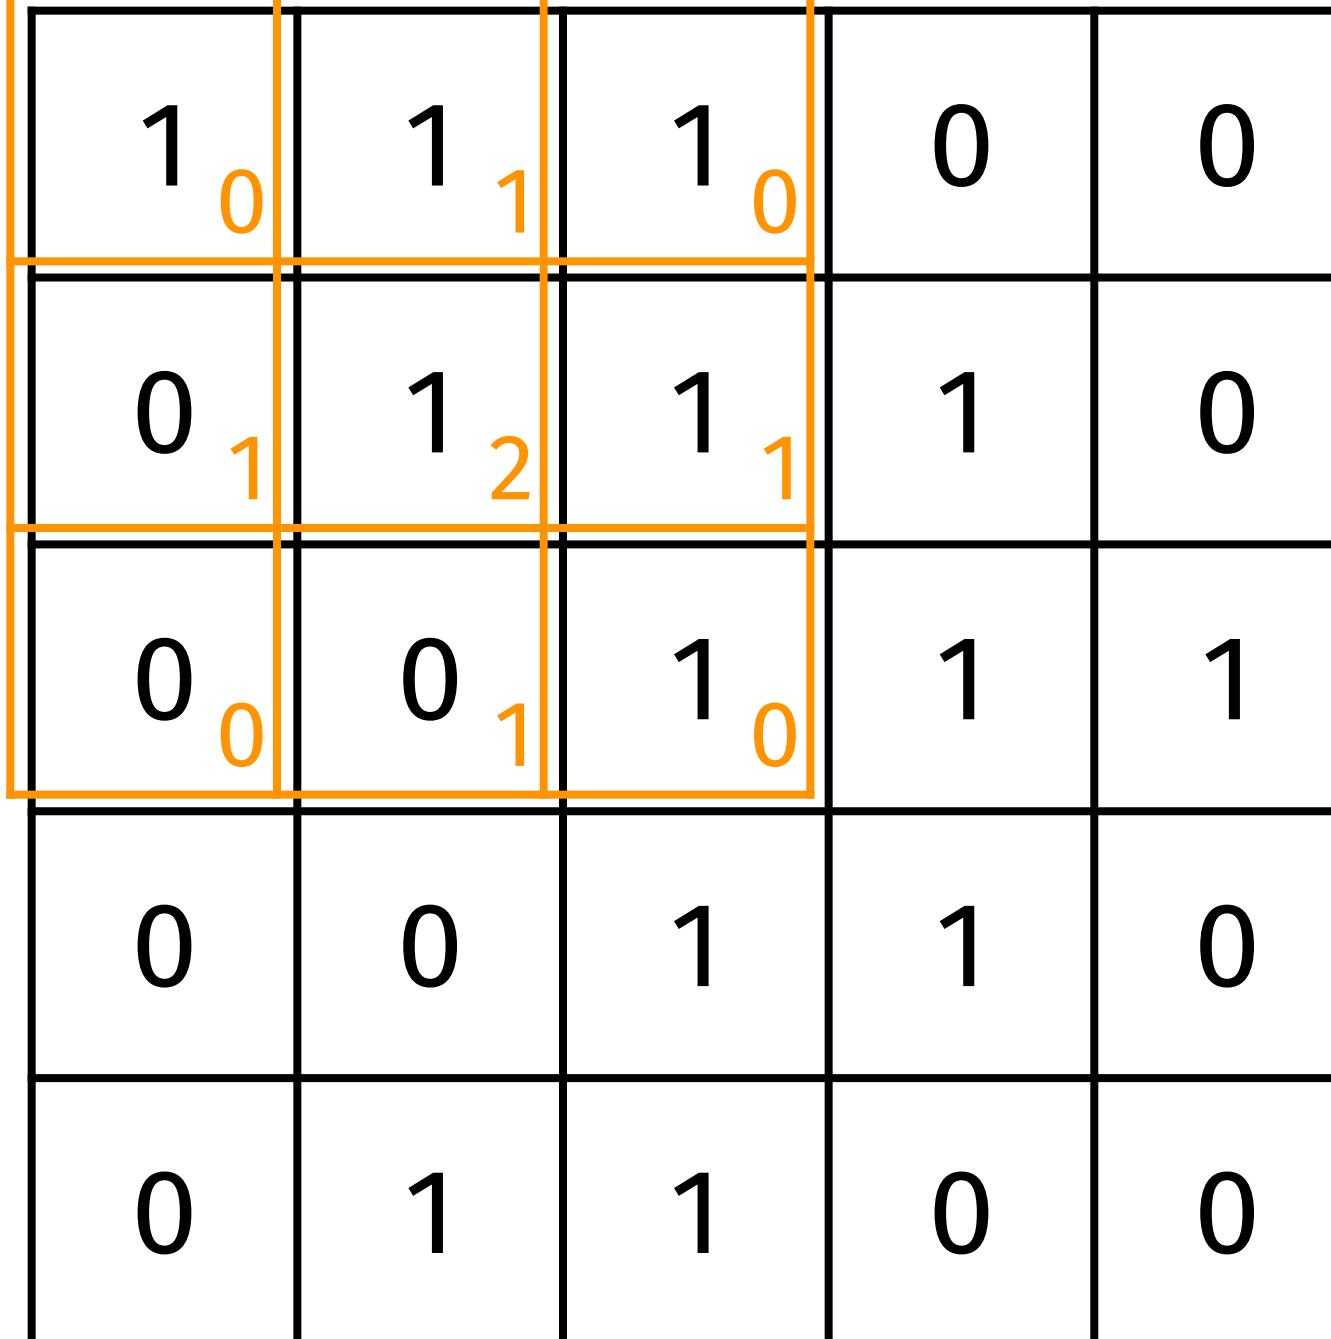
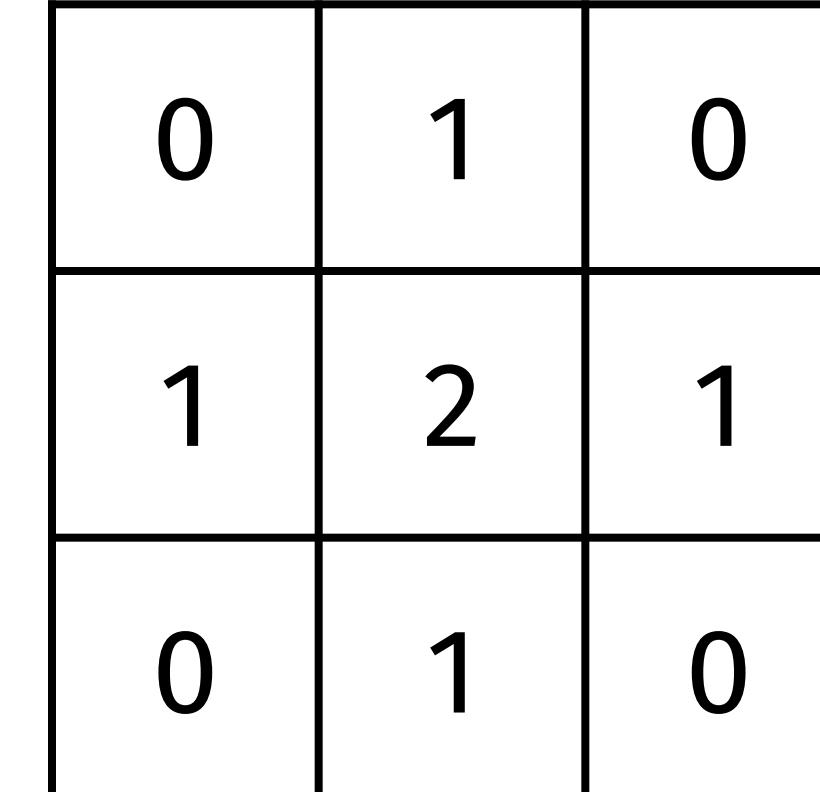
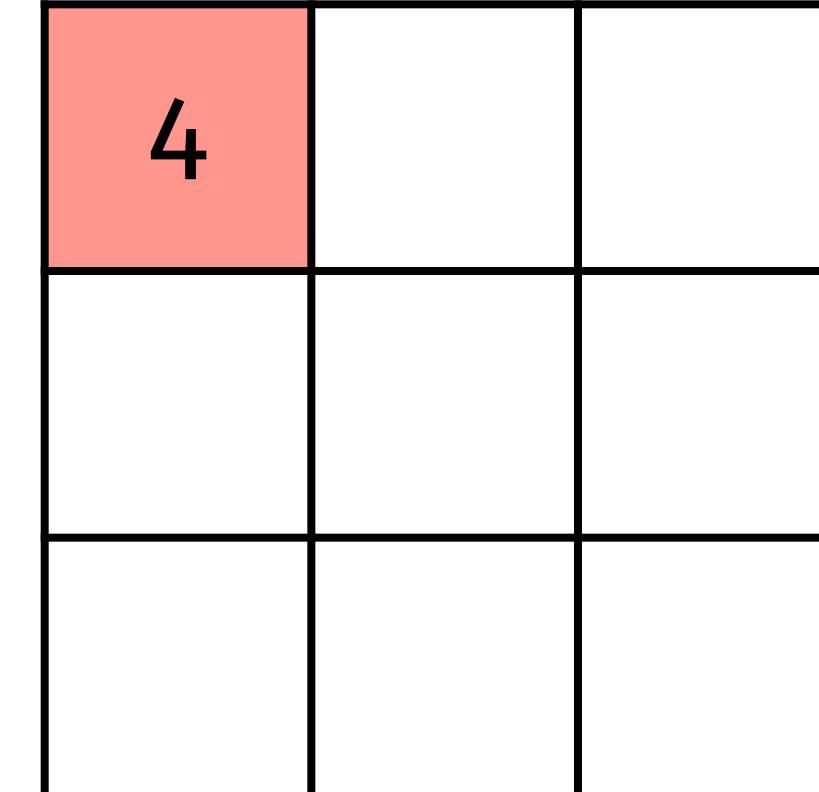
$$(I * K)(i, j) = \sum_{m=-\infty}^{\infty} \sum_{n=-\infty}^{\infty} I(m, n)K(i - m, j - n)$$

# Image convolutions

<b>Image <math>I</math></b>	<b>Kernel <math>K</math></b>	<b>Convolution <math>I * K</math></b>																																							
<table border="1" style="margin-left: auto; margin-right: auto;"><tr><td>1</td><td>0</td><td>1</td><td>1</td><td>0</td><td>0</td></tr><tr><td>0</td><td>1</td><td>1</td><td>2</td><td>1</td><td>0</td></tr><tr><td>0</td><td>0</td><td>0</td><td>1</td><td>1</td><td>1</td></tr><tr><td>0</td><td>0</td><td>1</td><td>1</td><td>1</td><td>0</td></tr><tr><td>0</td><td>1</td><td>1</td><td>0</td><td>0</td><td>0</td></tr></table>	1	0	1	1	0	0	0	1	1	2	1	0	0	0	0	1	1	1	0	0	1	1	1	0	0	1	1	0	0	0	*	<table border="1" style="margin-left: auto; margin-right: auto;"><tr><td>0</td><td>1</td><td>0</td></tr><tr><td>1</td><td>2</td><td>1</td></tr><tr><td>0</td><td>1</td><td>0</td></tr></table>	0	1	0	1	2	1	0	1	0
1	0	1	1	0	0																																				
0	1	1	2	1	0																																				
0	0	0	1	1	1																																				
0	0	1	1	1	0																																				
0	1	1	0	0	0																																				
0	1	0																																							
1	2	1																																							
0	1	0																																							
	=	<table border="1" style="margin-left: auto; margin-right: auto;"><tr><td style="background-color: #ff9999;"></td><td></td><td></td></tr><tr><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td></tr></table>																																							

$$(I * K)(i, j) = \sum_{m=-\infty}^{\infty} \sum_{n=-\infty}^{\infty} I(m, n)K(i - m, j - n)$$

# Image convolutions

<b>Image <math>I</math></b>		<b>Kernel <math>K</math></b>
	*	
	=	

$$(I * K)(i, j) = \sum_{m=-\infty}^{\infty} \sum_{n=-\infty}^{\infty} I(m, n)K(i - m, j - n)$$

# Image convolutions

<b>Image <math>I</math></b>																																
<table border="1" style="border-collapse: collapse; text-align: center;"><tr><td>1</td><td>1</td><td>0</td><td>1</td><td>0</td><td>0</td></tr><tr><td>0</td><td>1</td><td>1</td><td>2</td><td>1</td><td>0</td></tr><tr><td>0</td><td>0</td><td>0</td><td>1</td><td>1</td><td>1</td></tr><tr><td>0</td><td>0</td><td>1</td><td>1</td><td>1</td><td>0</td></tr><tr><td>0</td><td>1</td><td>1</td><td>0</td><td>0</td><td>0</td></tr></table>	1	1	0	1	0	0	0	1	1	2	1	0	0	0	0	1	1	1	0	0	1	1	1	0	0	1	1	0	0	0	$*$	
1	1	0	1	0	0																											
0	1	1	2	1	0																											
0	0	0	1	1	1																											
0	0	1	1	1	0																											
0	1	1	0	0	0																											
<b>Kernel <math>K</math></b>																																
<table border="1" style="border-collapse: collapse; text-align: center;"><tr><td>0</td><td>1</td><td>0</td></tr><tr><td>1</td><td>2</td><td>1</td></tr><tr><td>0</td><td>1</td><td>0</td></tr></table>	0	1	0	1	2	1	0	1	0	$=$																						
0	1	0																														
1	2	1																														
0	1	0																														
<b>Convolution <math>I * K</math></b>																																
<table border="1" style="border-collapse: collapse; text-align: center;"><tr><td>4</td><td></td><td></td></tr><tr><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td></tr></table>	4																															
4																																

$$(I * K)(i, j) = \sum_{m=-\infty}^{\infty} \sum_{n=-\infty}^{\infty} I(m, n)K(i - m, j - n)$$

# Image convolutions

Image $I$					Kernel $K$			Convolution $I * K$		
1	1	0	1	0	0	1	0	0	0	
0	1	1	2	1	0	1	2	1	0	
0	0	0	1	1	1	0	1	0	1	
0	0	1	1	0	0	1	1	0	0	
0	1	1	0	0	0	1	1	0	0	

\*

0	1	0
1	2	1
0	1	0

=

4	6	0
0	0	0
0	0	0

$$(I * K)(i, j) = \sum_{m=-\infty}^{\infty} \sum_{n=-\infty}^{\infty} I(m, n)K(i - m, j - n)$$

# Image convolutions

Image $I$				
1	1	1 0	0 1	0 0
0	1	1 1	1 2	0 1
0	0	1 0	1 1	1 0
0	0	1	1	0
0	1	1	0	0

\*

Kernel $K$		
0	1	0
1	2	1
0	1	0

=

4	6	

$$(I * K)(i, j) = \sum_{m=-\infty}^{\infty} \sum_{n=-\infty}^{\infty} I(m, n)K(i - m, j - n)$$

# Image convolutions

<b>Image <math>I</math></b>	*	<b>Kernel <math>K</math></b>	<b>Convolution <math>I * K</math></b>																																																
<table border="1" style="border-collapse: collapse; width: 100px; height: 100px;"> <tr><td>1</td><td>1</td><td>1</td><td>0</td><td>0</td><td>0</td></tr> <tr><td>0</td><td>1</td><td>1</td><td>1</td><td>2</td><td>0</td></tr> <tr><td>0</td><td>0</td><td>1</td><td>1</td><td>1</td><td>0</td></tr> <tr><td>0</td><td>0</td><td>1</td><td>1</td><td>0</td><td></td></tr> <tr><td>0</td><td>1</td><td>1</td><td>0</td><td></td><td></td></tr> </table>	1	1	1	0	0	0	0	1	1	1	2	0	0	0	1	1	1	0	0	0	1	1	0		0	1	1	0			*	<table border="1" style="border-collapse: collapse; width: 100px; height: 100px;"> <tr><td>0</td><td>1</td><td>0</td></tr> <tr><td>1</td><td>2</td><td>1</td></tr> <tr><td>0</td><td>1</td><td>0</td></tr> </table>	0	1	0	1	2	1	0	1	0	<table border="1" style="border-collapse: collapse; width: 100px; height: 100px;"> <tr><td>4</td><td>6</td><td>4</td></tr> <tr><td></td><td></td><td></td></tr> <tr><td></td><td></td><td></td></tr> </table>	4	6	4						
1	1	1	0	0	0																																														
0	1	1	1	2	0																																														
0	0	1	1	1	0																																														
0	0	1	1	0																																															
0	1	1	0																																																
0	1	0																																																	
1	2	1																																																	
0	1	0																																																	
4	6	4																																																	

The diagram illustrates the convolution operation  $I * K$ . The input image  $I$  is a 5x5 matrix with values [1,1,1,0,0], [0,1,1,1,2], [0,0,1,1,1], [0,0,1,1,0], and [0,1,1,0,0]. The kernel  $K$  is a 3x3 matrix with values [0,1,0], [1,2,1], and [0,1,0]. The convolution result is shown as a 3x3 matrix with values [4,6,4], [0,0,0], and [0,0,0]. The highlighted 3x3 submatrix in  $I$  represents the receptive field of the central element in the result matrix. The intermediate steps show the calculation of the result for each element in the output matrix.

$$(I * K)(i, j) = \sum_{m=-\infty}^{\infty} \sum_{n=-\infty}^{\infty} I(m, n)K(i - m, j - n)$$

# Image convolutions

<b>Image <math>I</math></b>	*	<b>Kernel <math>K</math></b>	<b>Convolution <math>I * K</math></b>																																											
<table border="1" style="border-collapse: collapse; text-align: center;"> <tr><td>1</td><td>1</td><td>1</td><td>0</td><td>0</td></tr> <tr><td>0</td><td>1</td><td>1</td><td>1</td><td>0</td></tr> <tr><td>0</td><td>0</td><td>1</td><td>1</td><td>1</td></tr> <tr><td>0</td><td>0</td><td>1</td><td>1</td><td>0</td></tr> <tr><td>0</td><td>1</td><td>1</td><td>0</td><td>0</td></tr> </table>	1	1	1	0	0	0	1	1	1	0	0	0	1	1	1	0	0	1	1	0	0	1	1	0	0	*	<table border="1" style="border-collapse: collapse; text-align: center;"> <tr><td>0</td><td>1</td><td>0</td></tr> <tr><td>1</td><td>2</td><td>1</td></tr> <tr><td>0</td><td>1</td><td>0</td></tr> </table>	0	1	0	1	2	1	0	1	0	$=$ <table border="1" style="border-collapse: collapse; text-align: center;"> <tr><td>4</td><td>6</td><td>4</td></tr> <tr><td></td><td></td><td></td></tr> <tr><td></td><td></td><td></td></tr> </table>	4	6	4						
1	1	1	0	0																																										
0	1	1	1	0																																										
0	0	1	1	1																																										
0	0	1	1	0																																										
0	1	1	0	0																																										
0	1	0																																												
1	2	1																																												
0	1	0																																												
4	6	4																																												

$$(I * K)(i, j) = \sum_{m=-\infty}^{\infty} \sum_{n=-\infty}^{\infty} I(m, n)K(i - m, j - n)$$

# Image convolutions

<b>Image <math>I</math></b>	<b>Kernel <math>K</math></b>	<b>Convolution <math>I * K</math></b>																																		
<table border="1" style="width: 100%; border-collapse: collapse;"><tr><td>1</td><td>1</td><td>1</td><td>0</td><td>0</td></tr><tr><td>0 0</td><td>1 1</td><td>1 0</td><td>1</td><td>0</td></tr><tr><td>0 1</td><td>0 2</td><td>1 1</td><td>1</td><td>1</td></tr><tr><td>0 0</td><td>0 1</td><td>1 0</td><td>1</td><td>0</td></tr><tr><td>0</td><td>1</td><td>1</td><td>0</td><td>0</td></tr></table>	1	1	1	0	0	0 0	1 1	1 0	1	0	0 1	0 2	1 1	1	1	0 0	0 1	1 0	1	0	0	1	1	0	0	*	<table border="1" style="width: 100%; border-collapse: collapse;"><tr><td>0</td><td>1</td><td>0</td></tr><tr><td>1</td><td>2</td><td>1</td></tr><tr><td>0</td><td>1</td><td>0</td></tr></table>	0	1	0	1	2	1	0	1	0
1	1	1	0	0																																
0 0	1 1	1 0	1	0																																
0 1	0 2	1 1	1	1																																
0 0	0 1	1 0	1	0																																
0	1	1	0	0																																
0	1	0																																		
1	2	1																																		
0	1	0																																		
	=	<table border="1" style="width: 100%; border-collapse: collapse;"><tr><td>4</td><td>6</td><td>4</td></tr><tr><td>2</td><td></td><td></td></tr><tr><td></td><td></td><td></td></tr></table>	4	6	4	2																														
4	6	4																																		
2																																				

$$(I * K)(i, j) = \sum_{m=-\infty}^{\infty} \sum_{n=-\infty}^{\infty} I(m, n)K(i - m, j - n)$$

# Image convolutions

<b>Image <math>I</math></b>		<b>Kernel <math>K</math></b>																																		
<table border="1" style="border-collapse: collapse; text-align: center;"><tr><td>1</td><td>1</td><td>1</td><td>0</td><td>0</td></tr><tr><td>0</td><td>1</td><td>1</td><td>1</td><td>0</td></tr><tr><td>0</td><td>0</td><td>1</td><td>1</td><td>1</td></tr><tr><td>0</td><td>0</td><td>1</td><td>1</td><td>0</td></tr><tr><td>0</td><td>1</td><td>1</td><td>0</td><td>0</td></tr></table>	1	1	1	0	0	0	1	1	1	0	0	0	1	1	1	0	0	1	1	0	0	1	1	0	0	*	<table border="1" style="border-collapse: collapse; text-align: center;"><tr><td>0</td><td>1</td><td>0</td></tr><tr><td>1</td><td>2</td><td>1</td></tr><tr><td>0</td><td>1</td><td>0</td></tr></table>	0	1	0	1	2	1	0	1	0
1	1	1	0	0																																
0	1	1	1	0																																
0	0	1	1	1																																
0	0	1	1	0																																
0	1	1	0	0																																
0	1	0																																		
1	2	1																																		
0	1	0																																		
		=																																		
		<table border="1" style="border-collapse: collapse; text-align: center;"><tr><td>4</td><td>6</td><td>4</td></tr><tr><td>2</td><td>5</td><td>6</td></tr><tr><td>2</td><td>5</td><td>4</td></tr></table>	4	6	4	2	5	6	2	5	4																									
4	6	4																																		
2	5	6																																		
2	5	4																																		

$$(I * K)(i, j) = \sum_{m=-\infty}^{\infty} \sum_{n=-\infty}^{\infty} I(m, n)K(i - m, j - n)$$

# Example: edge detection

$$\begin{array}{c} \text{Image } I \\ \hline \begin{matrix} 10 & 10 & 10 & 0 & 0 & 0 \\ 10 & 10 & 10 & 0 & 0 & 0 \\ 10 & 10 & 10 & 0 & 0 & 0 \\ 10 & 10 & 10 & 0 & 0 & 0 \\ 10 & 10 & 10 & 0 & 0 & 0 \\ 10 & 10 & 10 & 0 & 0 & 0 \end{matrix} \end{array} * \begin{array}{c} \text{Kernel } K \\ \hline \begin{matrix} & & \\ & & \\ \hline & & \\ & & \\ \hline & & \\ & & \end{matrix} \end{array} = \begin{array}{c} \text{Convolution } I * K \\ \hline \begin{matrix} & & & \\ & & & \\ & & & \\ & & & \\ & & & \\ & & & \end{matrix} \end{array}$$

# Example: edge detection

Image $I$					
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

Kernel  $K$

?	?	?
?	?	?
?	?	?

\*

Convolution  $I * K$


=

# Example: edge detection

<b>Image <math>I</math></b>					
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

<b>Kernel <math>K</math></b>					
1	0	-1			
1	0	-1			
1	0	-1			

<b>Convolution <math>I * K</math></b>					
0	30	30	0		
0	30	30	0		
0	30	30	0		
0	30	30	0		

\*

=

# Example: edge detection

**Image  $I$**

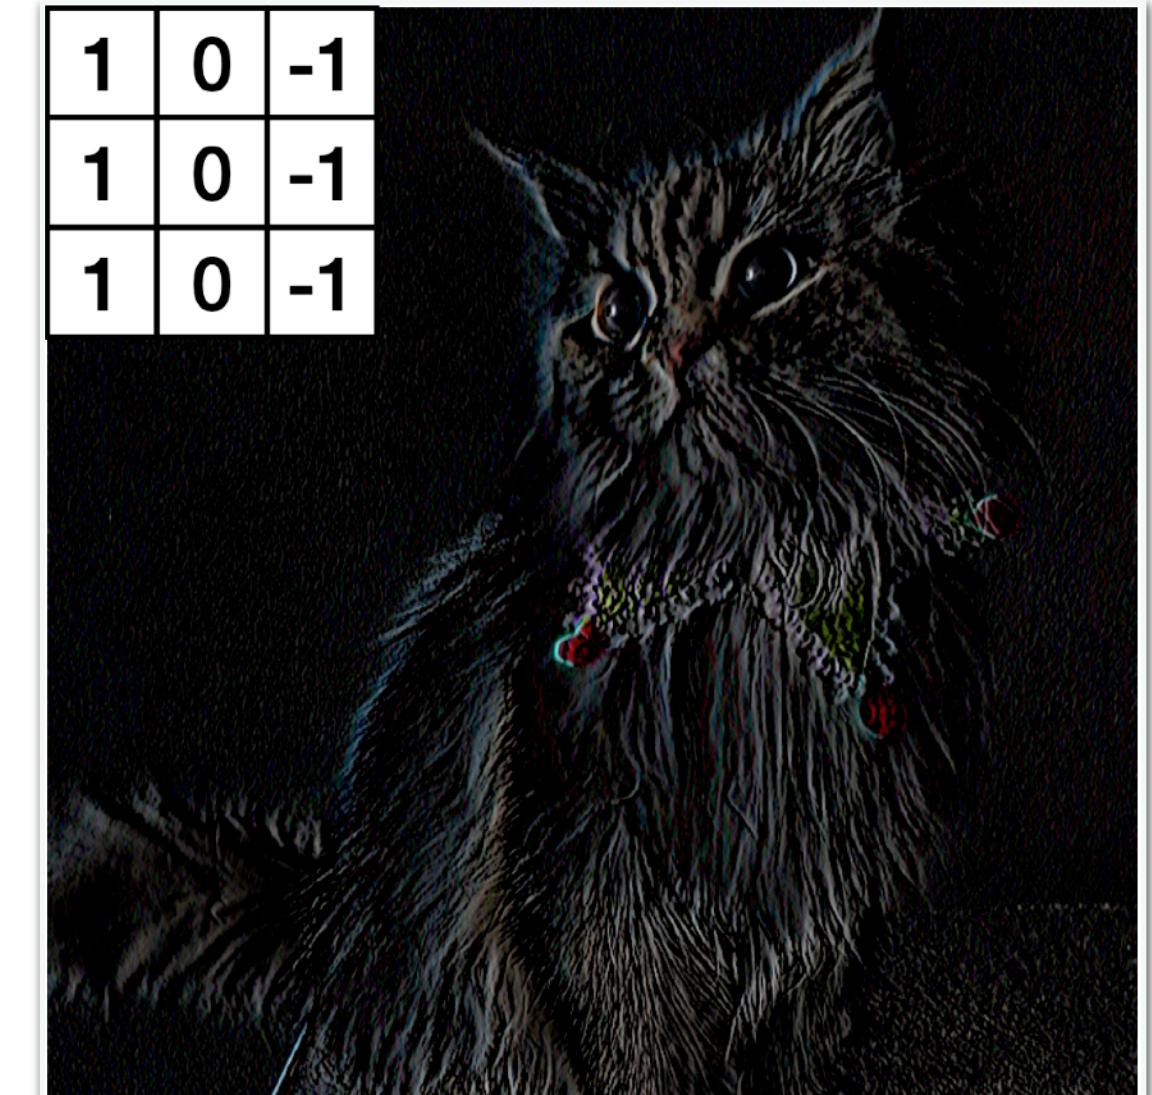


10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

**Kernel  $K$**

$$\begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 1 & 0 & -1 \\ \hline 1 & 0 & -1 \\ \hline \end{array}$$

**Convolution  $I * K$**

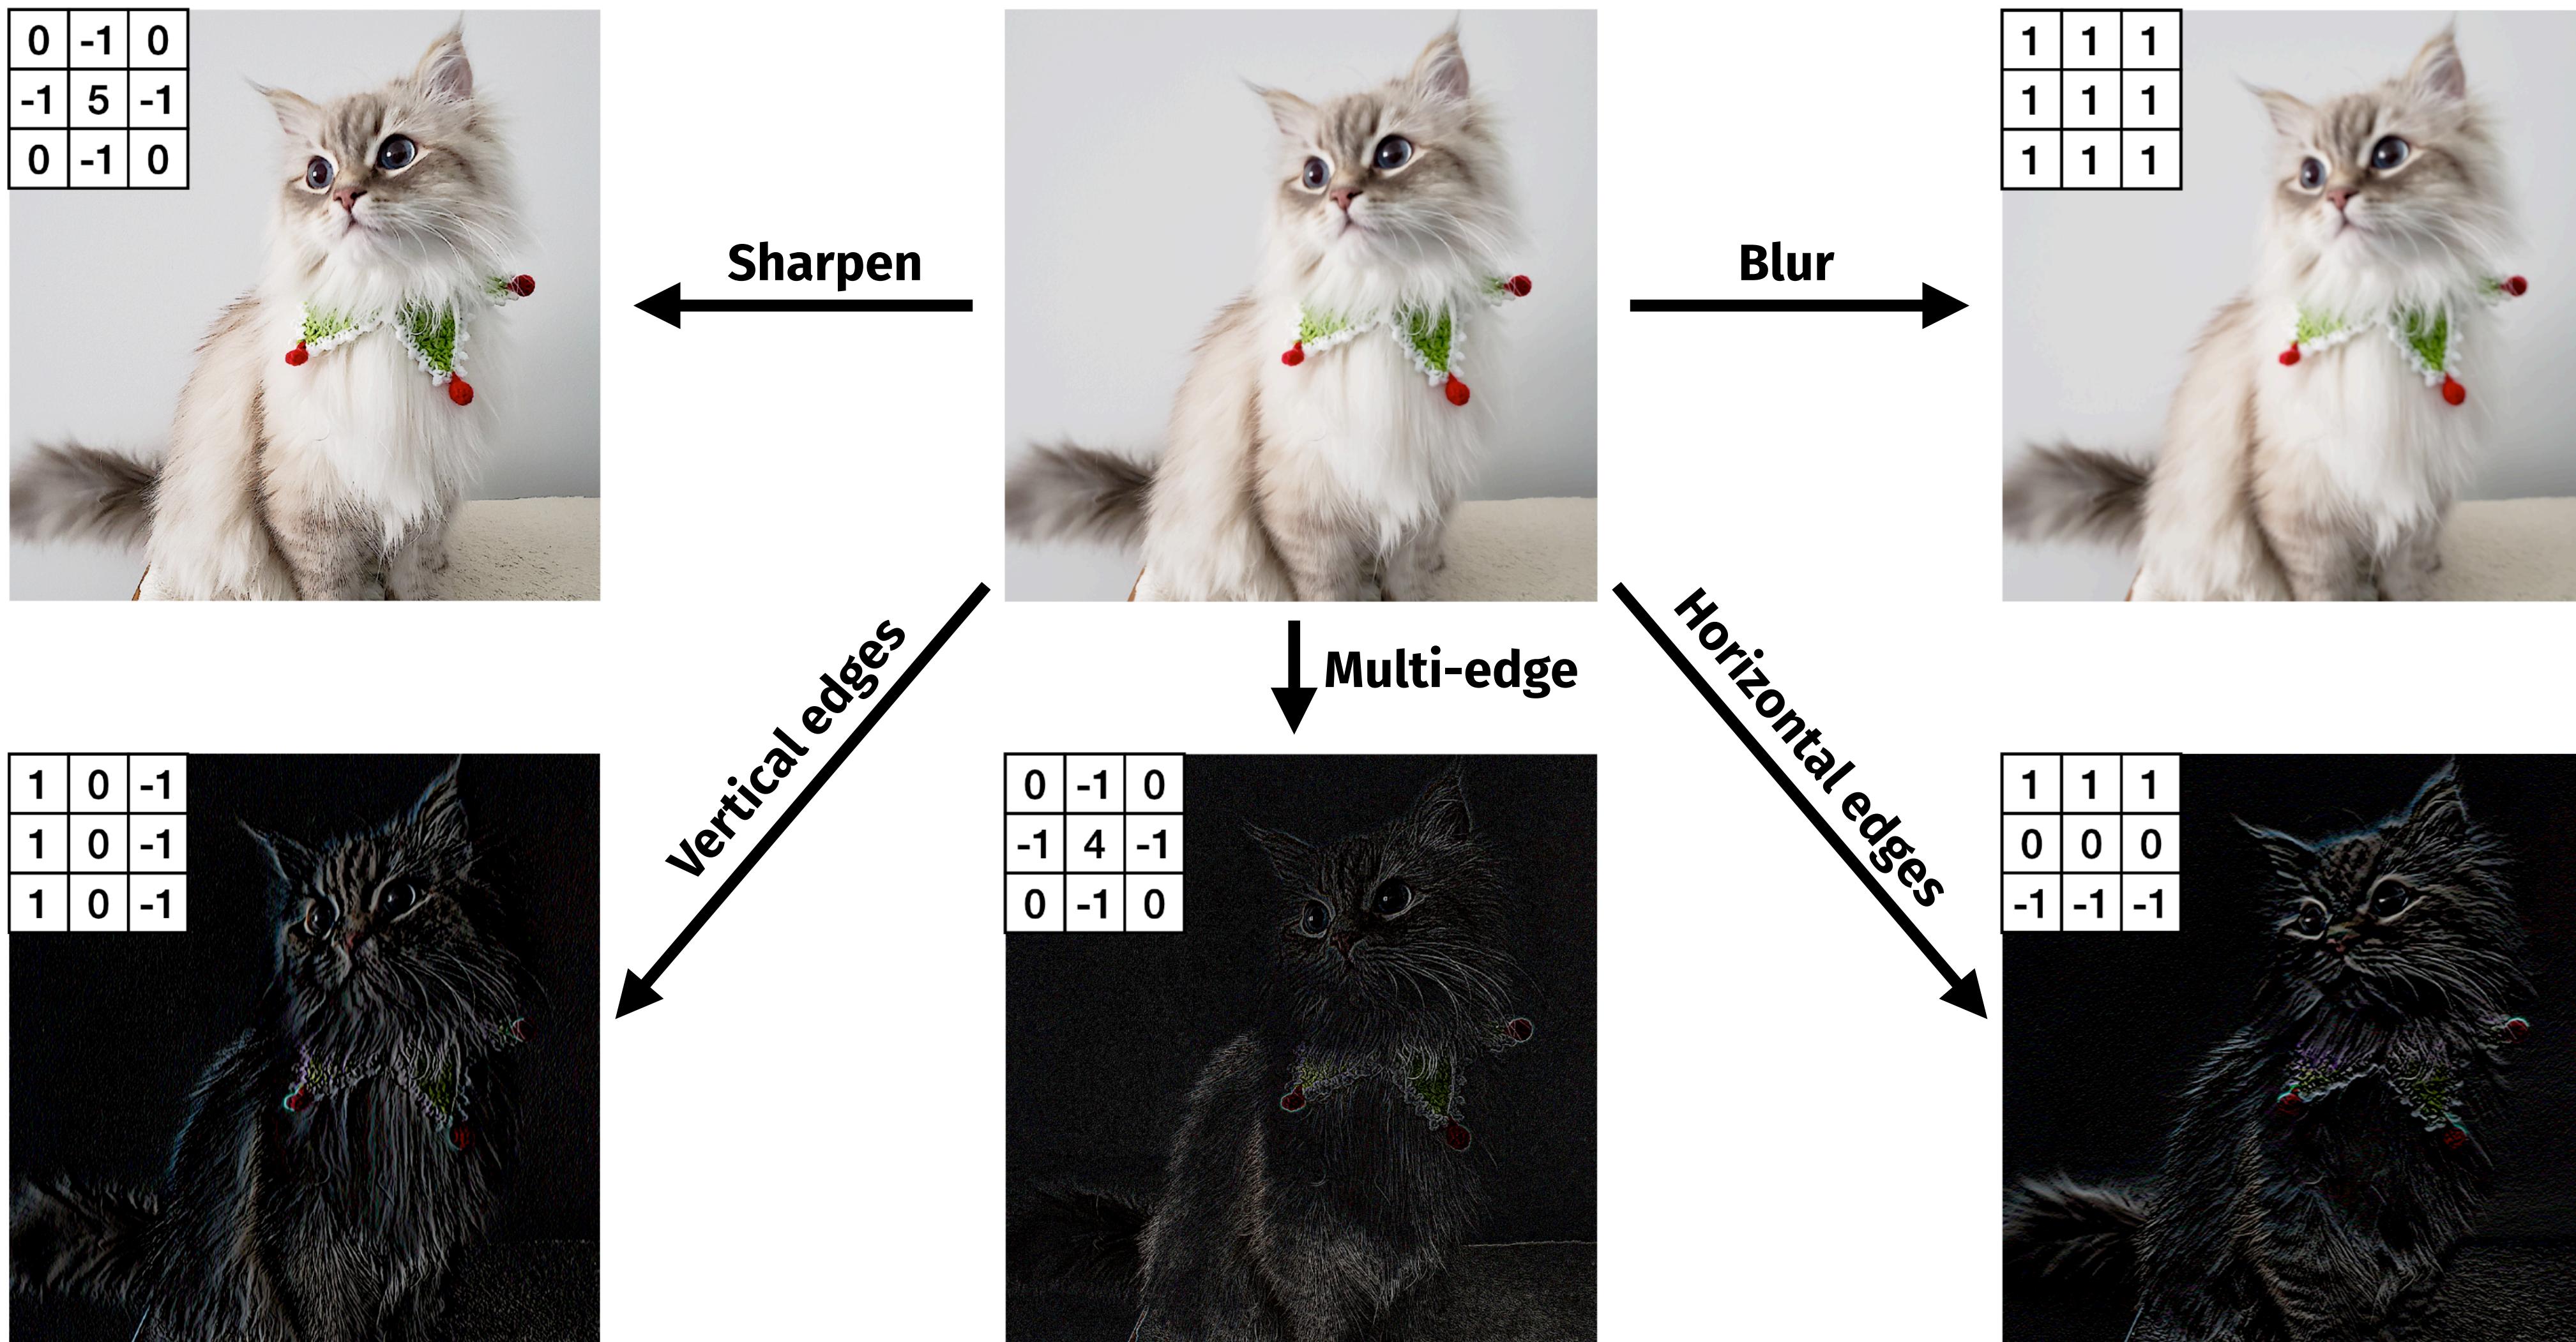


1	0	-1
1	0	-1
1	0	-1

$*$

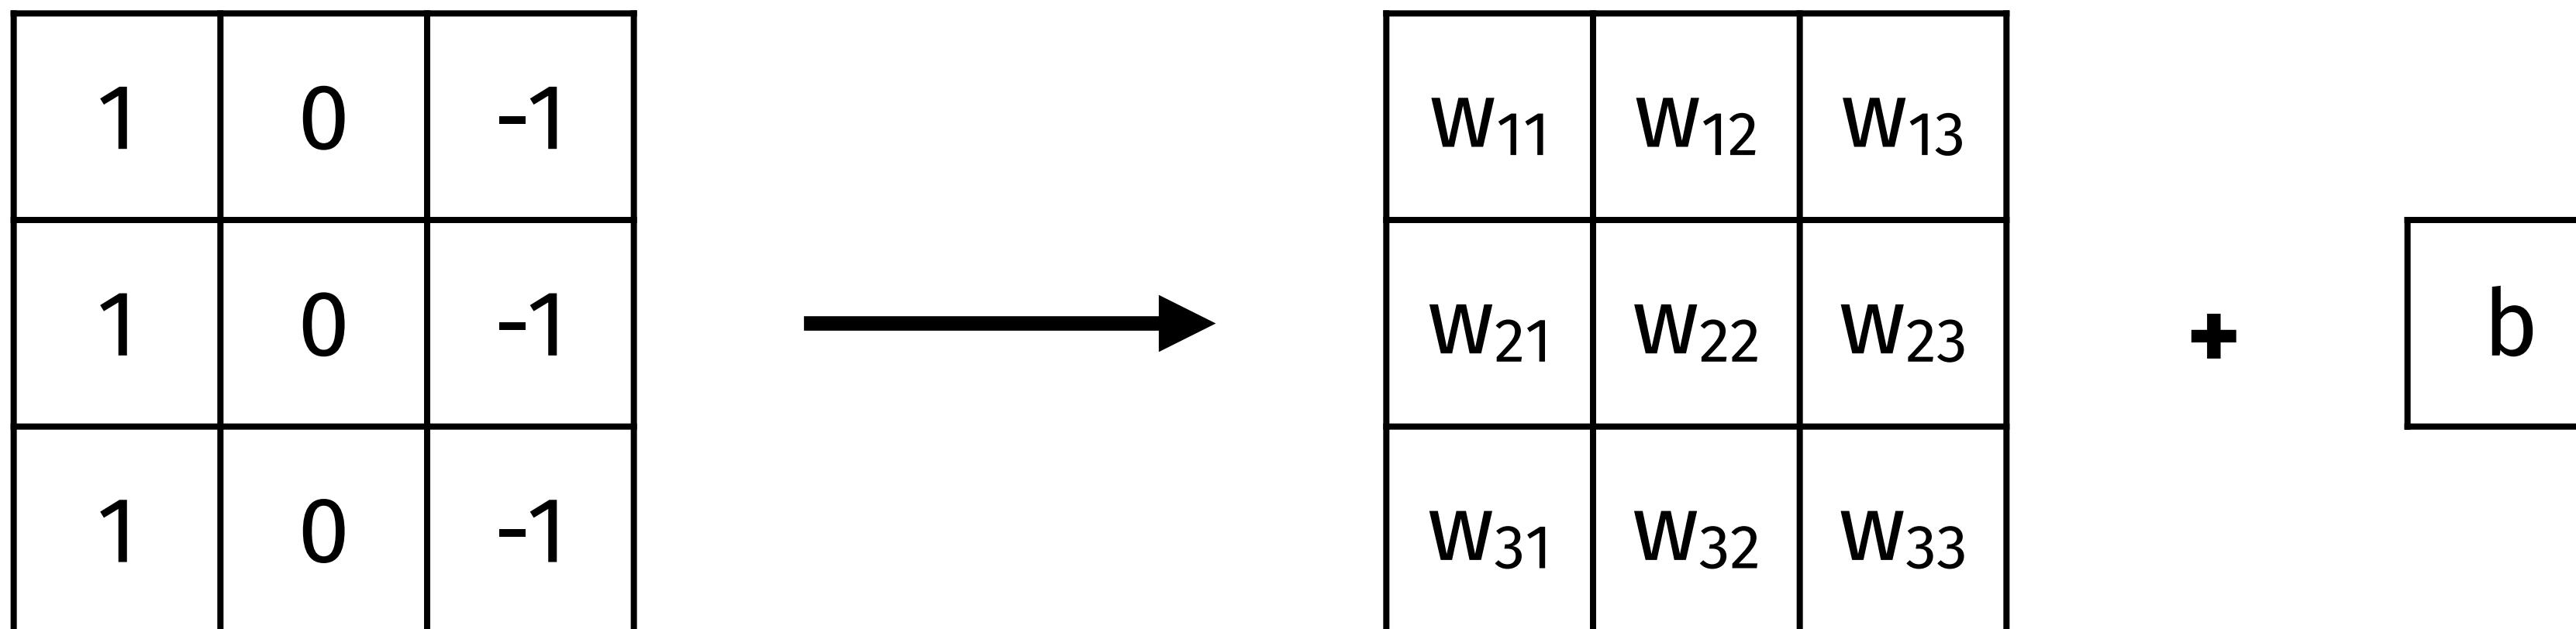
$=$

# Filters for specific purposes



# Machine-learned convolutions

**Idea: Generalize using machine learning**



**Convolutional neural network layer**  
Replace fixed kernel entries with  
**learnable weight and bias**

# Outline

- ① Challenges of image processing
- ② Introduction to convolutions
- ③ **Convolutional layers**
- ④ CNN building blocks
- ⑤ A short history of CNNs
- ⑥ Outlook

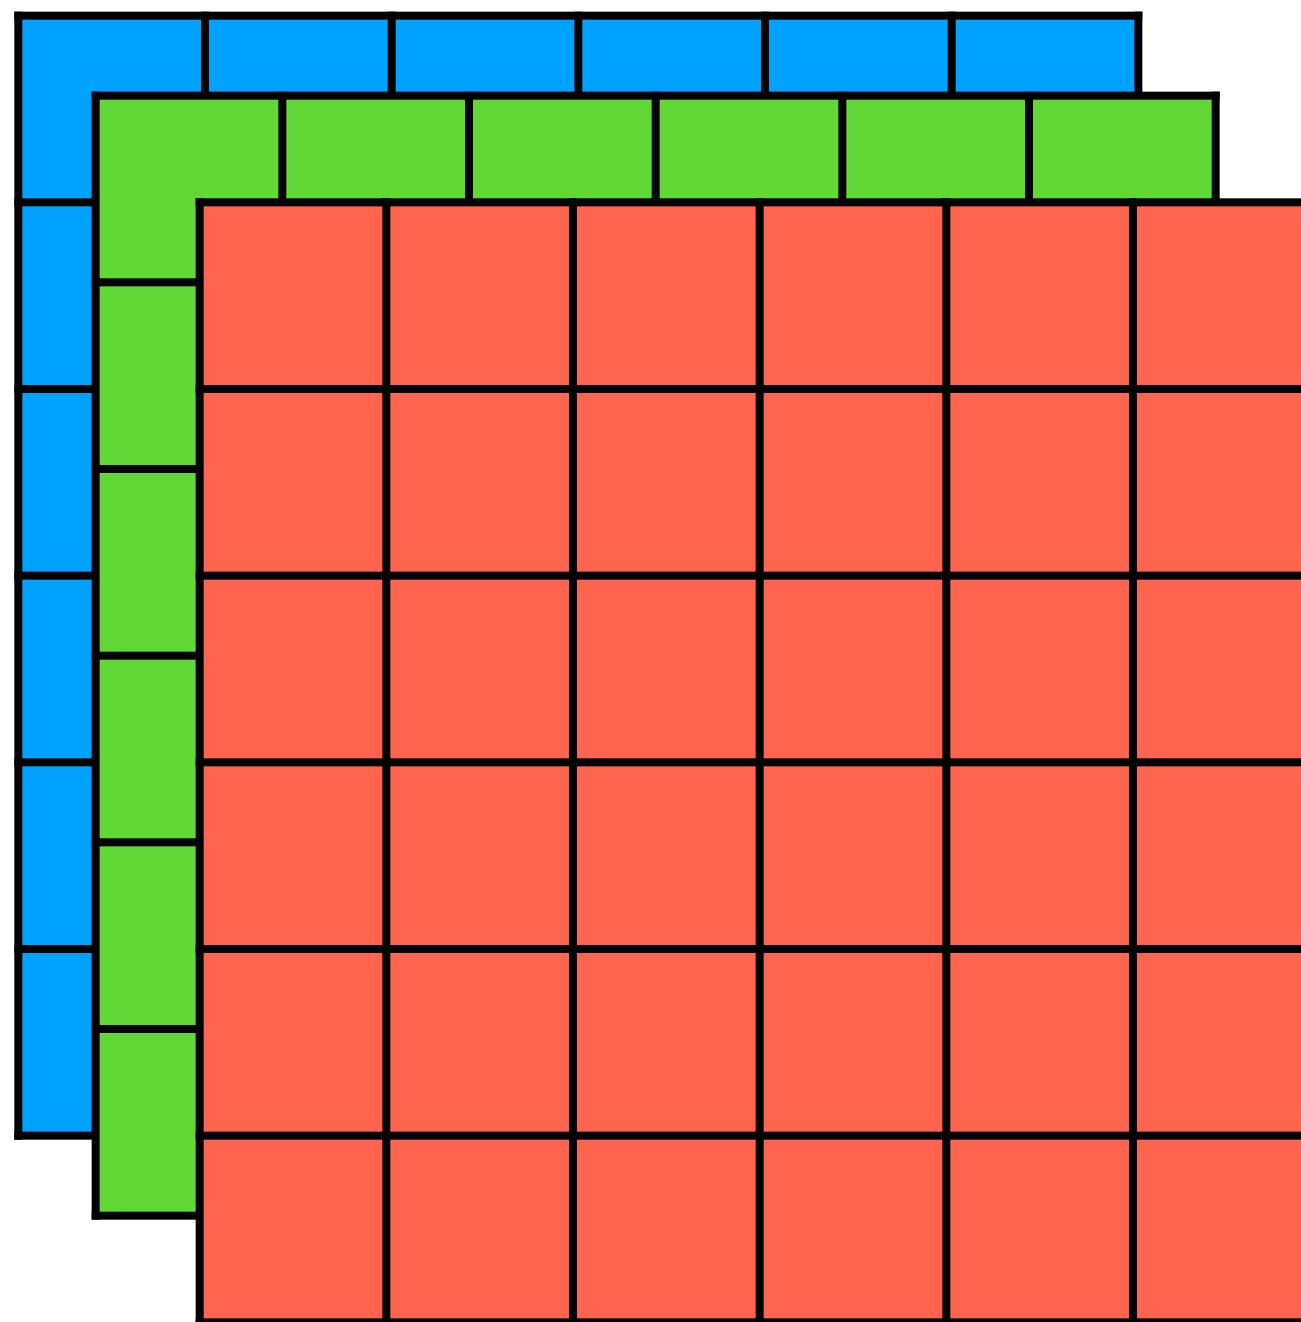
# CNN layers

```
torch.nn.Conv2d(  
    in_channels,  
    out_channels,  
    kernel_size,  
    stride=1,  
    padding=0,  
    dilation=1,  
    groups=1,  
    bias=True,  
    padding_mode='zeros',  
    device=None,  
    dtype=None  
)
```

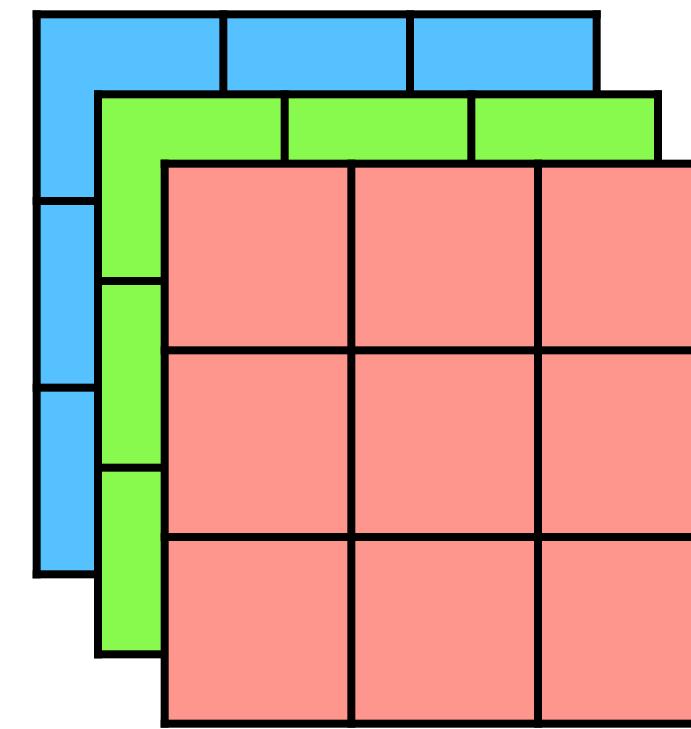
## Most important settings

- ① Channels
- ② Kernel size
- ③ Padding
- ④ Strides

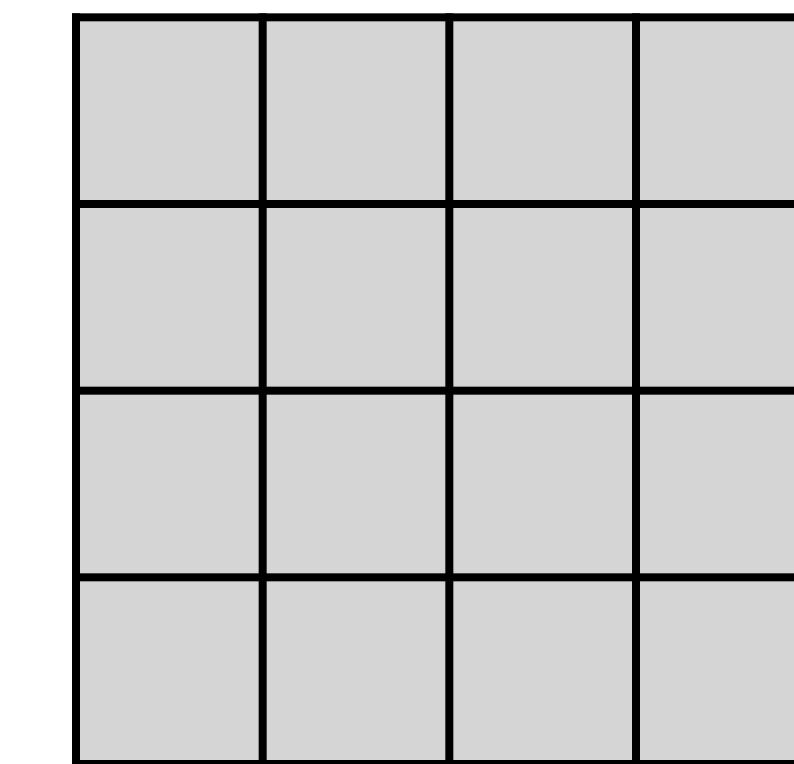
# Channels



\*

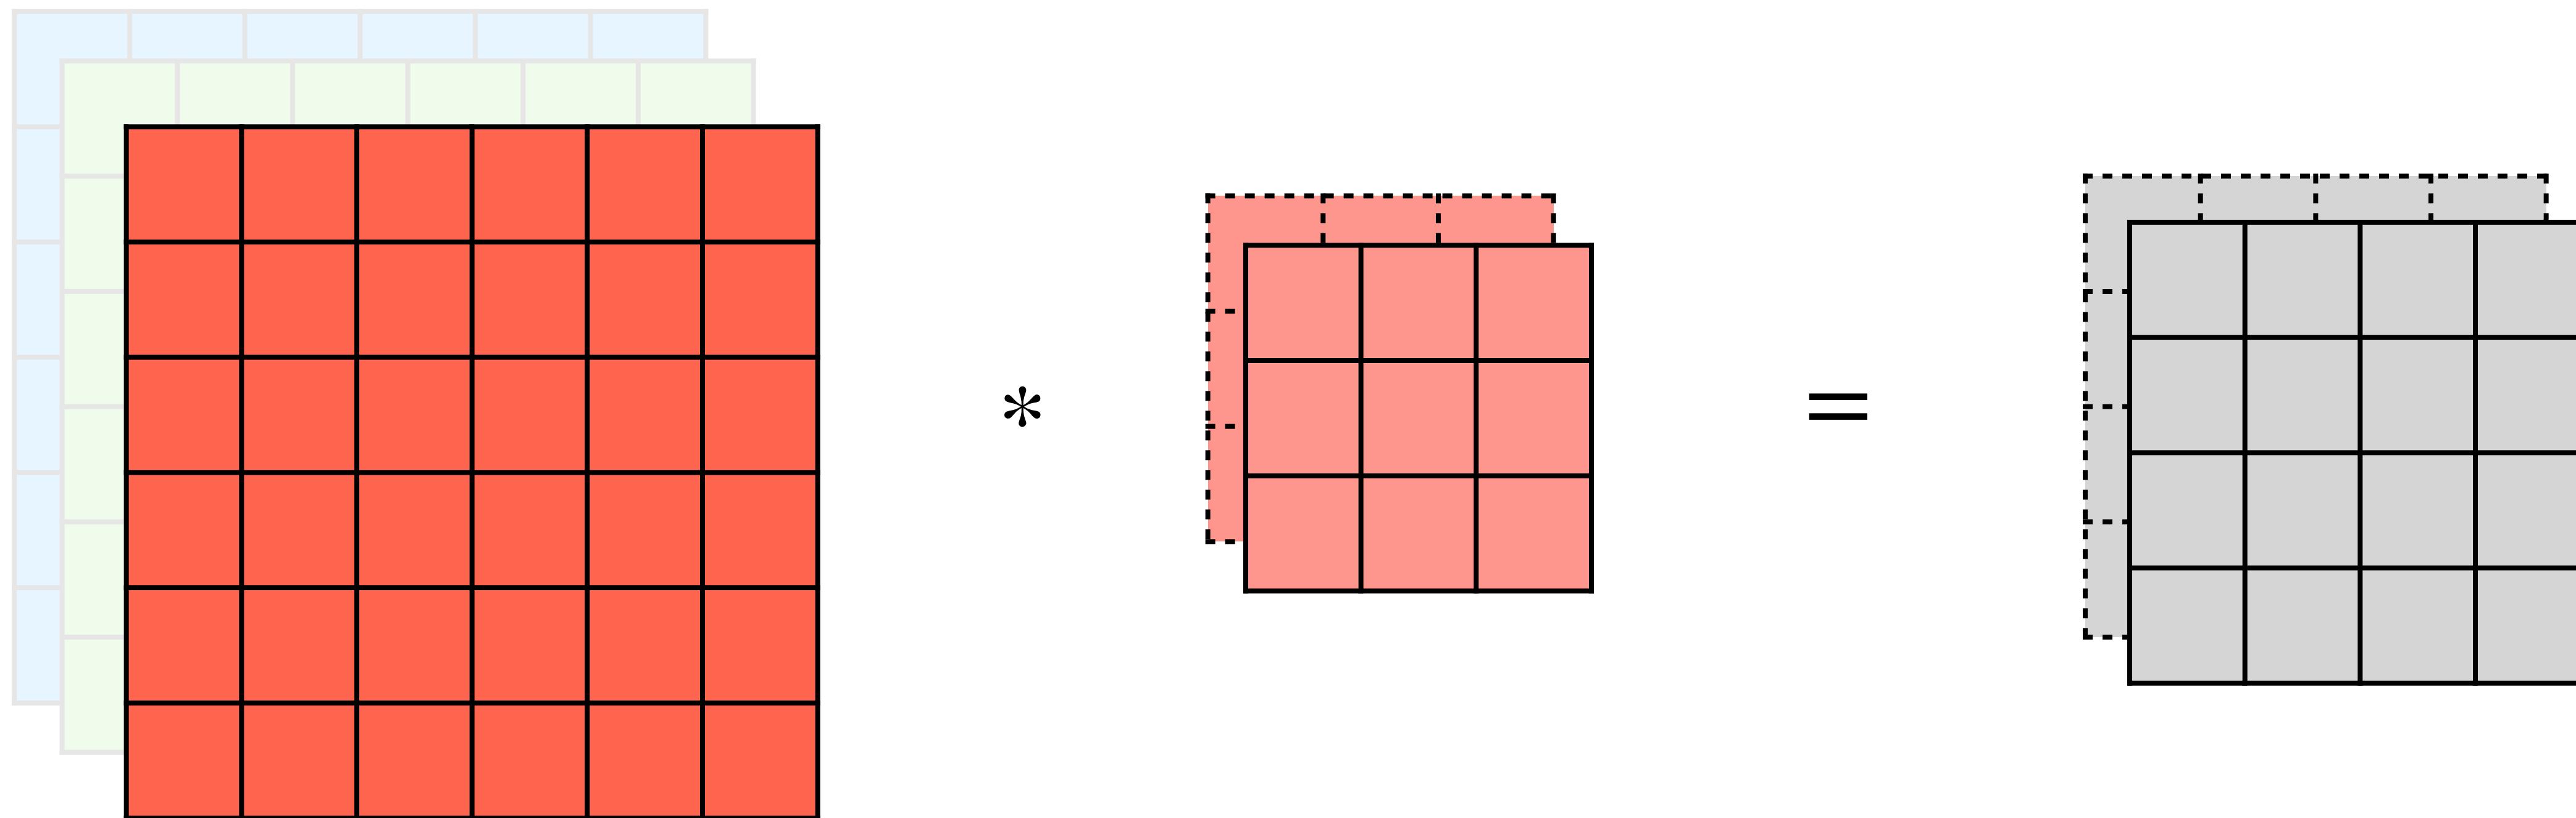


=



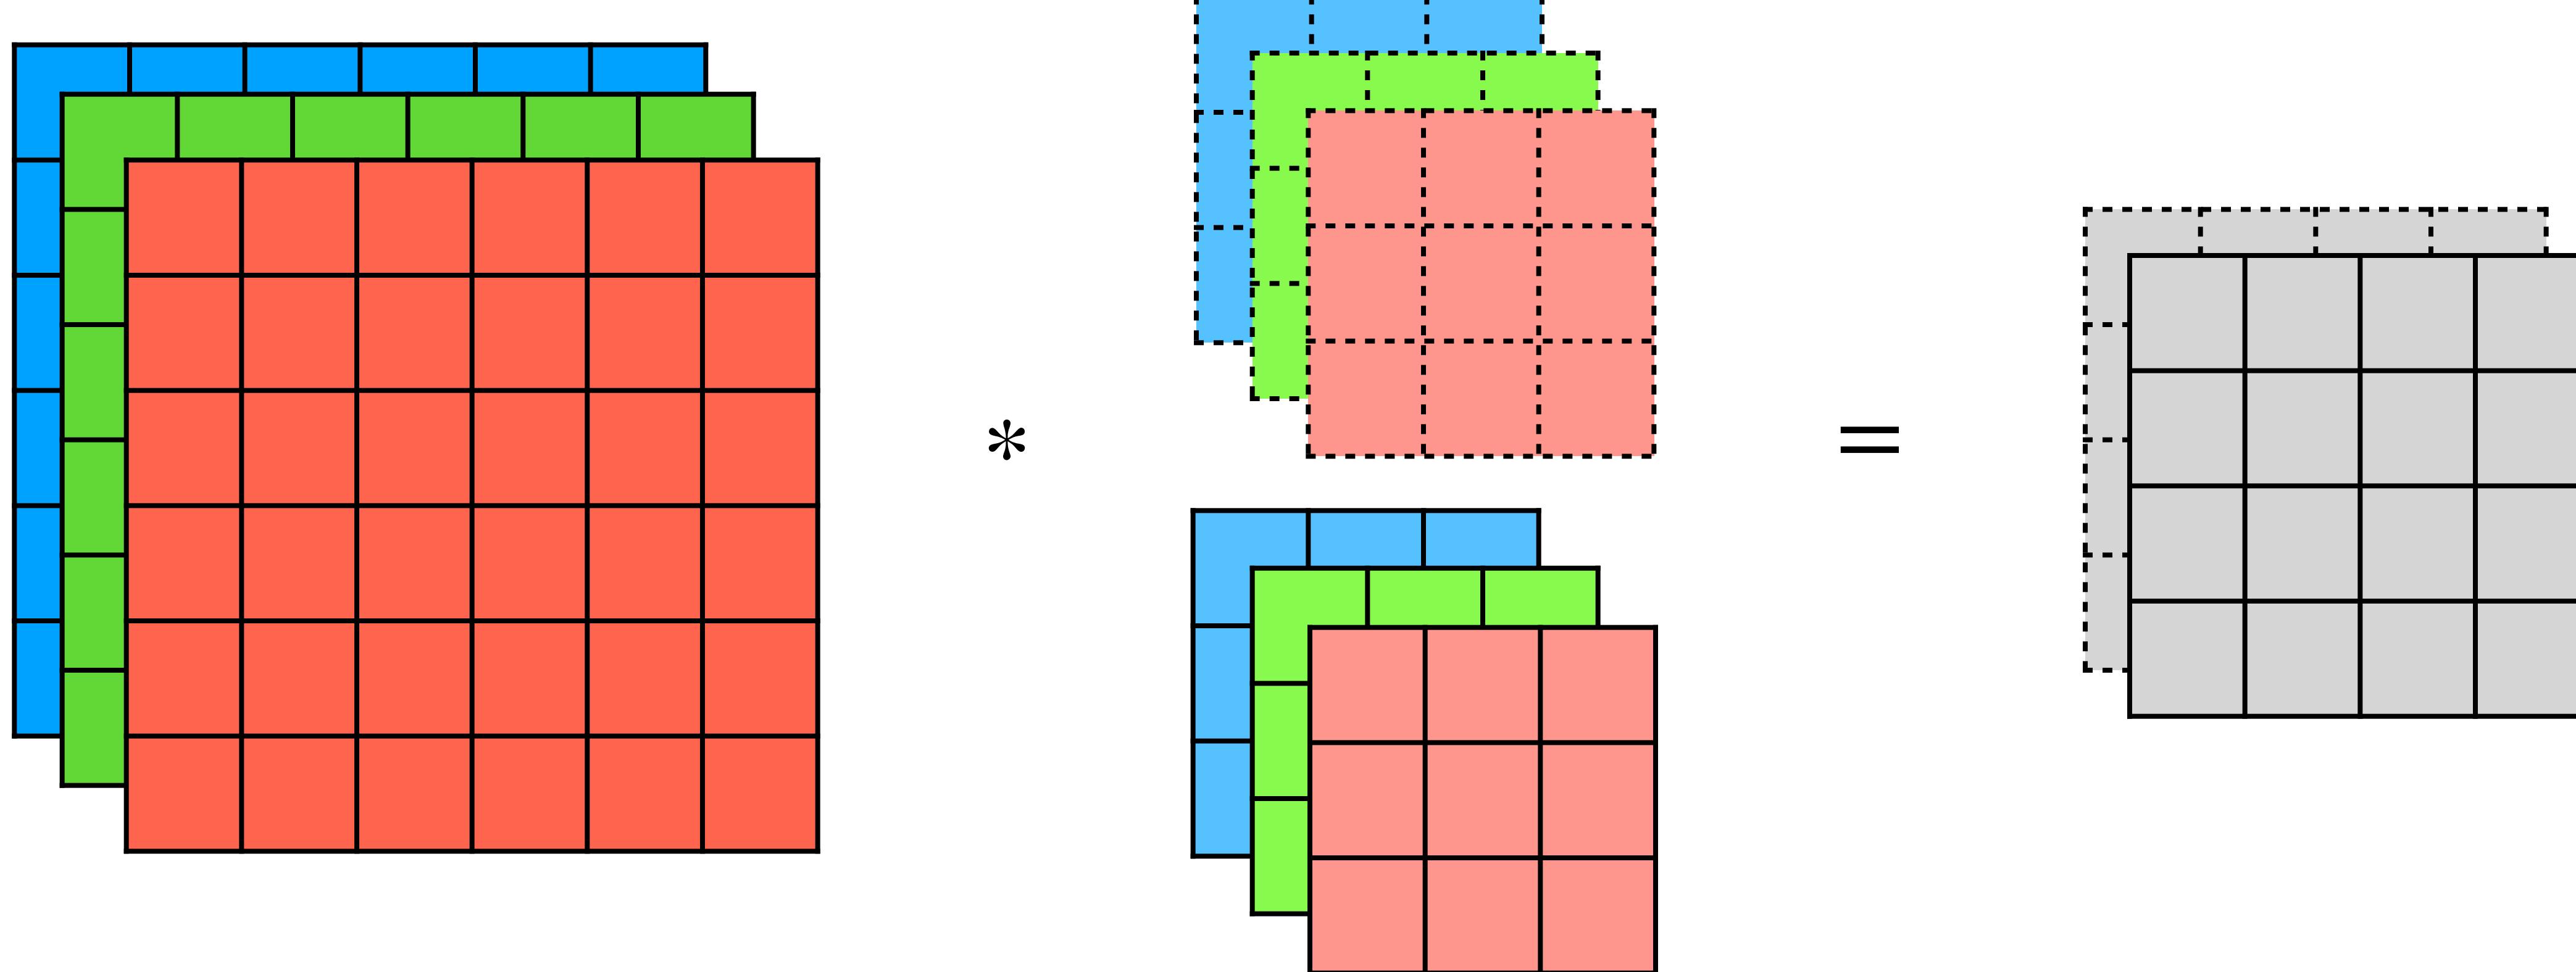
- can have multiple inputs, for example red, green, blue channels
- separate **kernel for each input, result is sum**

# Channels



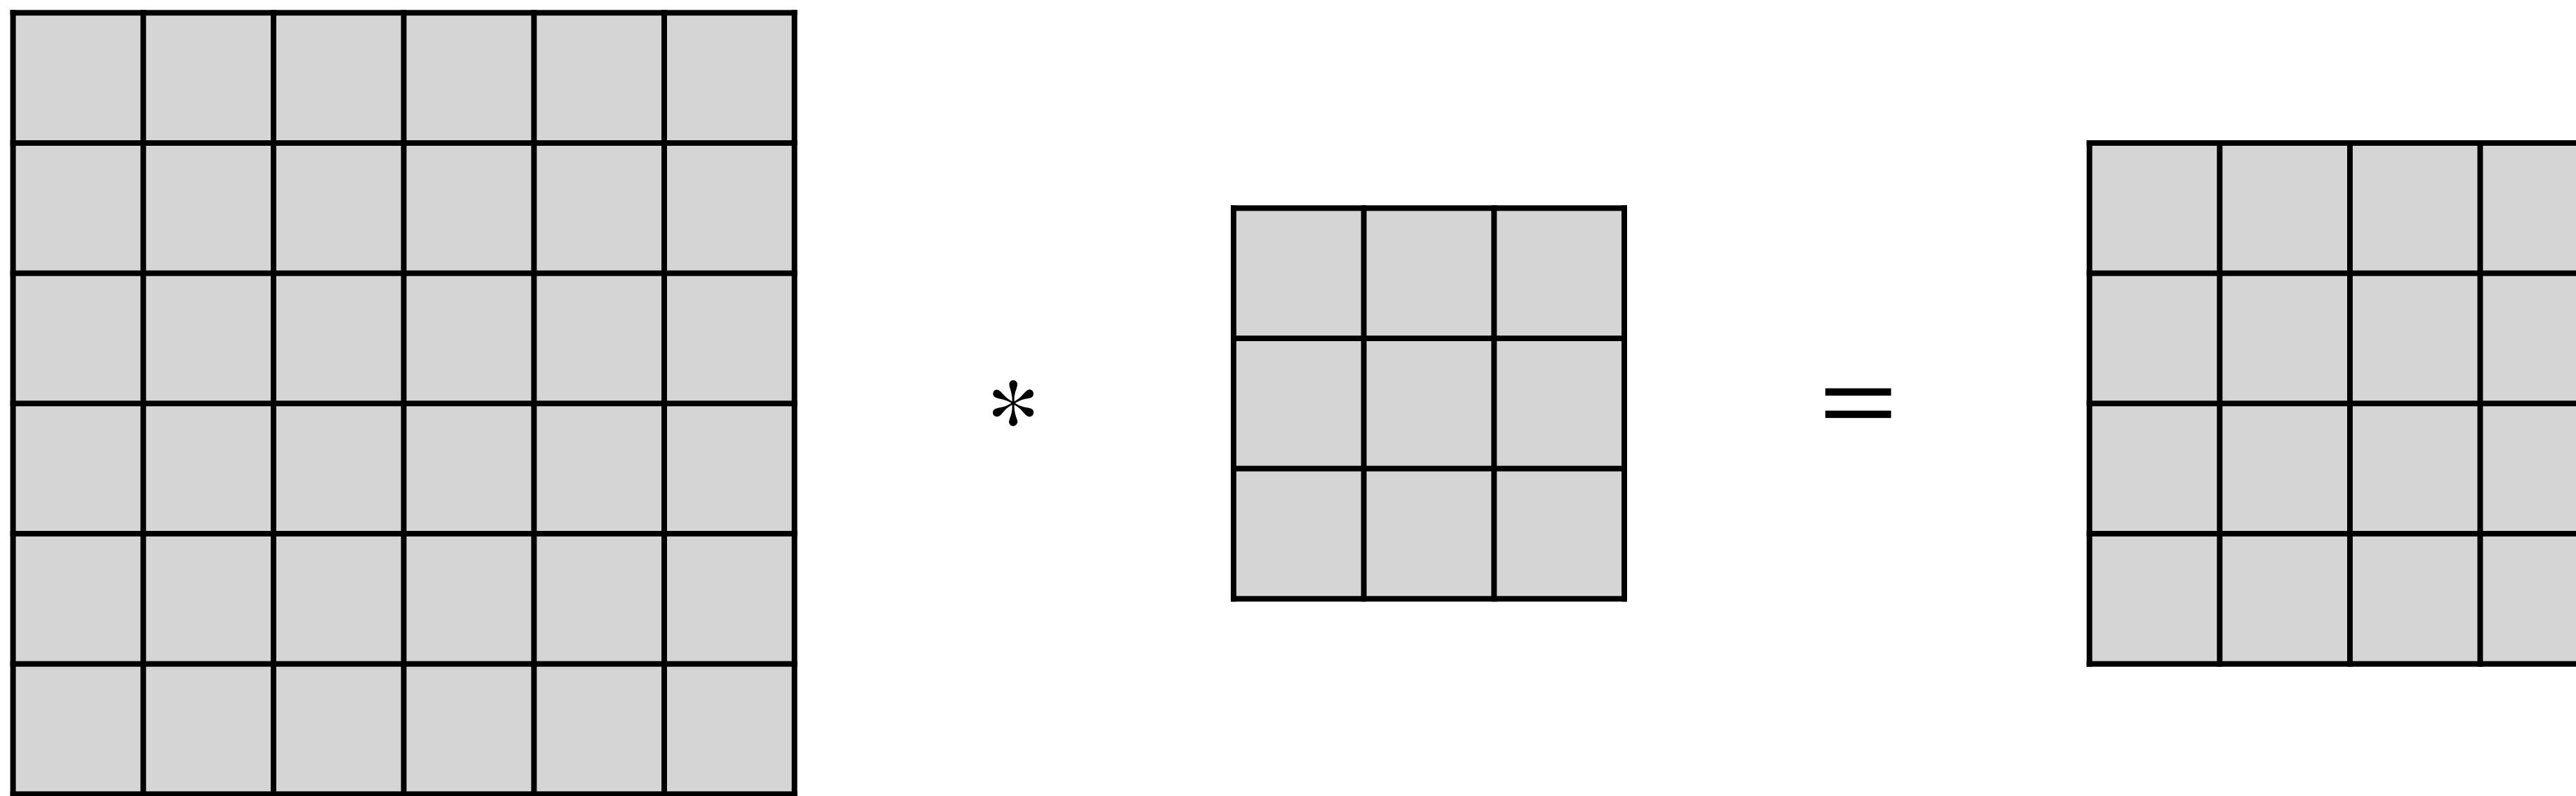
- one kernel may not be enough to capture all important features
- multiple **kernels with independent weights**
- new output channel for each kernel

# Channels



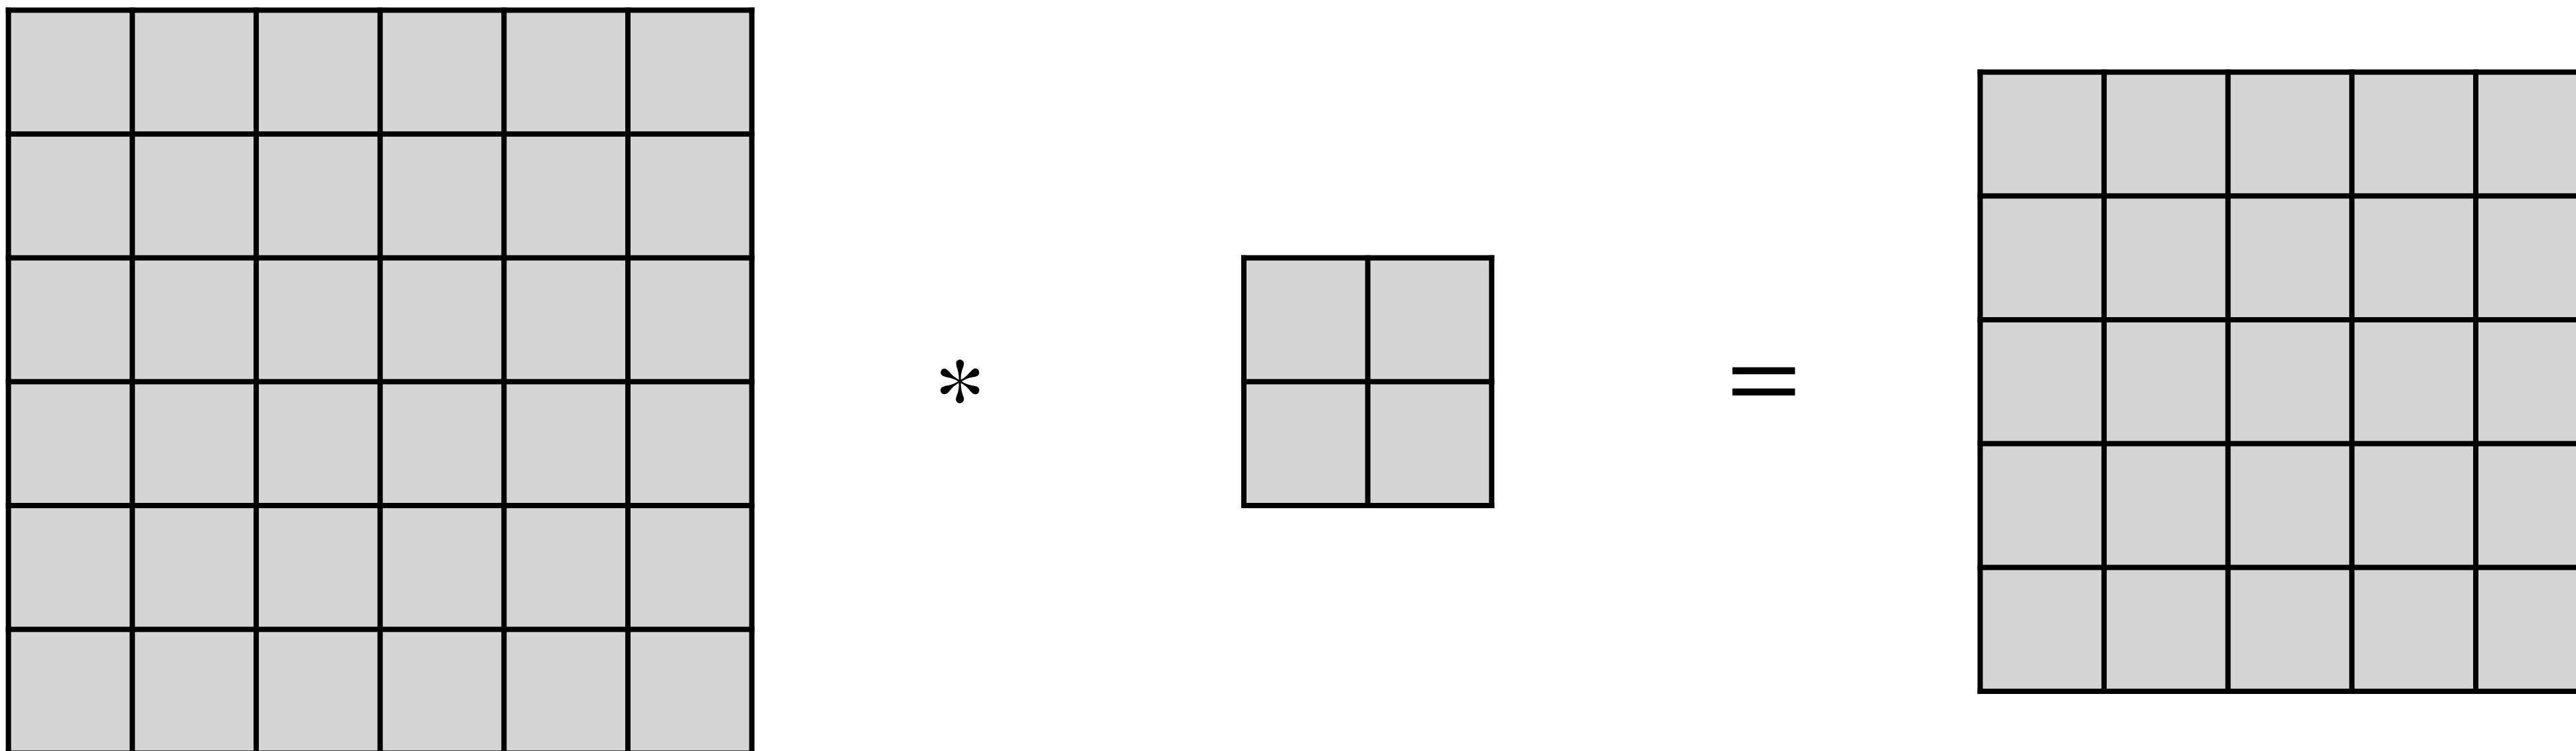
- Combine both:  $O_i = I_j * K_{ij}$  with  $N_{\text{in}} \times N_{\text{out}}$  independent kernels
- Sum over input channels  $j$

# Kernel size



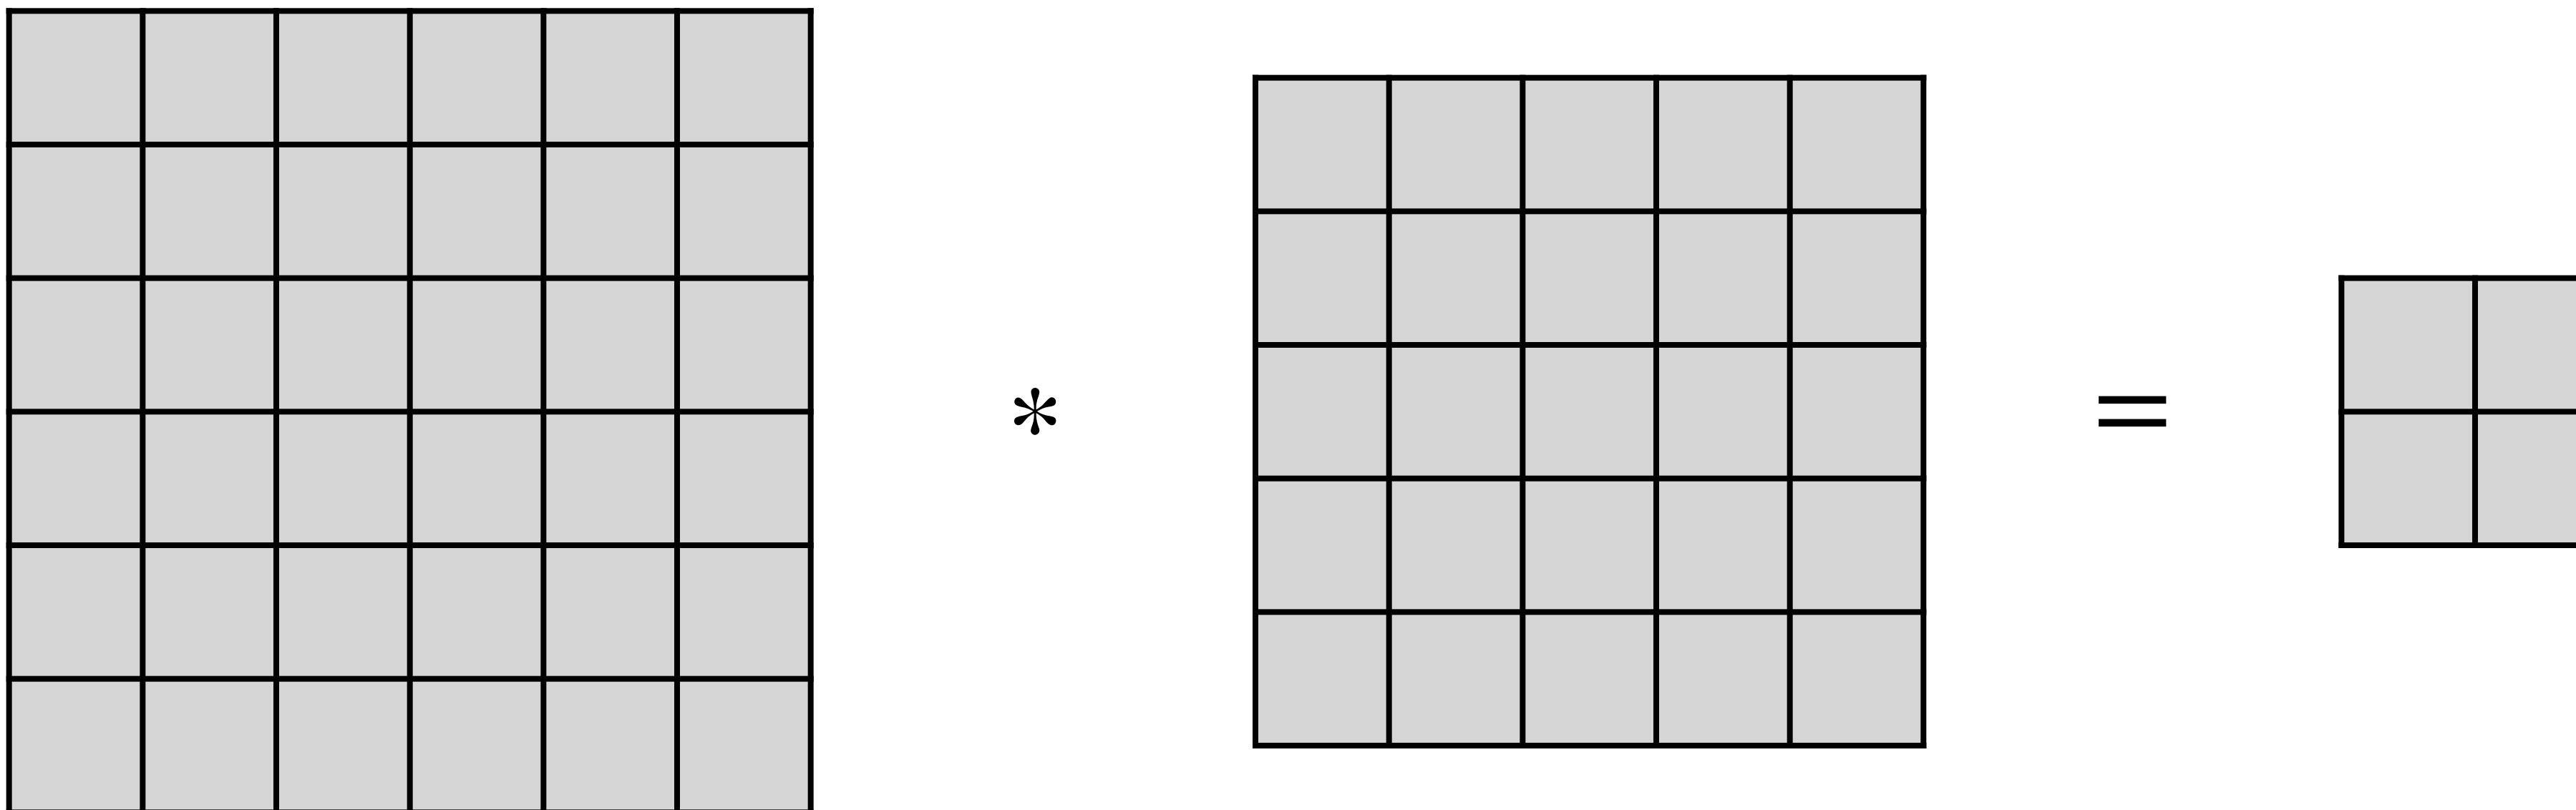
- kernel size **determines how far pixels are connected**

# Kernel size



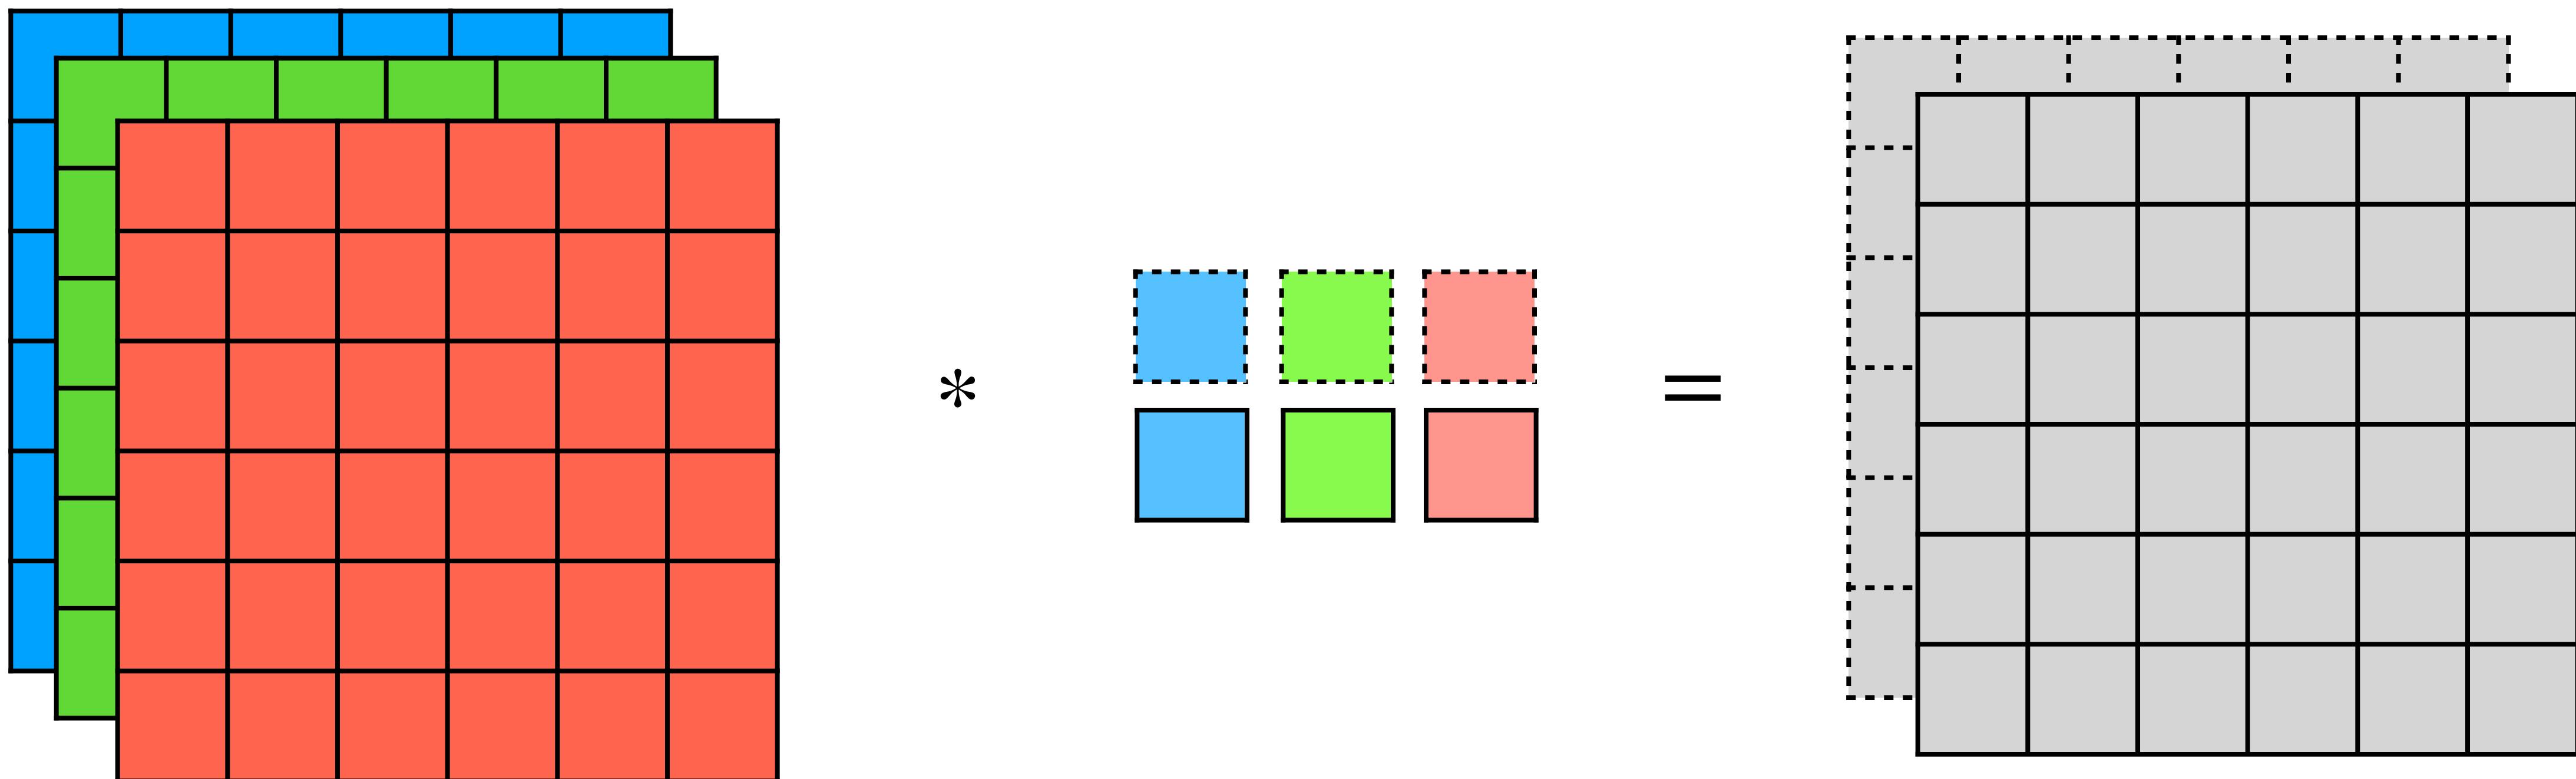
- kernel size **determines how far pixels are connected**
- too small: can't cover interesting features

# Kernel size



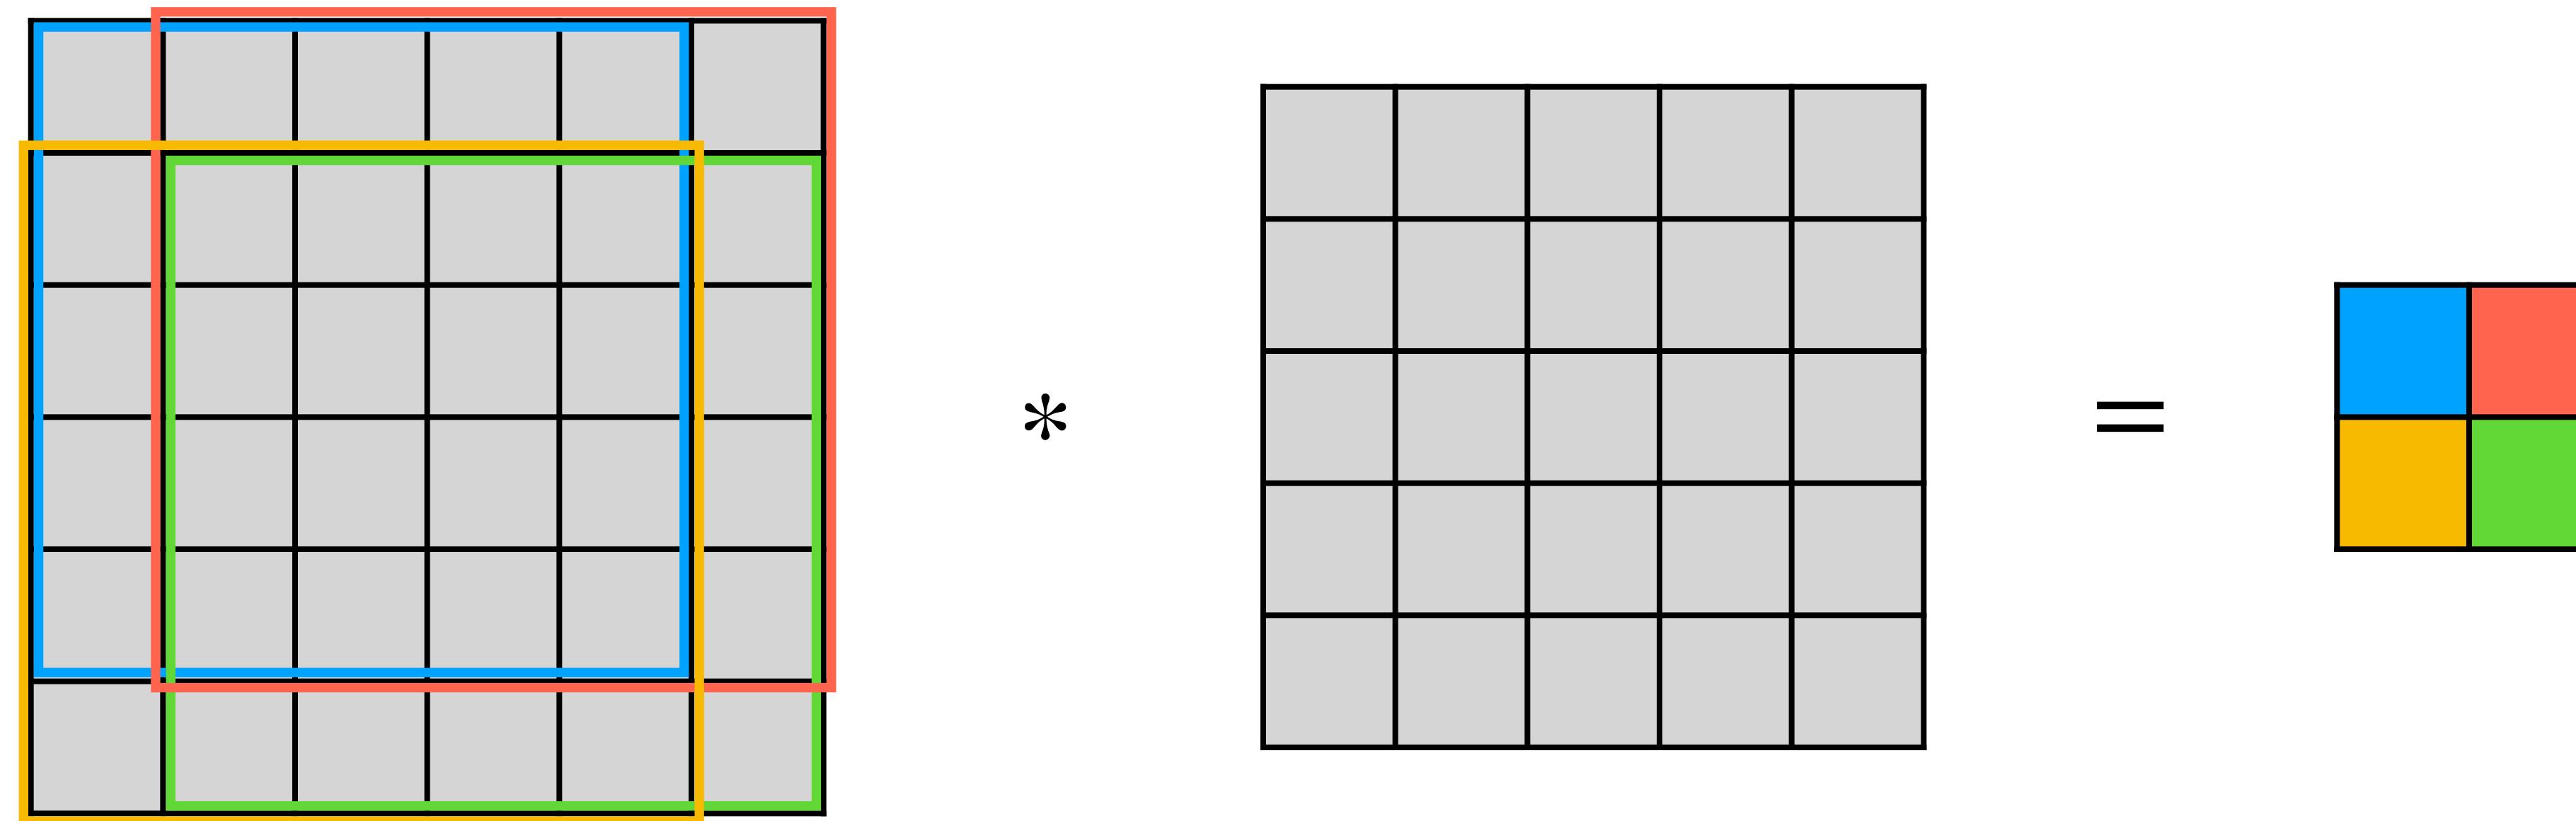
- kernel size **determines how far pixels are connected**
- too small: can't cover interesting features
- too large: don't profit from weight sharing any more

# 1x1 Convolutions



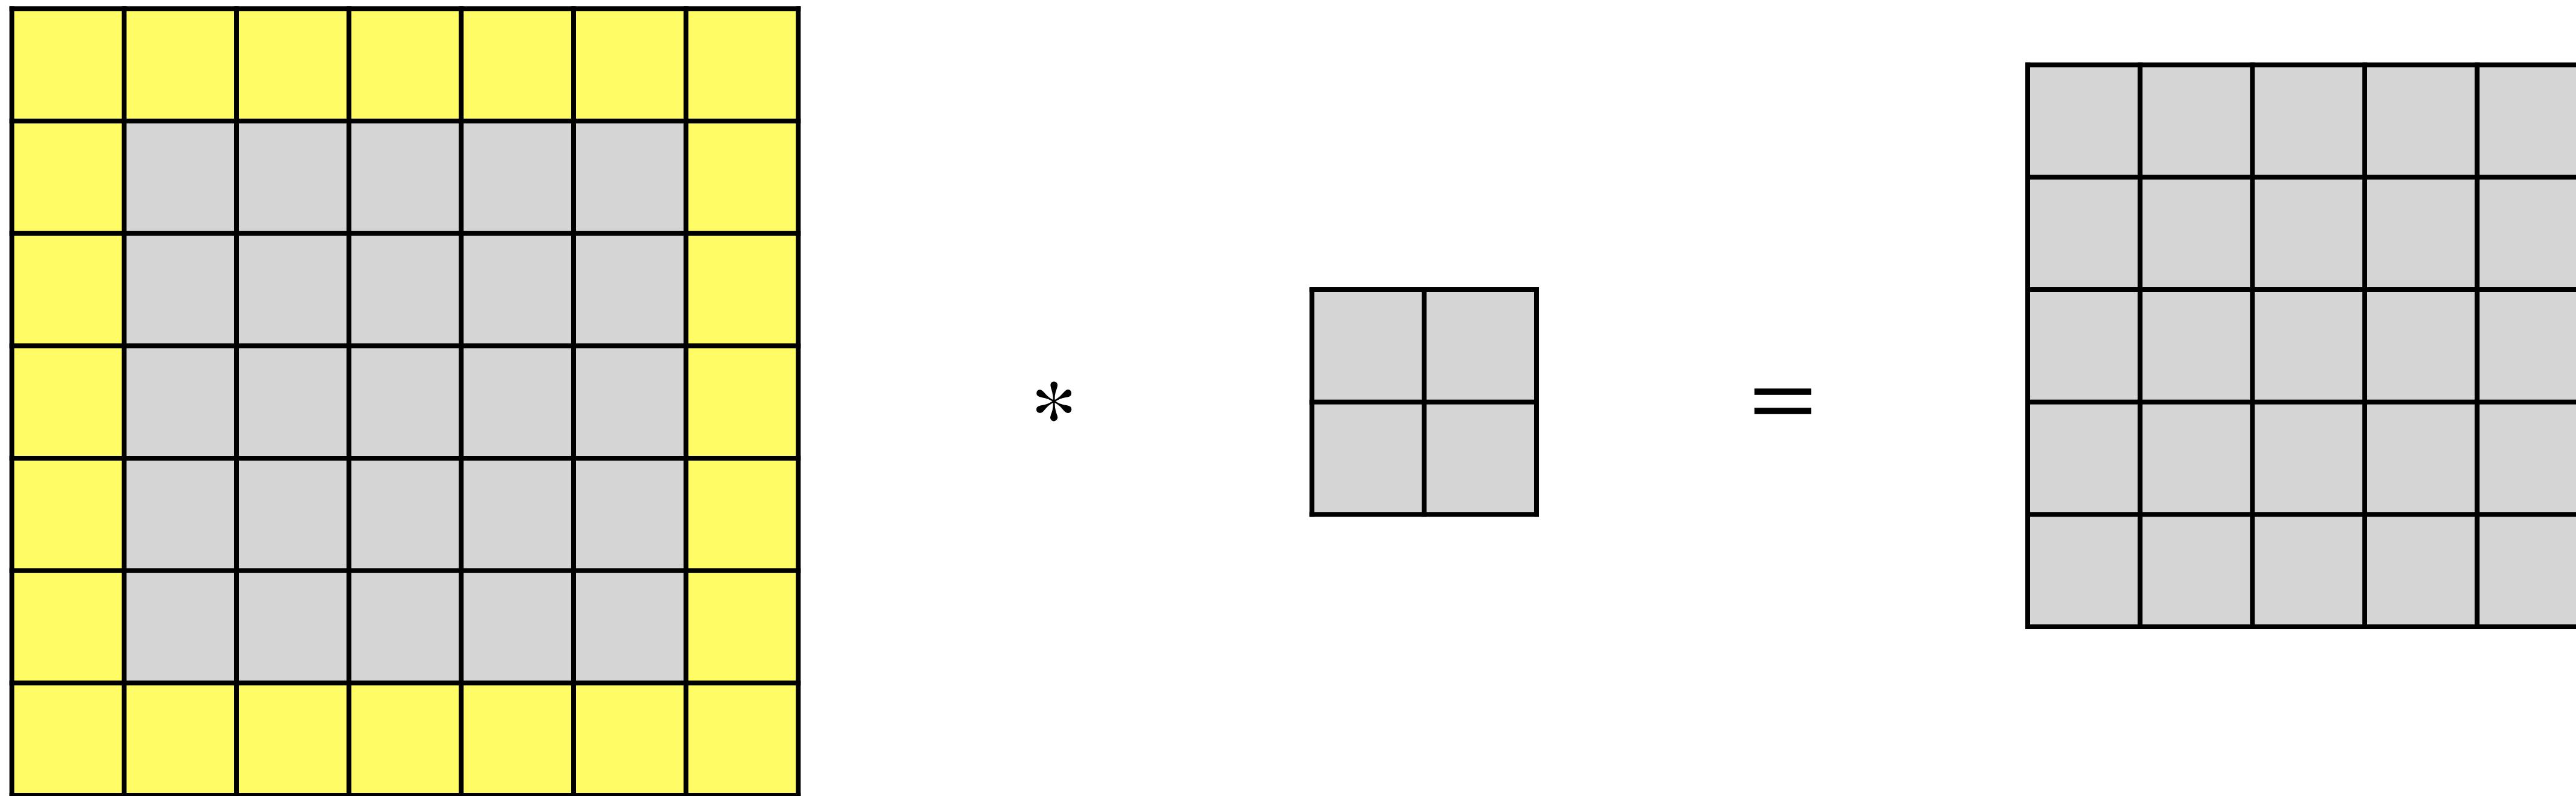
- Interesting special case: 1x1 convolution
- Acts like **linear layer applied to every pixel** individually

# Output size



- kernel size affects output size
- image and kernel need to fully overlap  
→ output always smaller than input
- output size = input size - kernel size + 1

# Padding

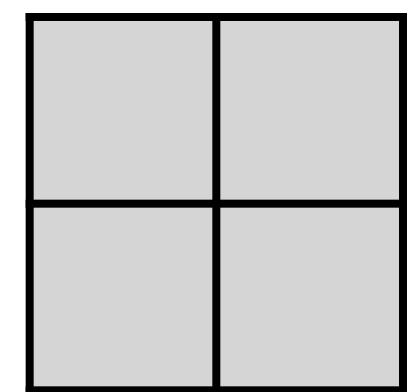


- use padding to **prevent reduction in output size**
- PyTorch: padding=“same” such that output size = input size
- Also possible: fixed value for padding size

# Padding

0	0	0	0	0	0	0
0						0
0						0
0						0
0						0
0						0
0	0	0	0	0	0	0

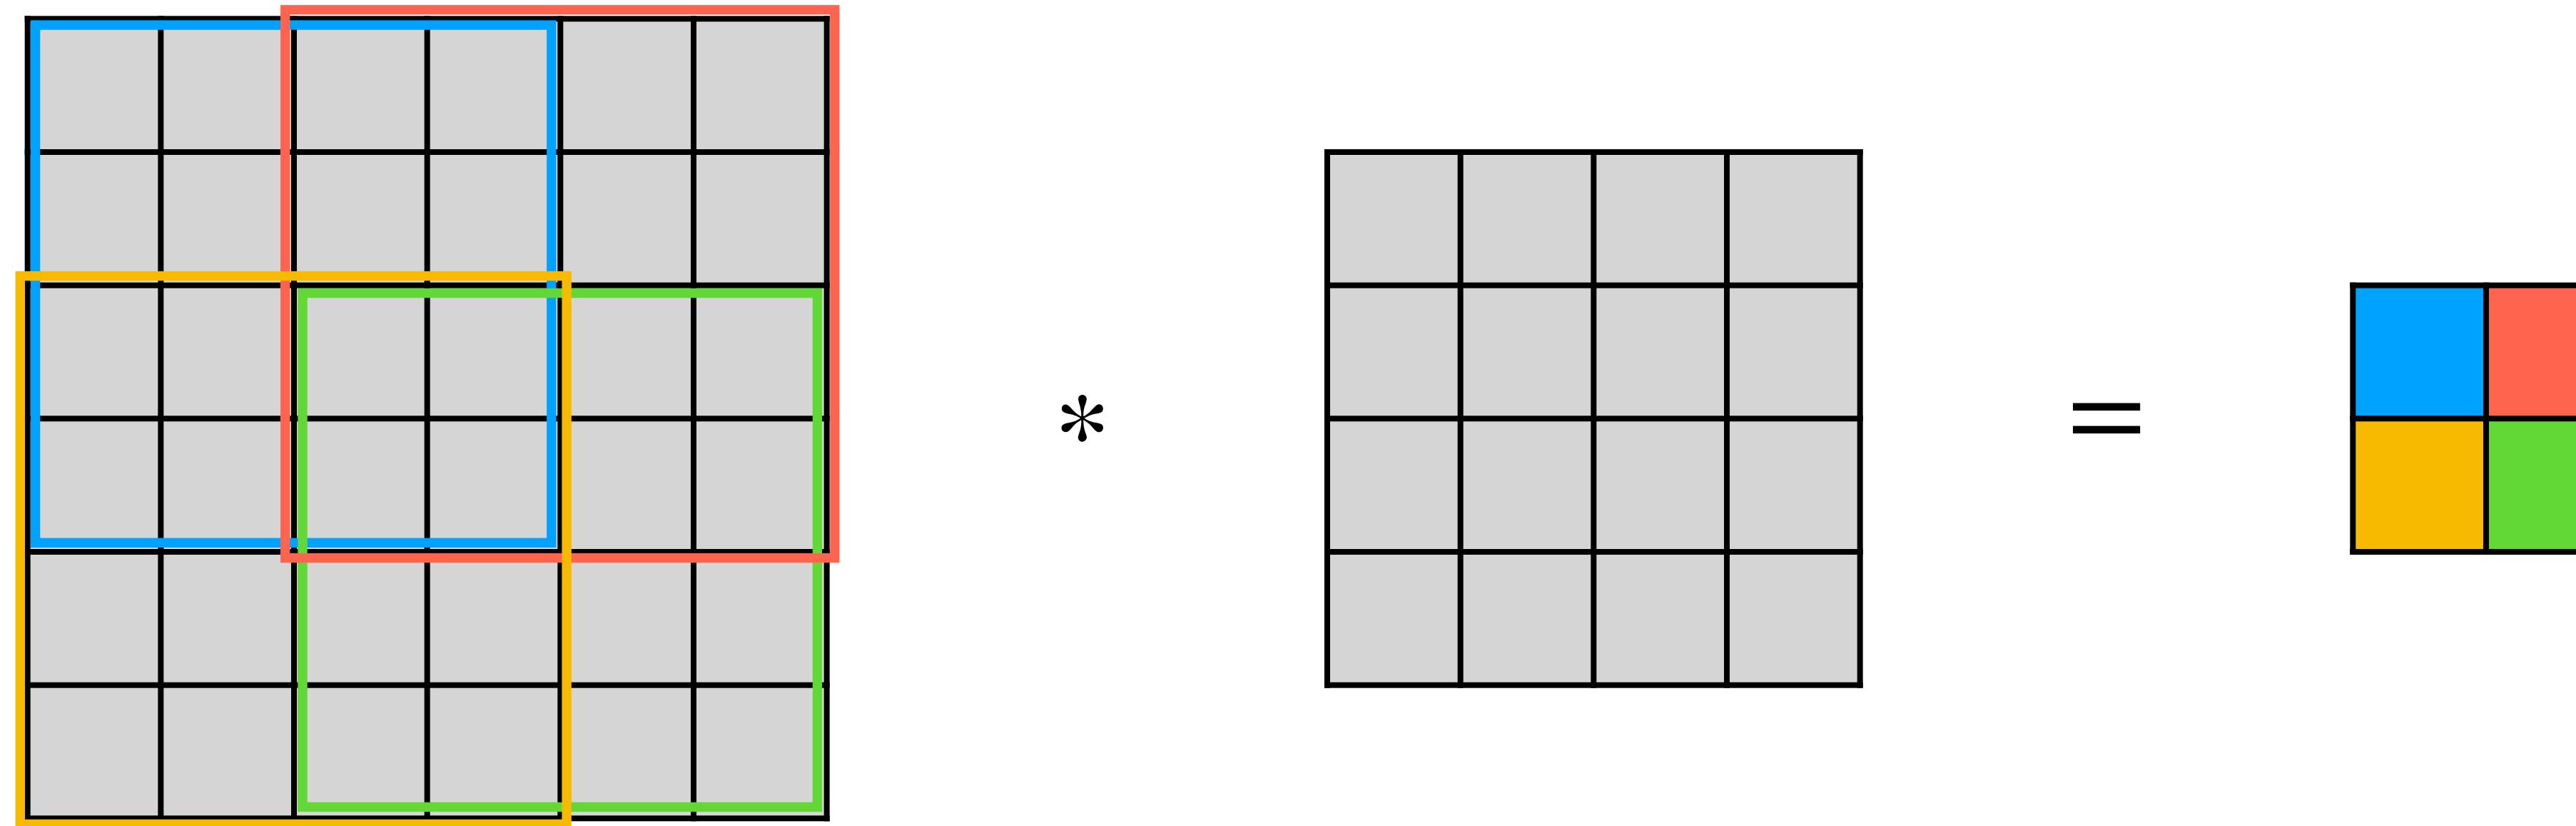
\*



=


- use padding to **prevent reduction in output size**
- PyTorch: padding=“same” such that output size = input size
- Also possible: fixed value for padding size
- Default mode: use zeros for padding

# Strides



- strides: amount that kernel moves per step
  - can be greater than one
- **combined convolution and downsampling**

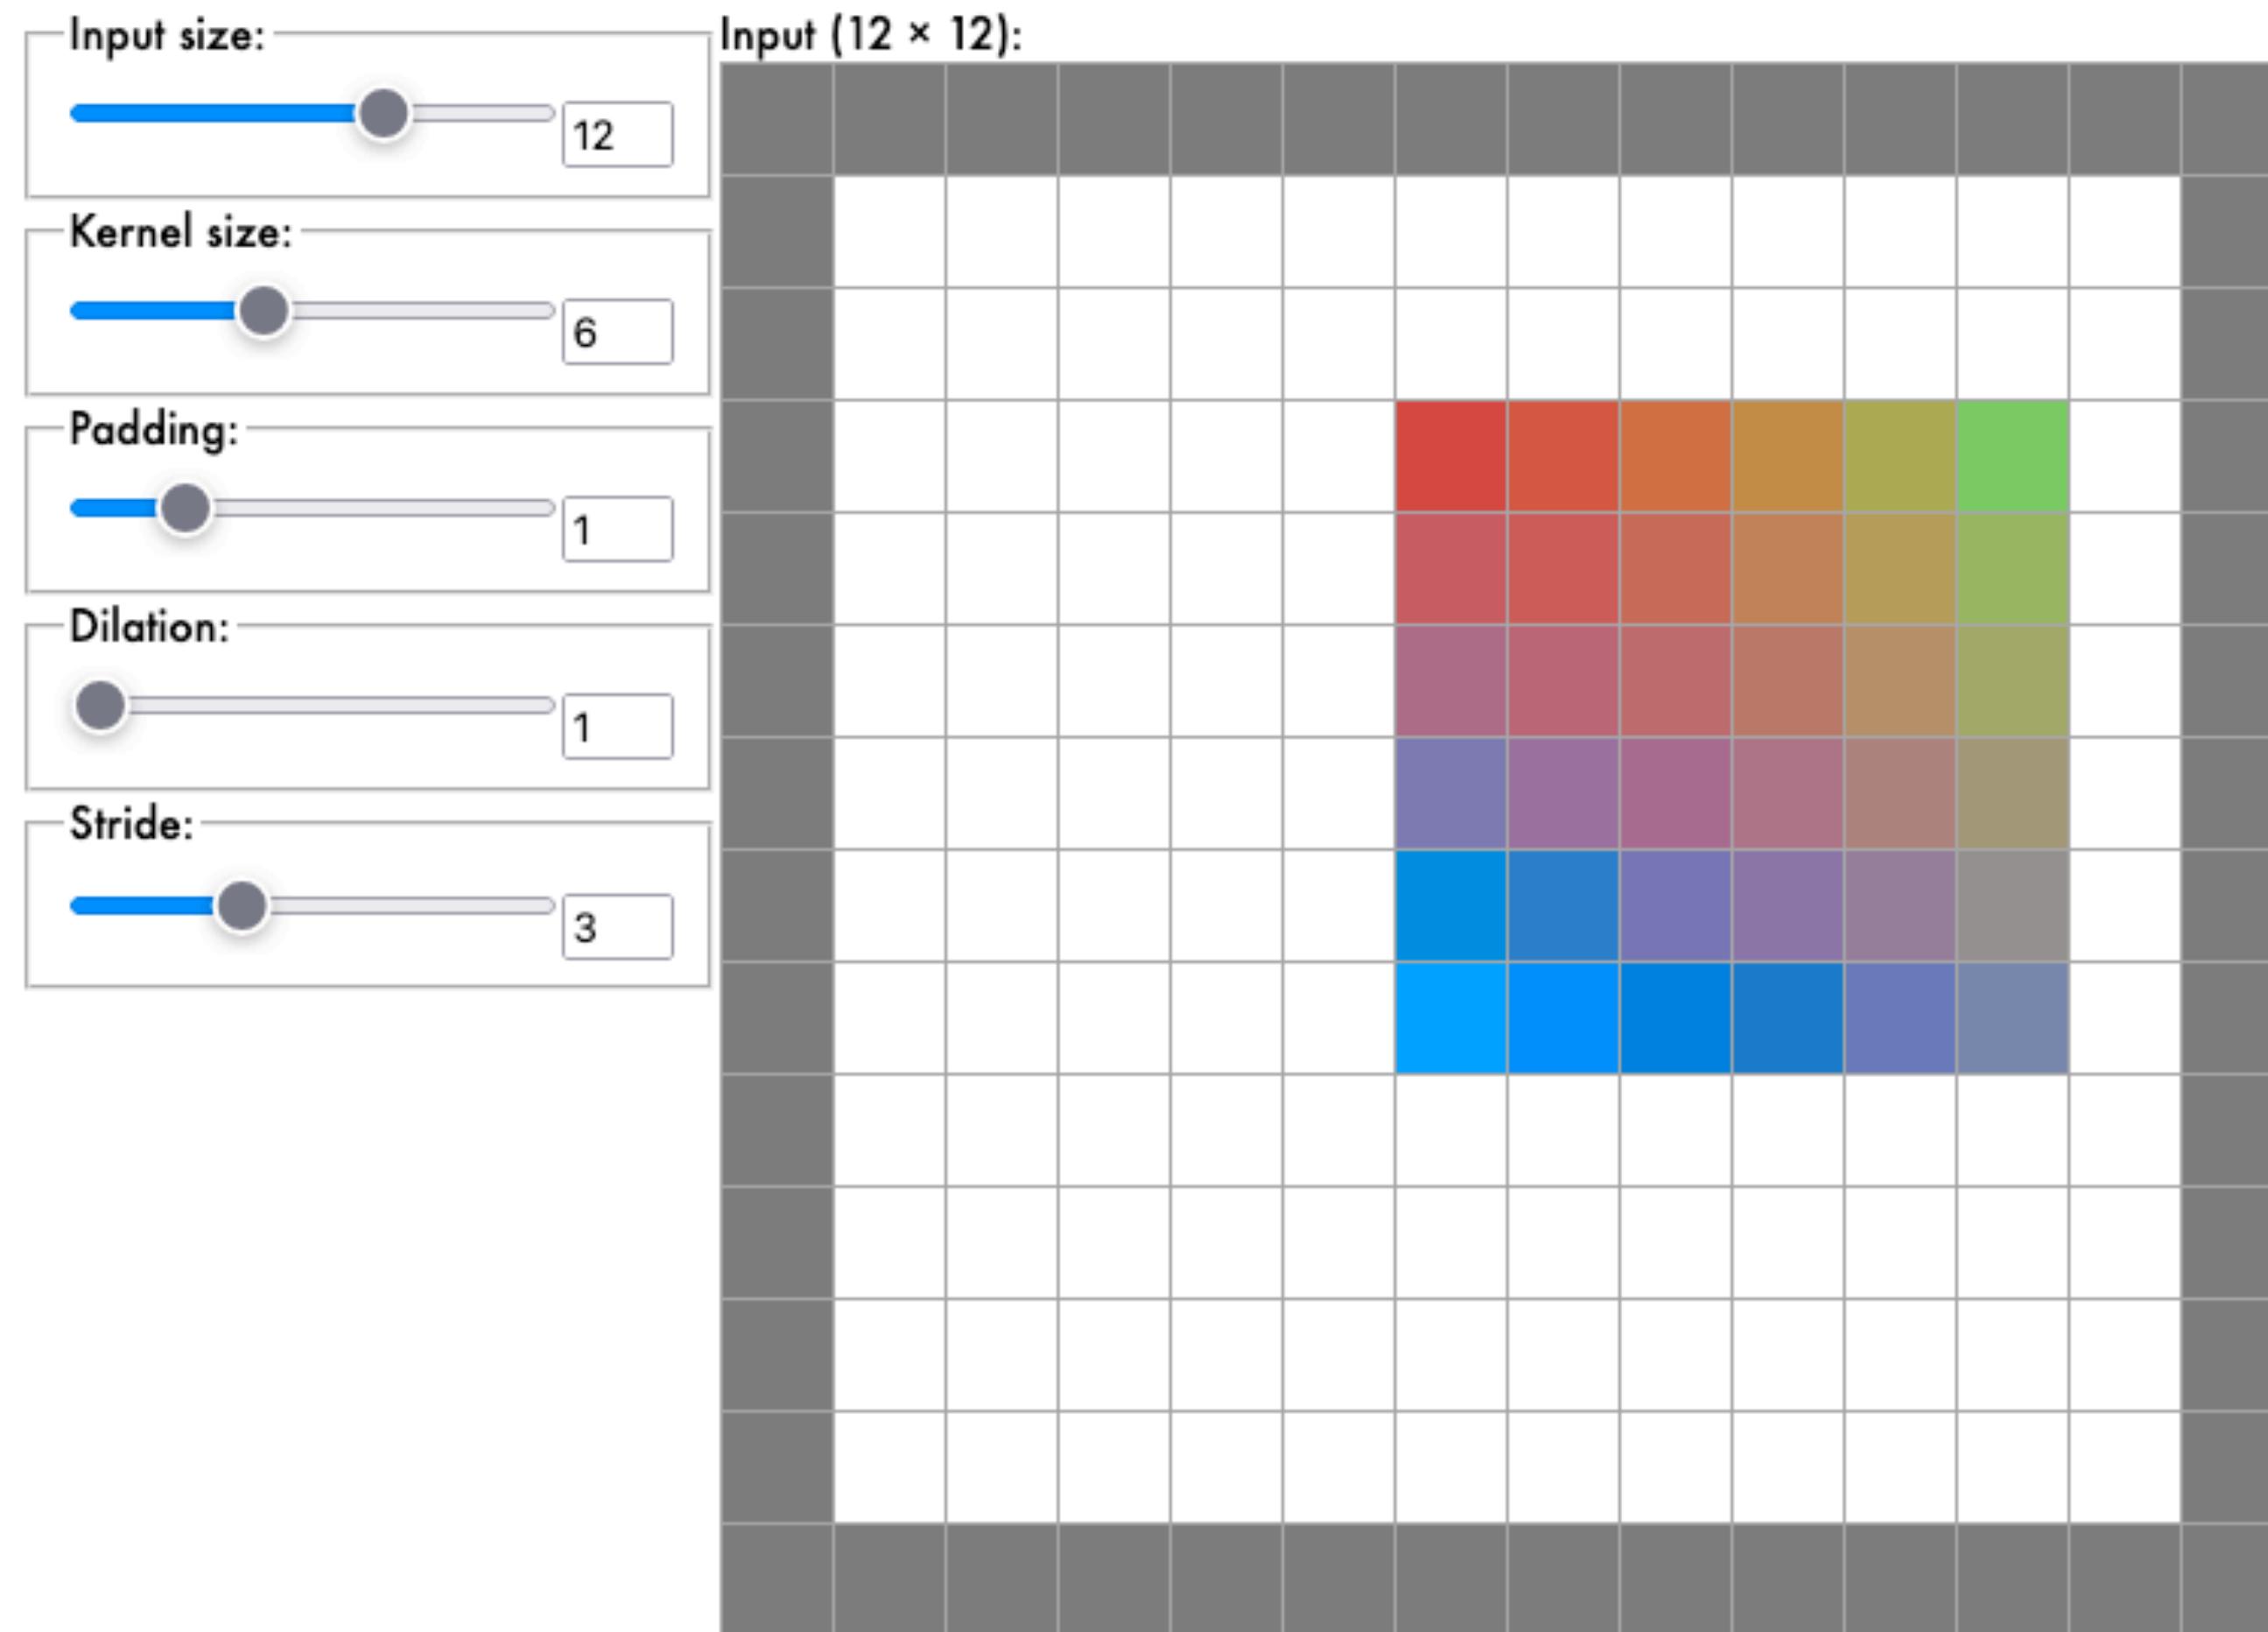
# CNN layers

```
torch.nn.Conv2d(  
    in_channels,  
    out_channels,  
    kernel_size,  
    stride=1,  
    padding=0,  
    dilation=1,  
    groups=1,  
    bias=True,  
    padding_mode='zeros',  
    device=None,  
    dtype=None  
)
```

## Most important settings

- ① **Channels**  
number of input/output feature maps
- ② **Kernel size**  
choose according to size of features
- ③ **Padding**  
add pixels to change output size
- ④ **Strides**  
step size, allows for downsampling

# Interactive visualization

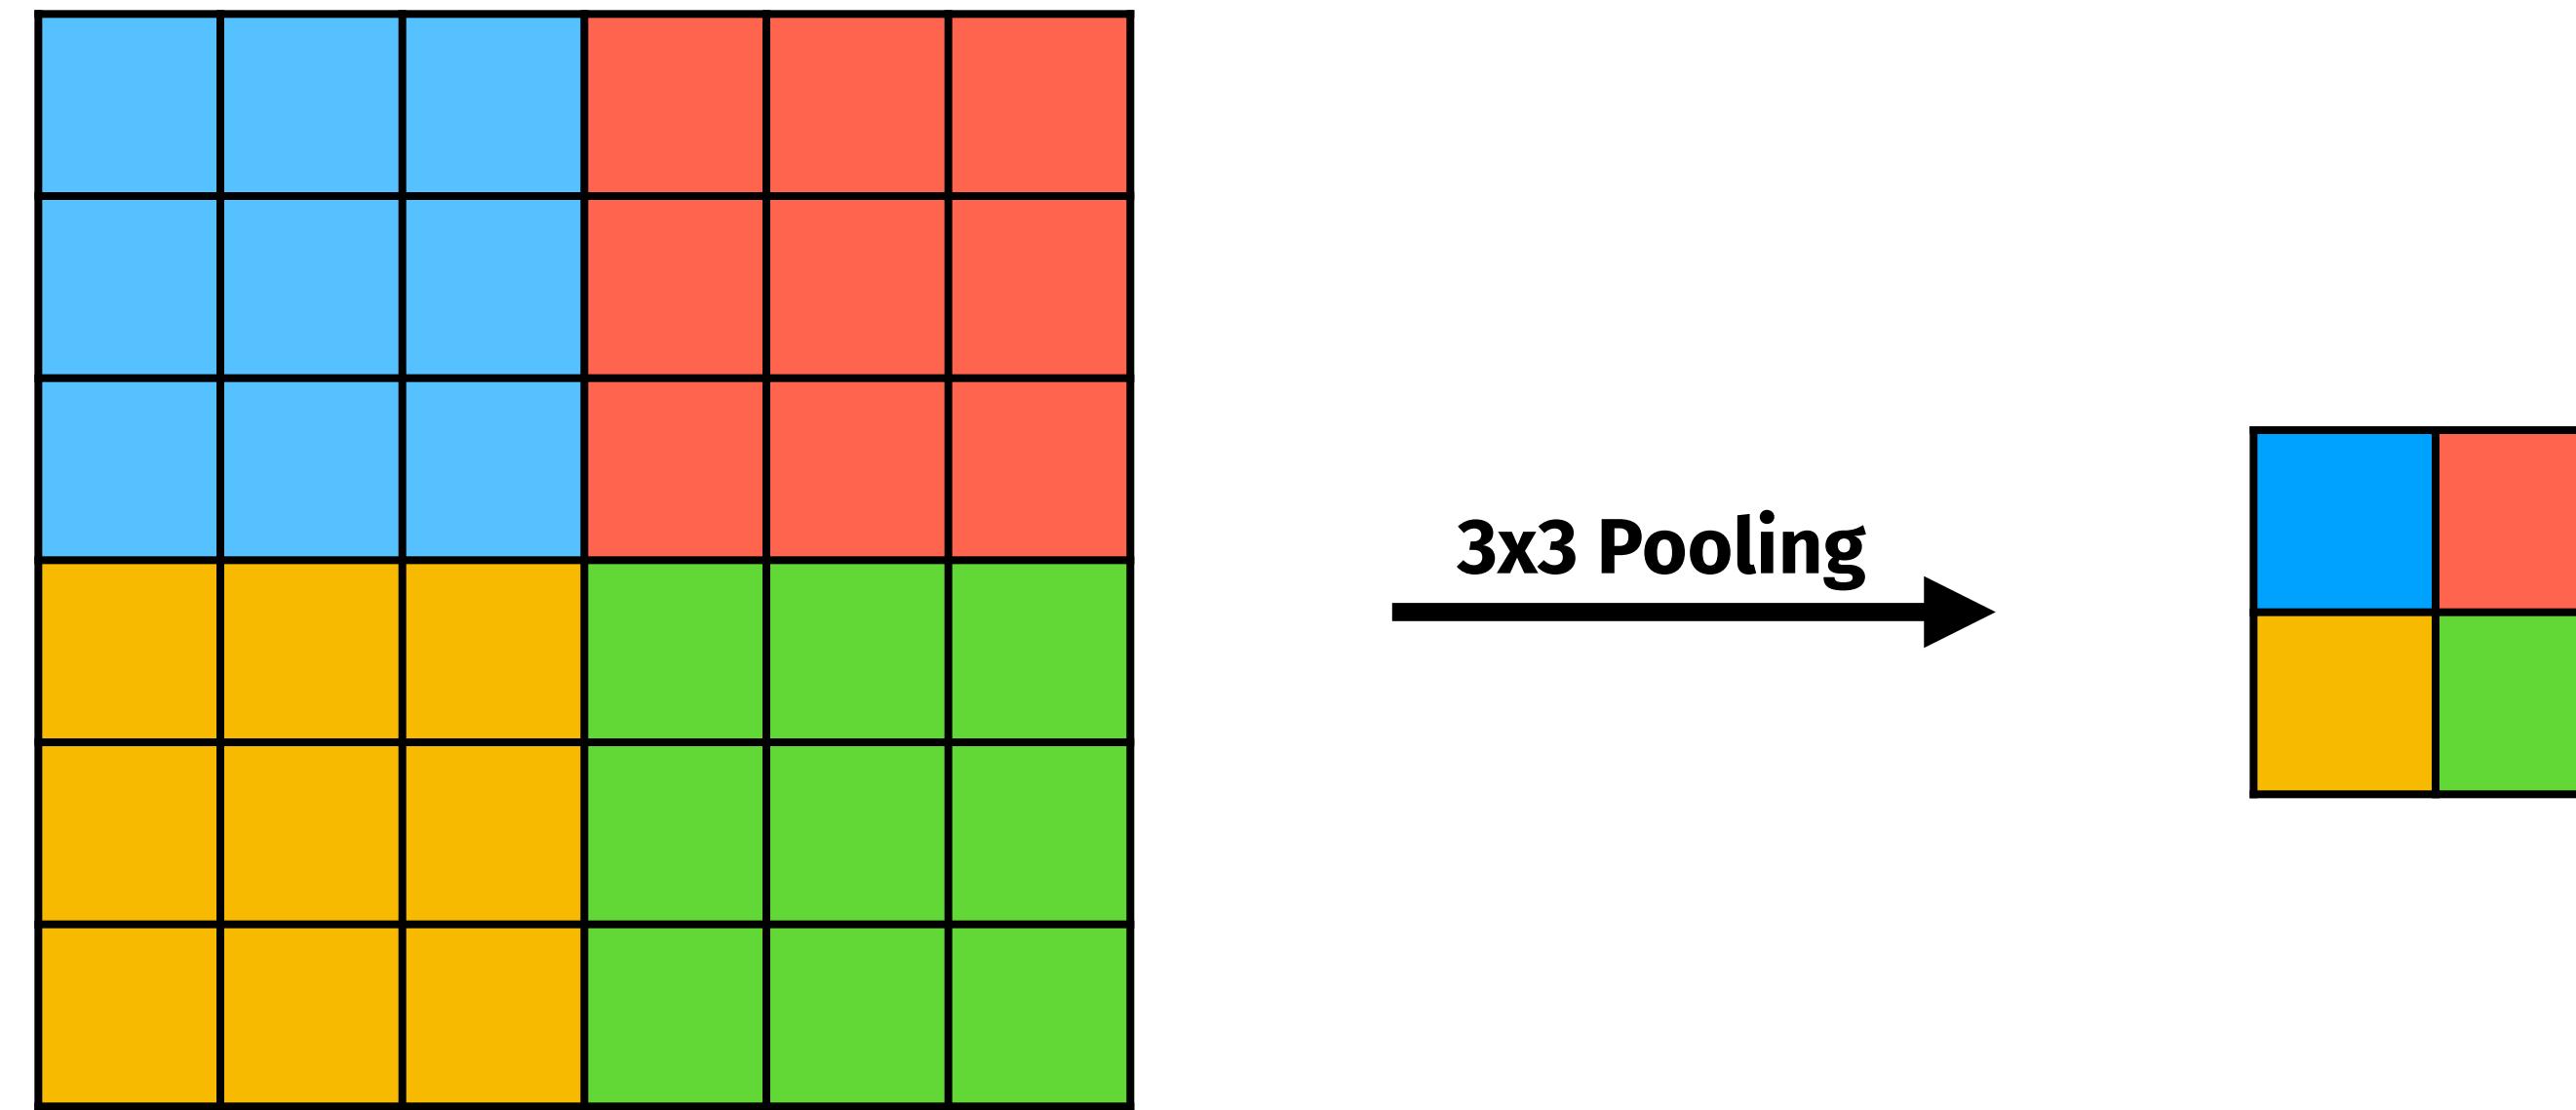


<https://ezyang.github.io/convolution-visualizer/>

# Outline

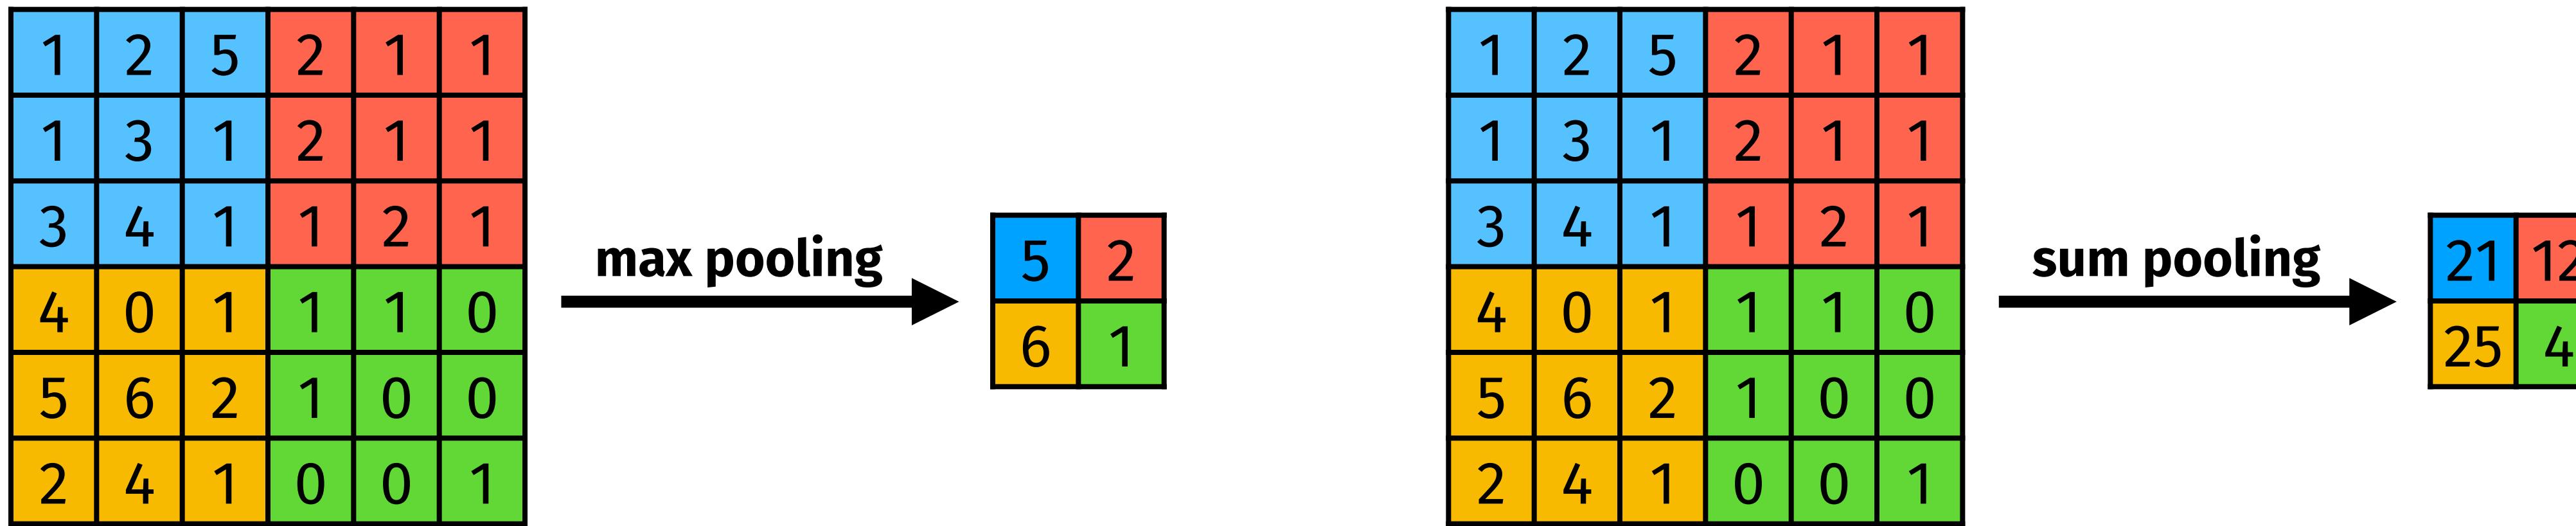
- ① Challenges of image processing
- ② Introduction to convolutions
- ③ Convolutional layers
- ④ **CNN building blocks**
- ⑤ A short history of CNNs
- ⑥ Outlook

# Pooling



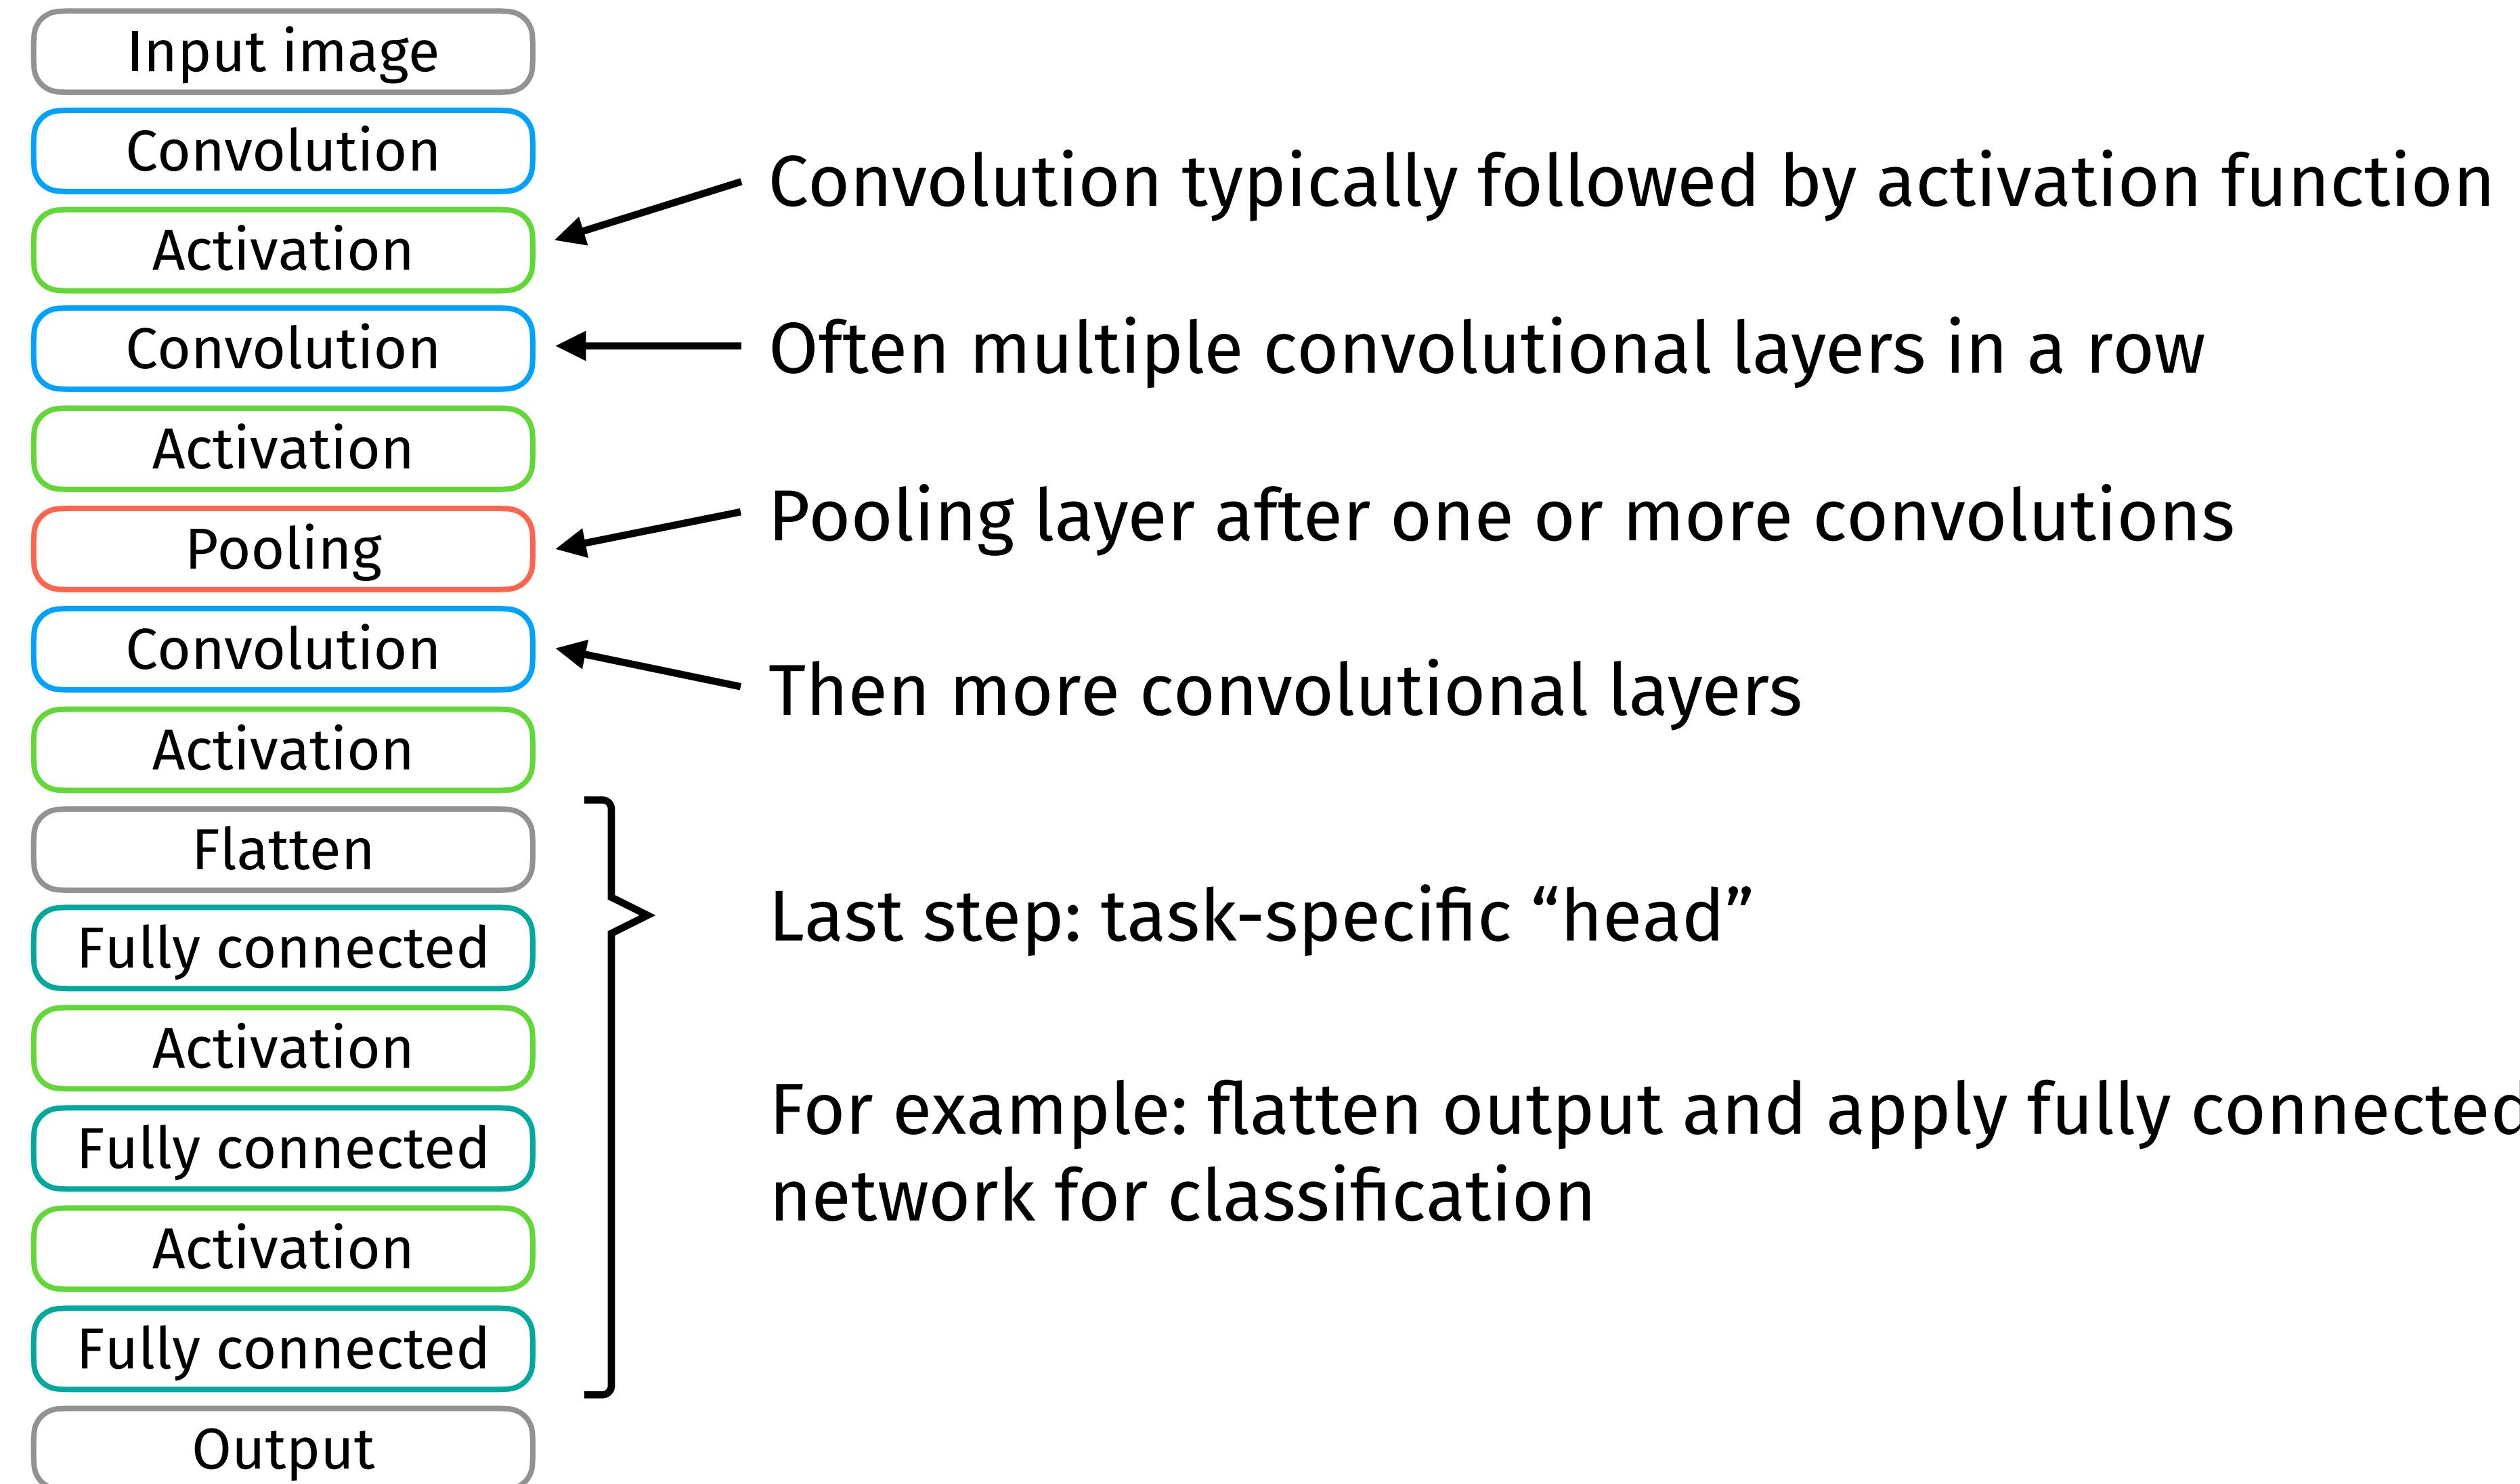
- Often need to **reduce image size**
- Seen before: convolution with strides
- Alternative: combine values from multiple pixels to one
- $\text{output size} = \text{input size} / \text{kernel size}$

# Pooling



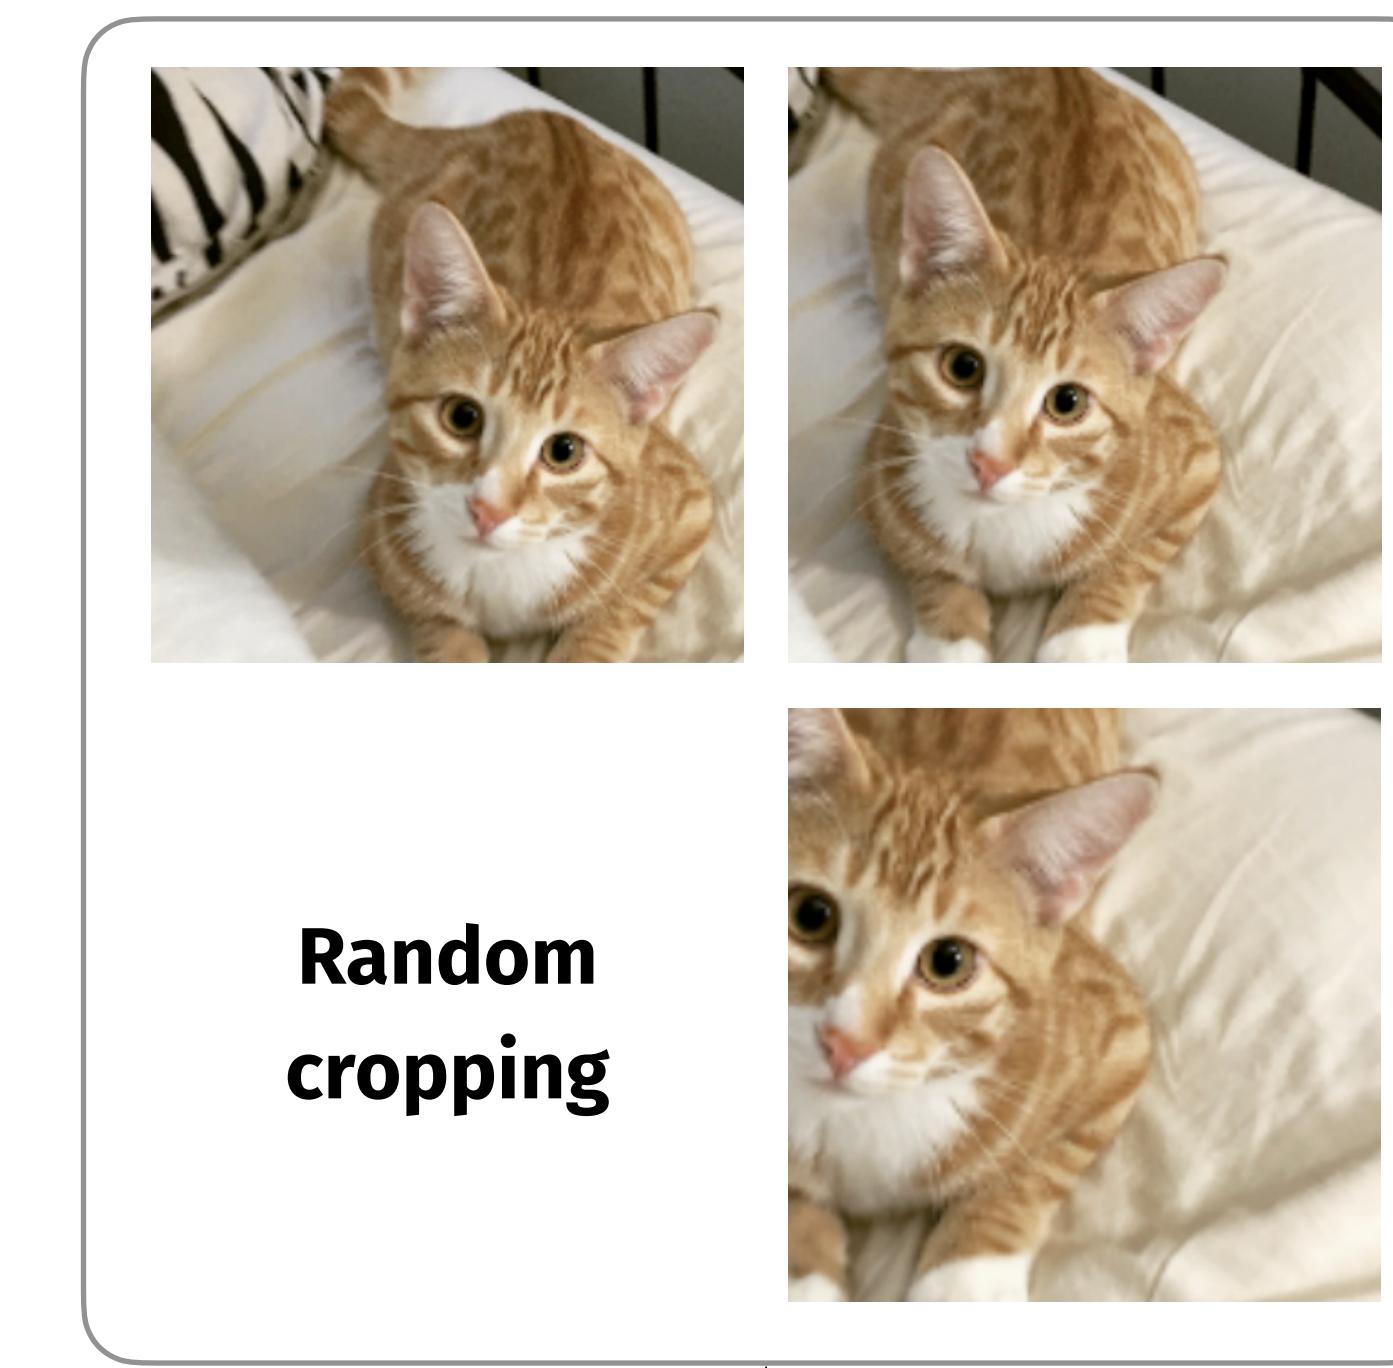
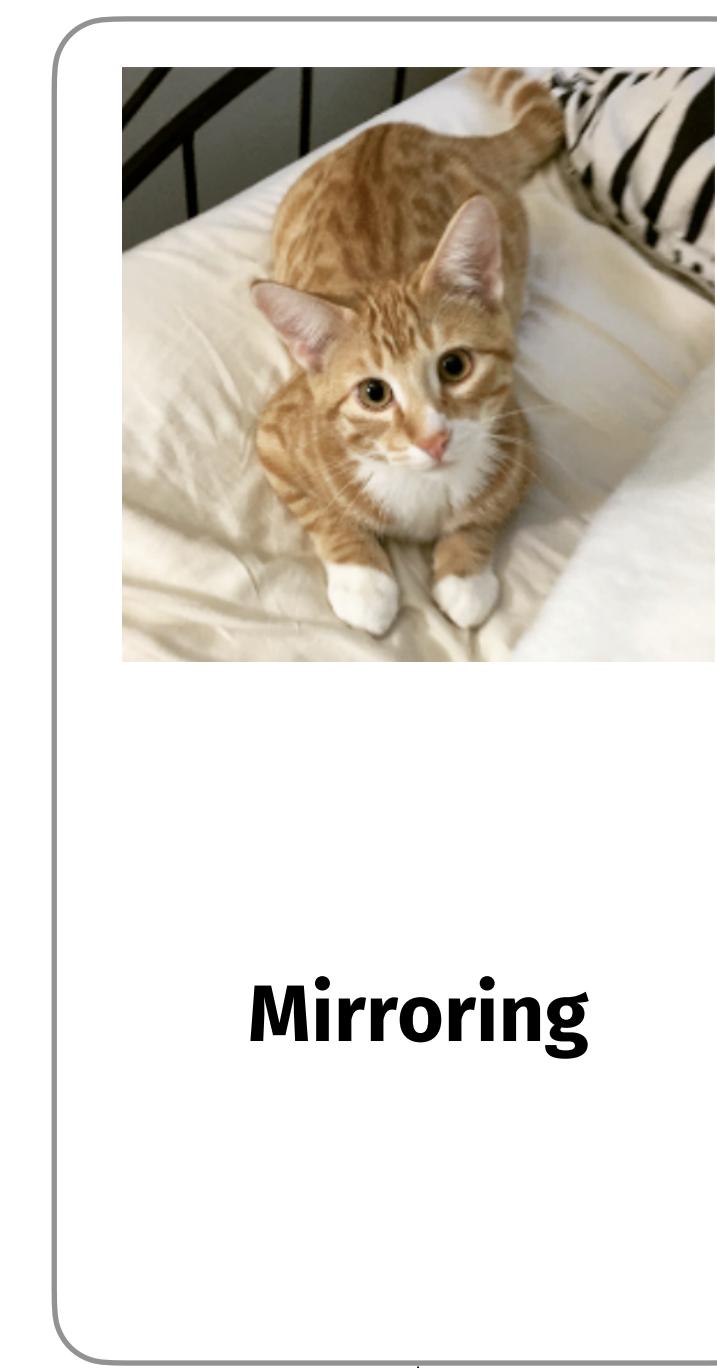
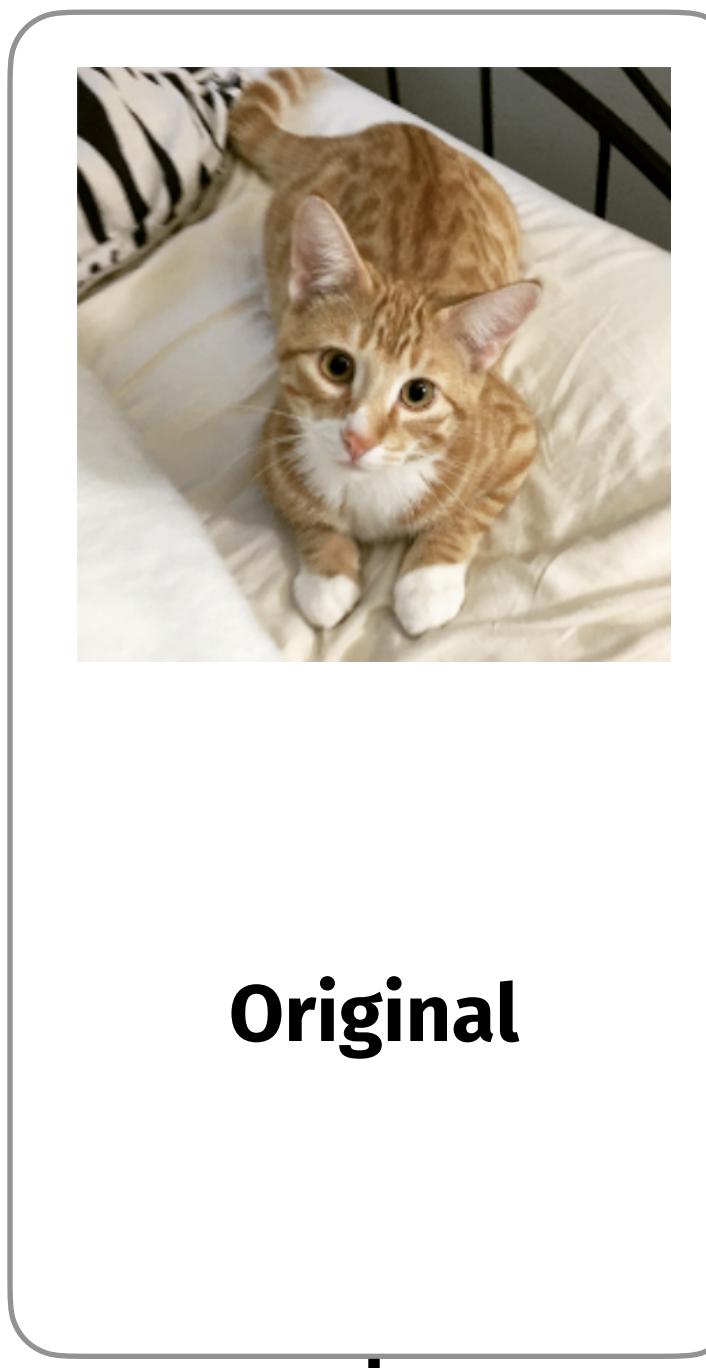
- Different functions can be used for pooling
  - max pooling: focus on high value pixels
  - sum pooling: preserve total intensity
  - average pooling: preserve average intensity
- PyTorch: `torch.nn.MaxPool2d`, `torch.nn.AvgPool2d`

# Typical CNN architecture



# Data augmentation

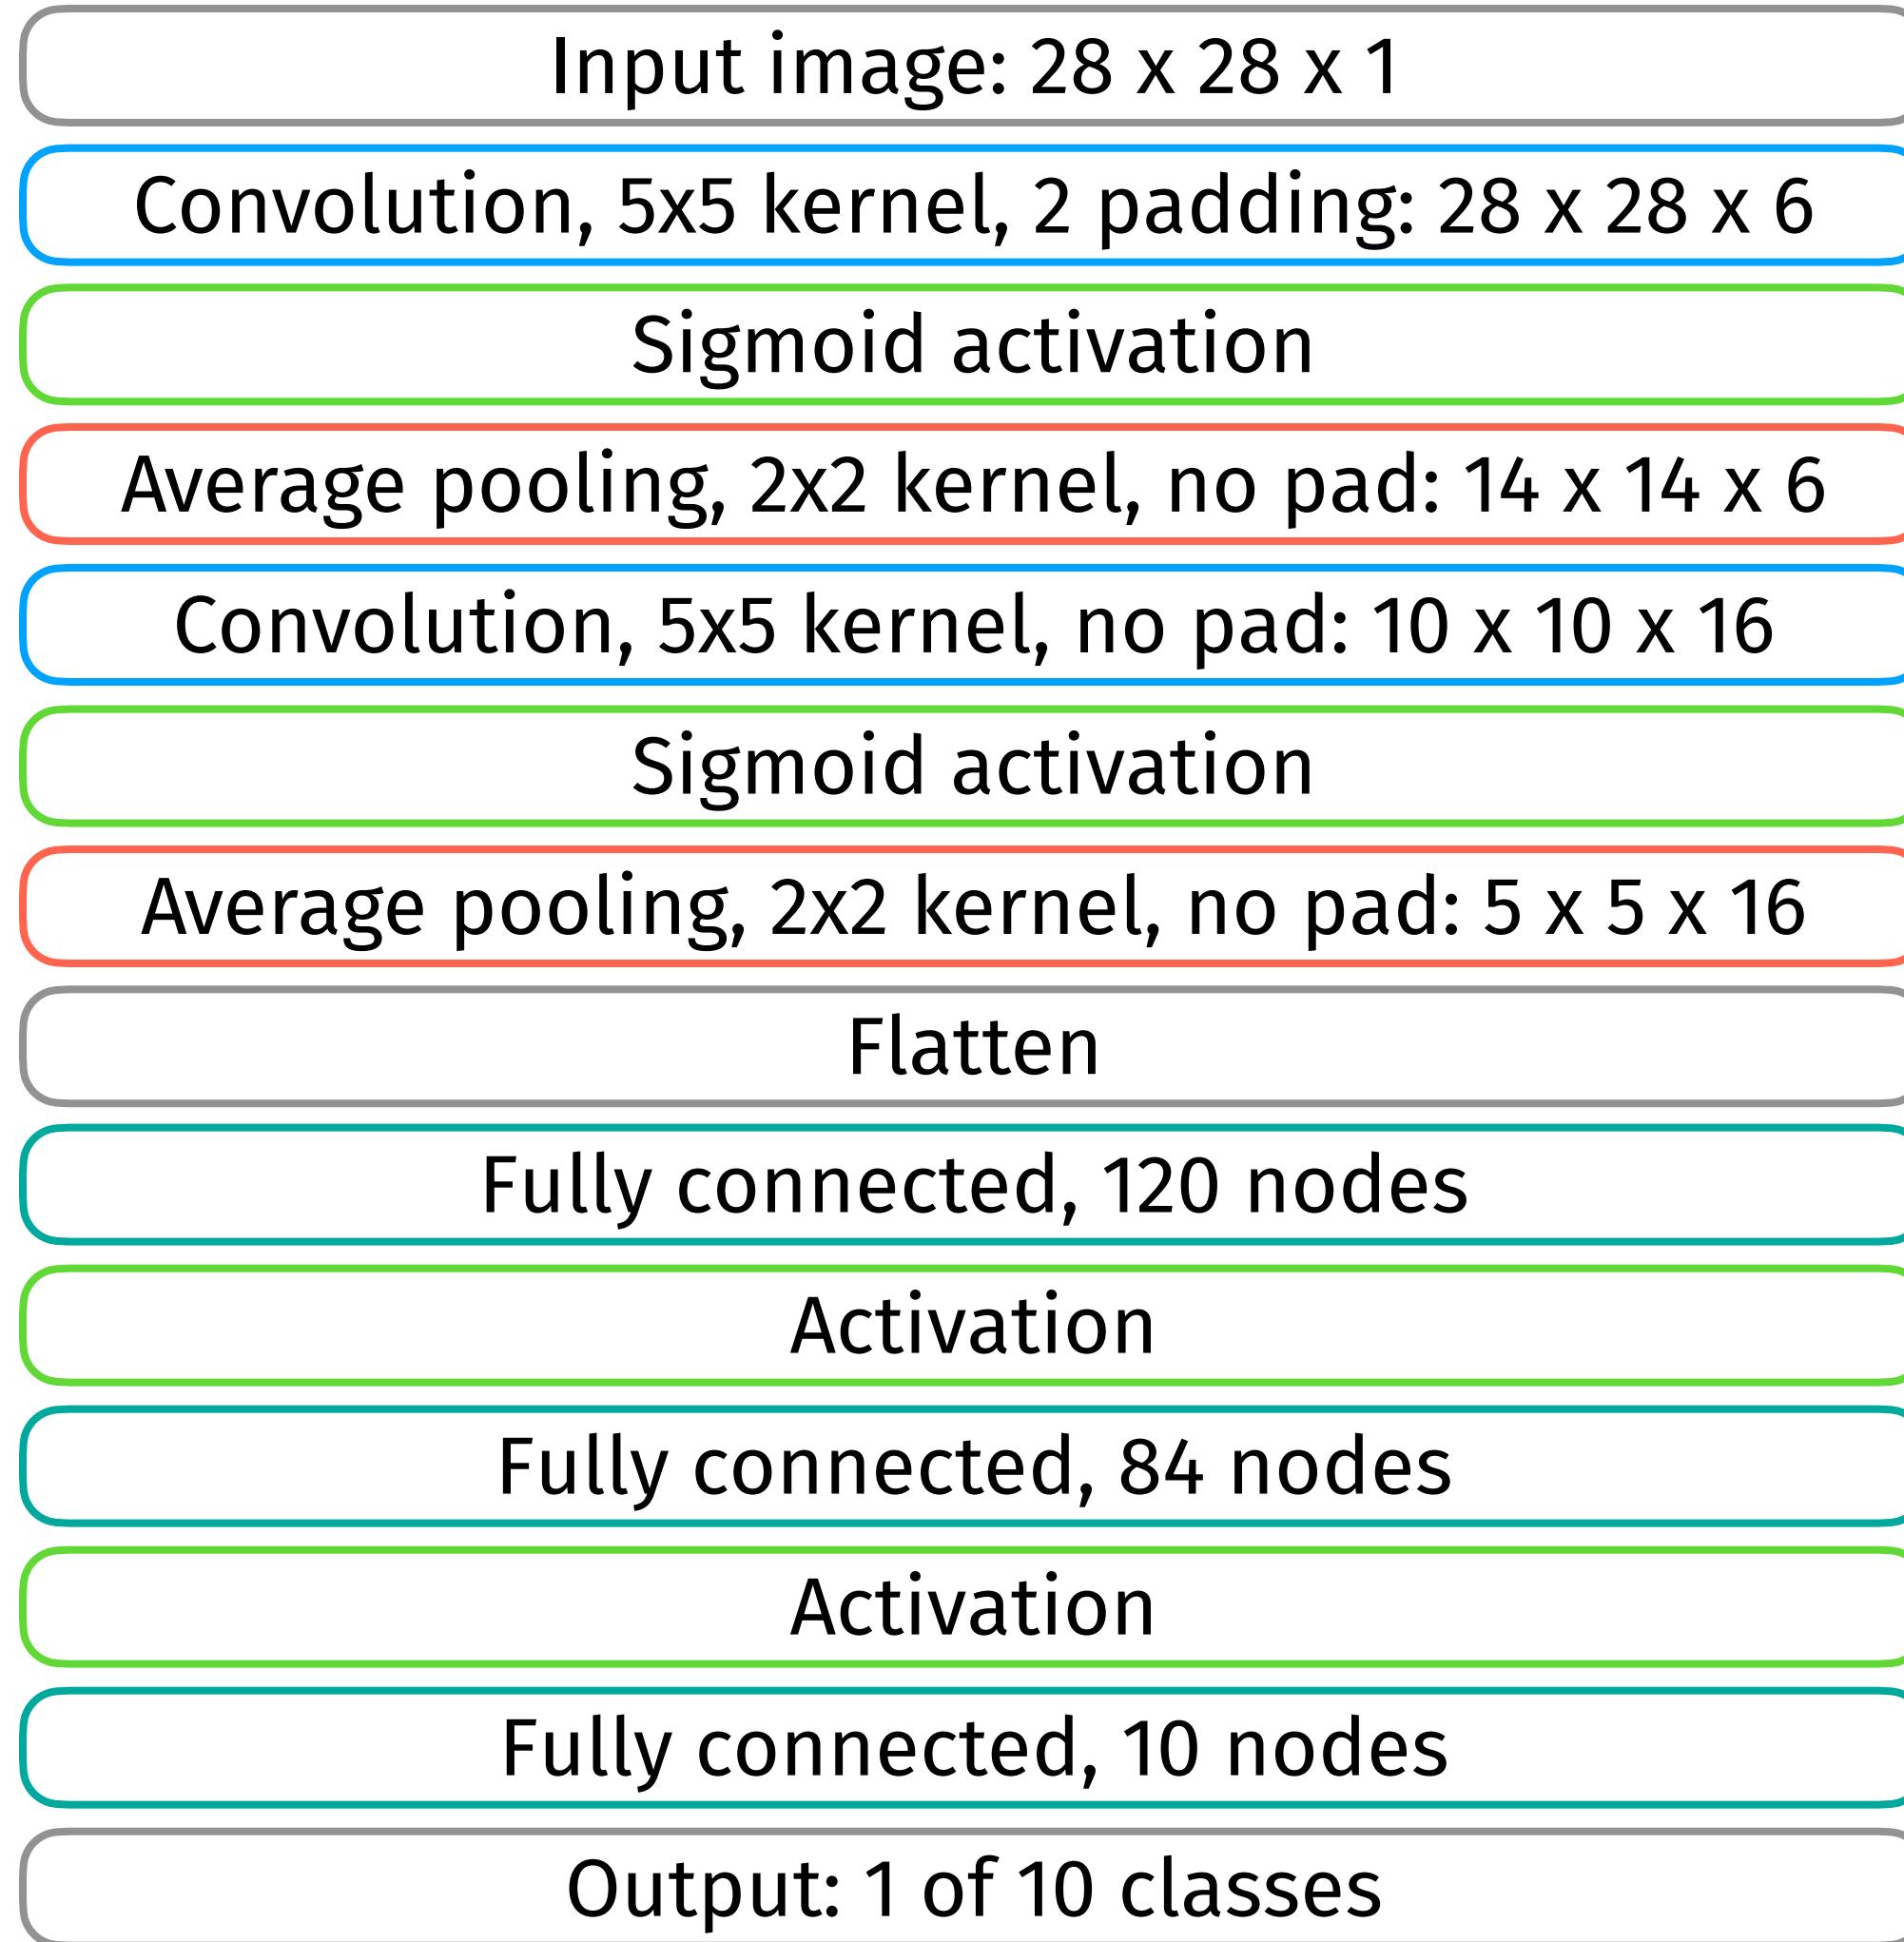
- CNNs: translation-invariant up to boundary effects
- want network output to be invariant under other transformations  
→ cropping, mirroring, rotations, shearing, warping, color shifting
- simple solution: **increase training statistics by augmenting data**



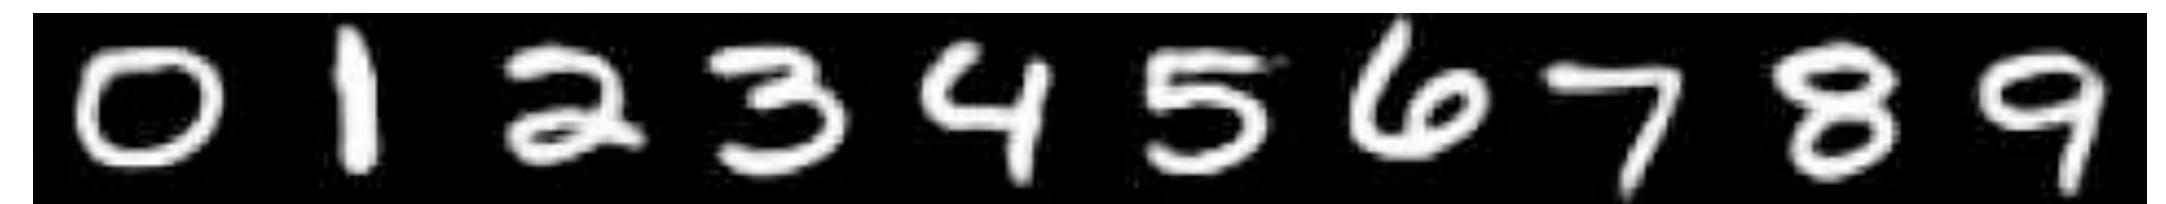
# Outline

- ① Challenges of image processing
- ② Introduction to convolutions
- ③ Convolutional layers
- ④ CNN building blocks
- ⑤ **A short history of CNNs**
- ⑥ Outlook

# Early CNN: LeNet

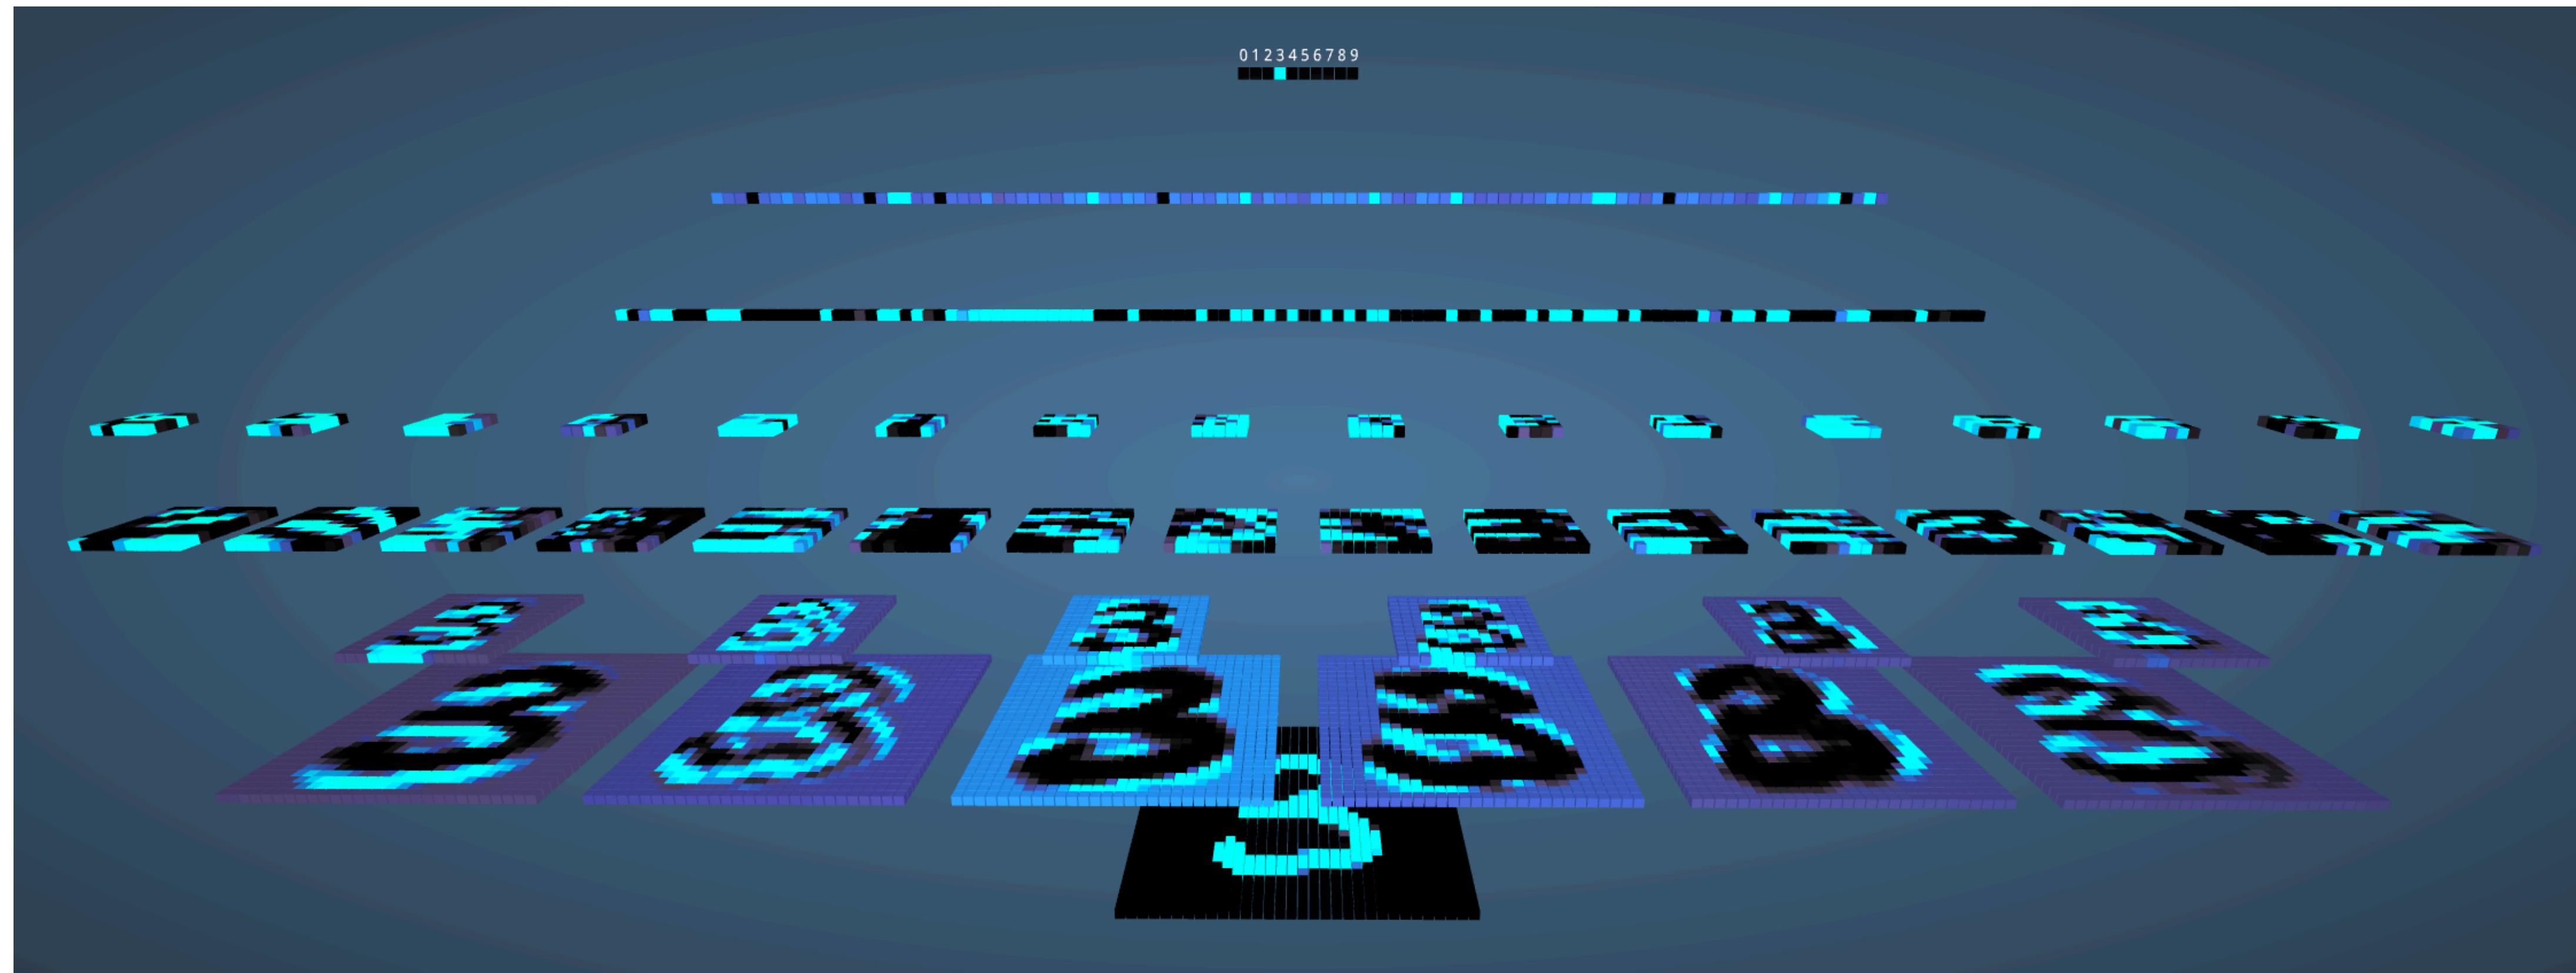


- Multiple versions developed by Yann LeCun between 1989 - 1995
- First **CNN trained using backpropagation**
- Classification of handwritten digits
- 1994: MNIST dataset created  
→ still common toy problem today



- Left: LeNet-5 architecture (1995)

# Interactive 3D visualization

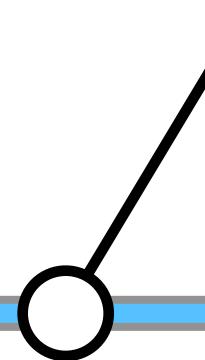


[https://adamharley.com/nn\\_vis/cnn/3d.html](https://adamharley.com/nn_vis/cnn/3d.html)

# A short history of CNNs

**LeNet**

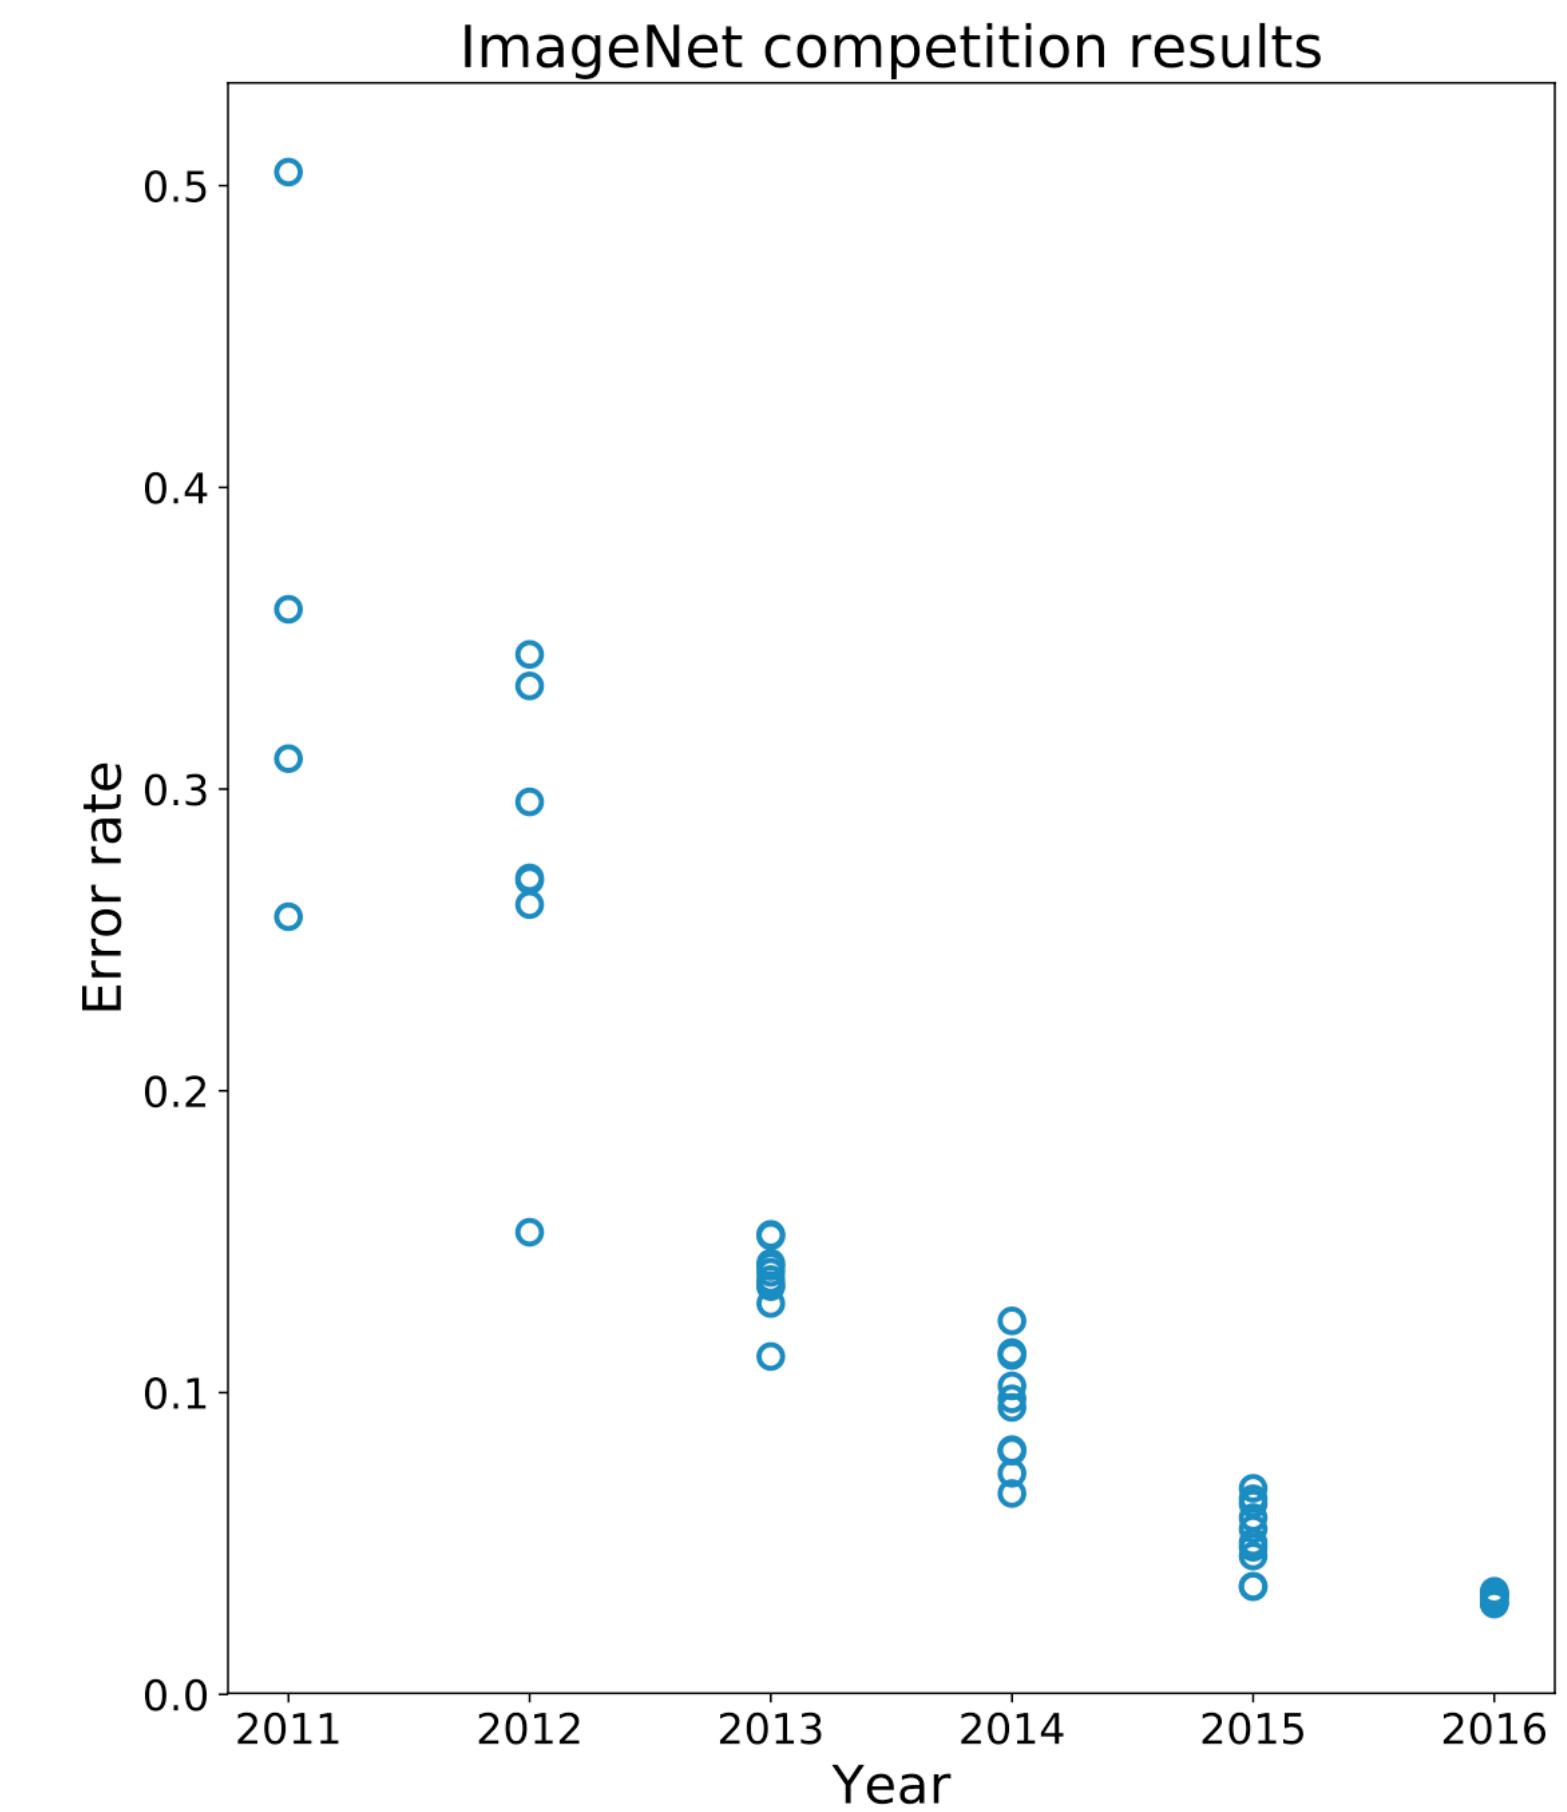
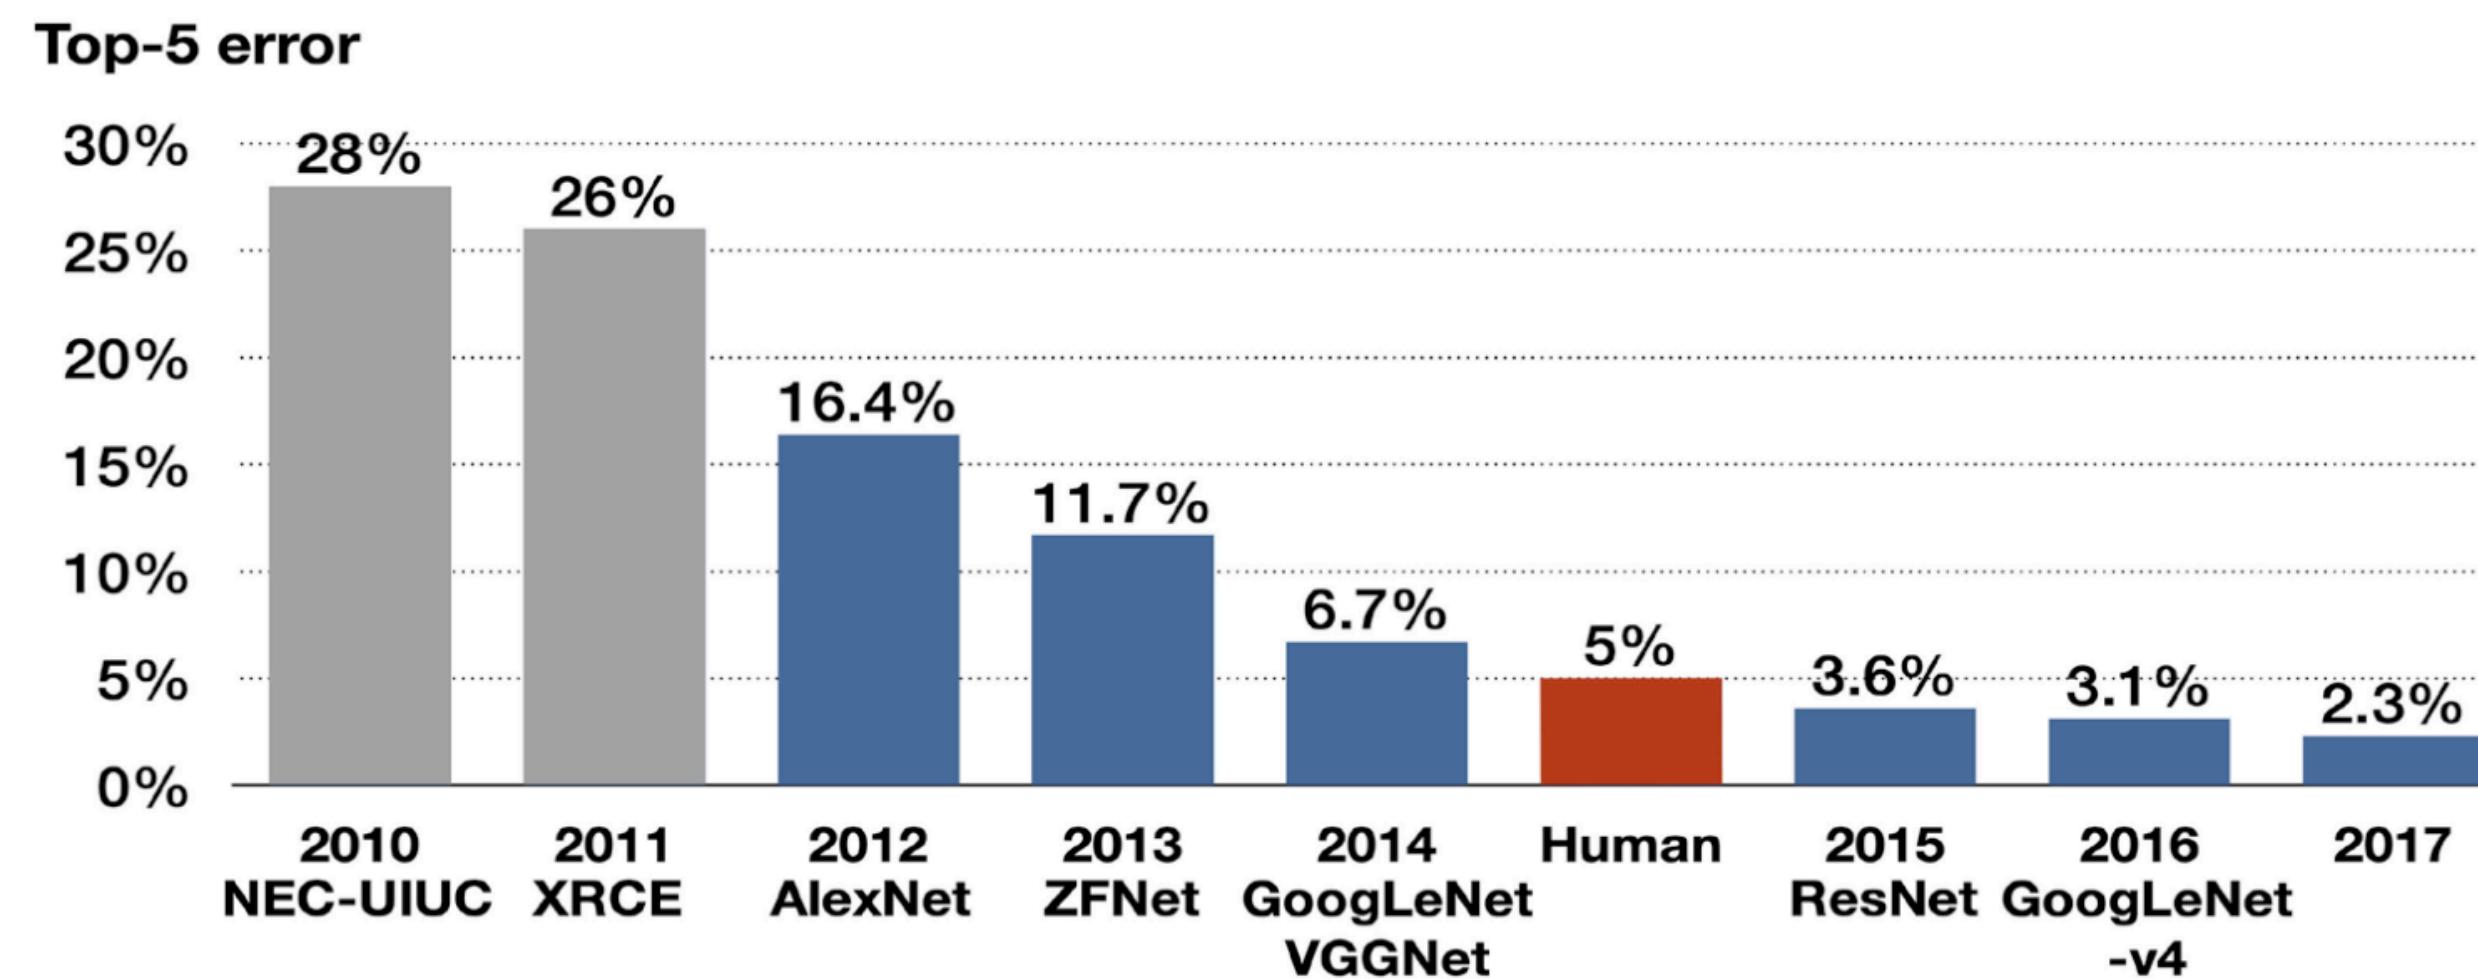
First CNN trained  
using backpropagation



1989 - 1995

# ImageNet challenge

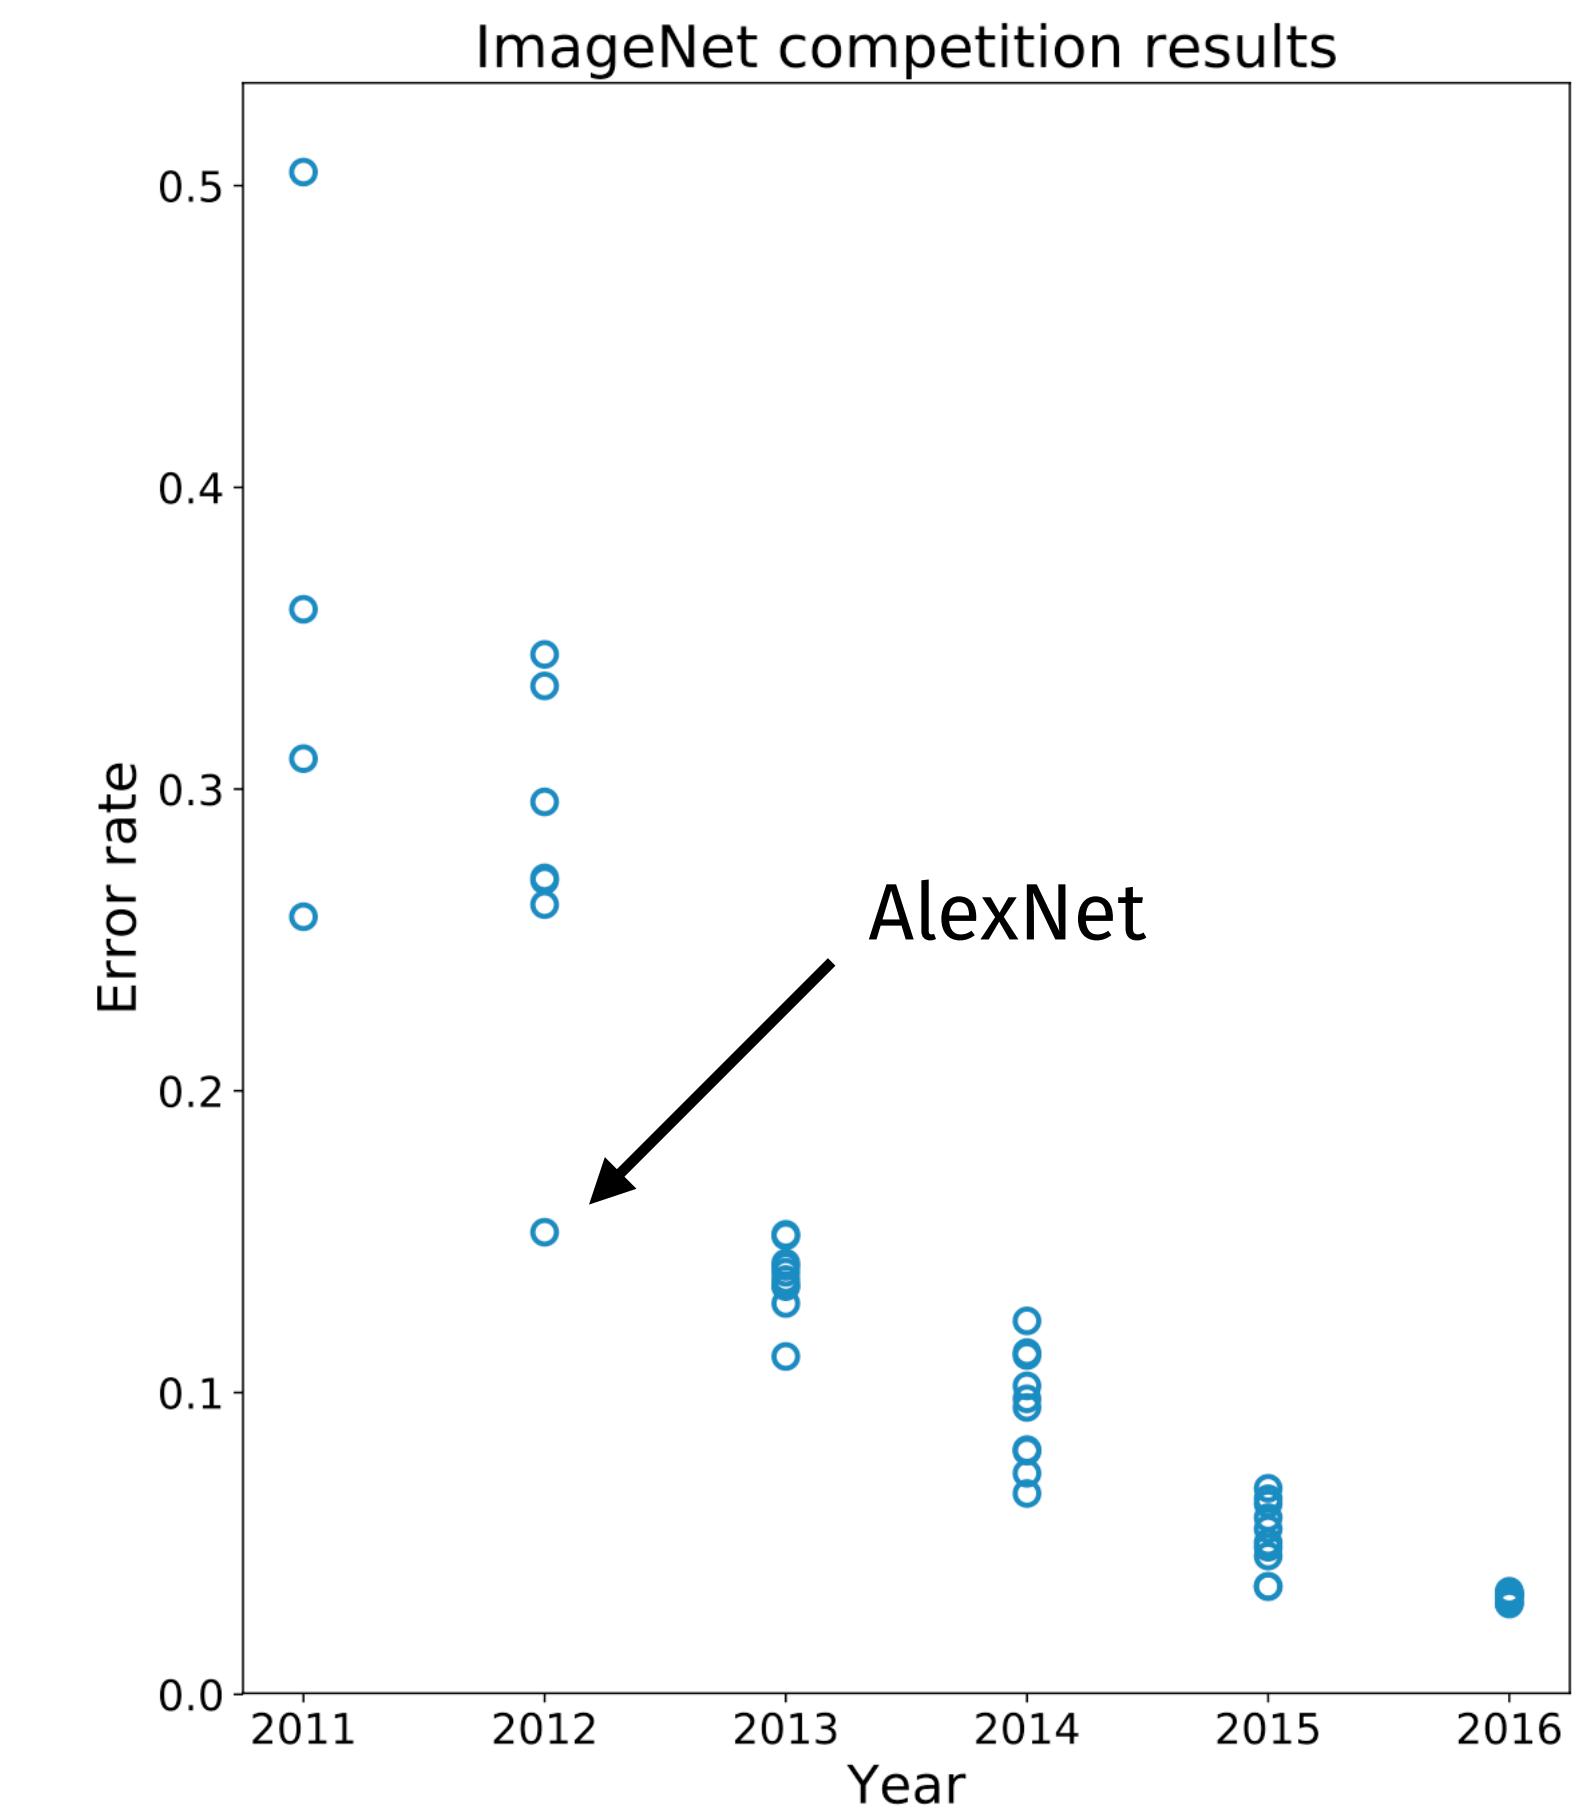
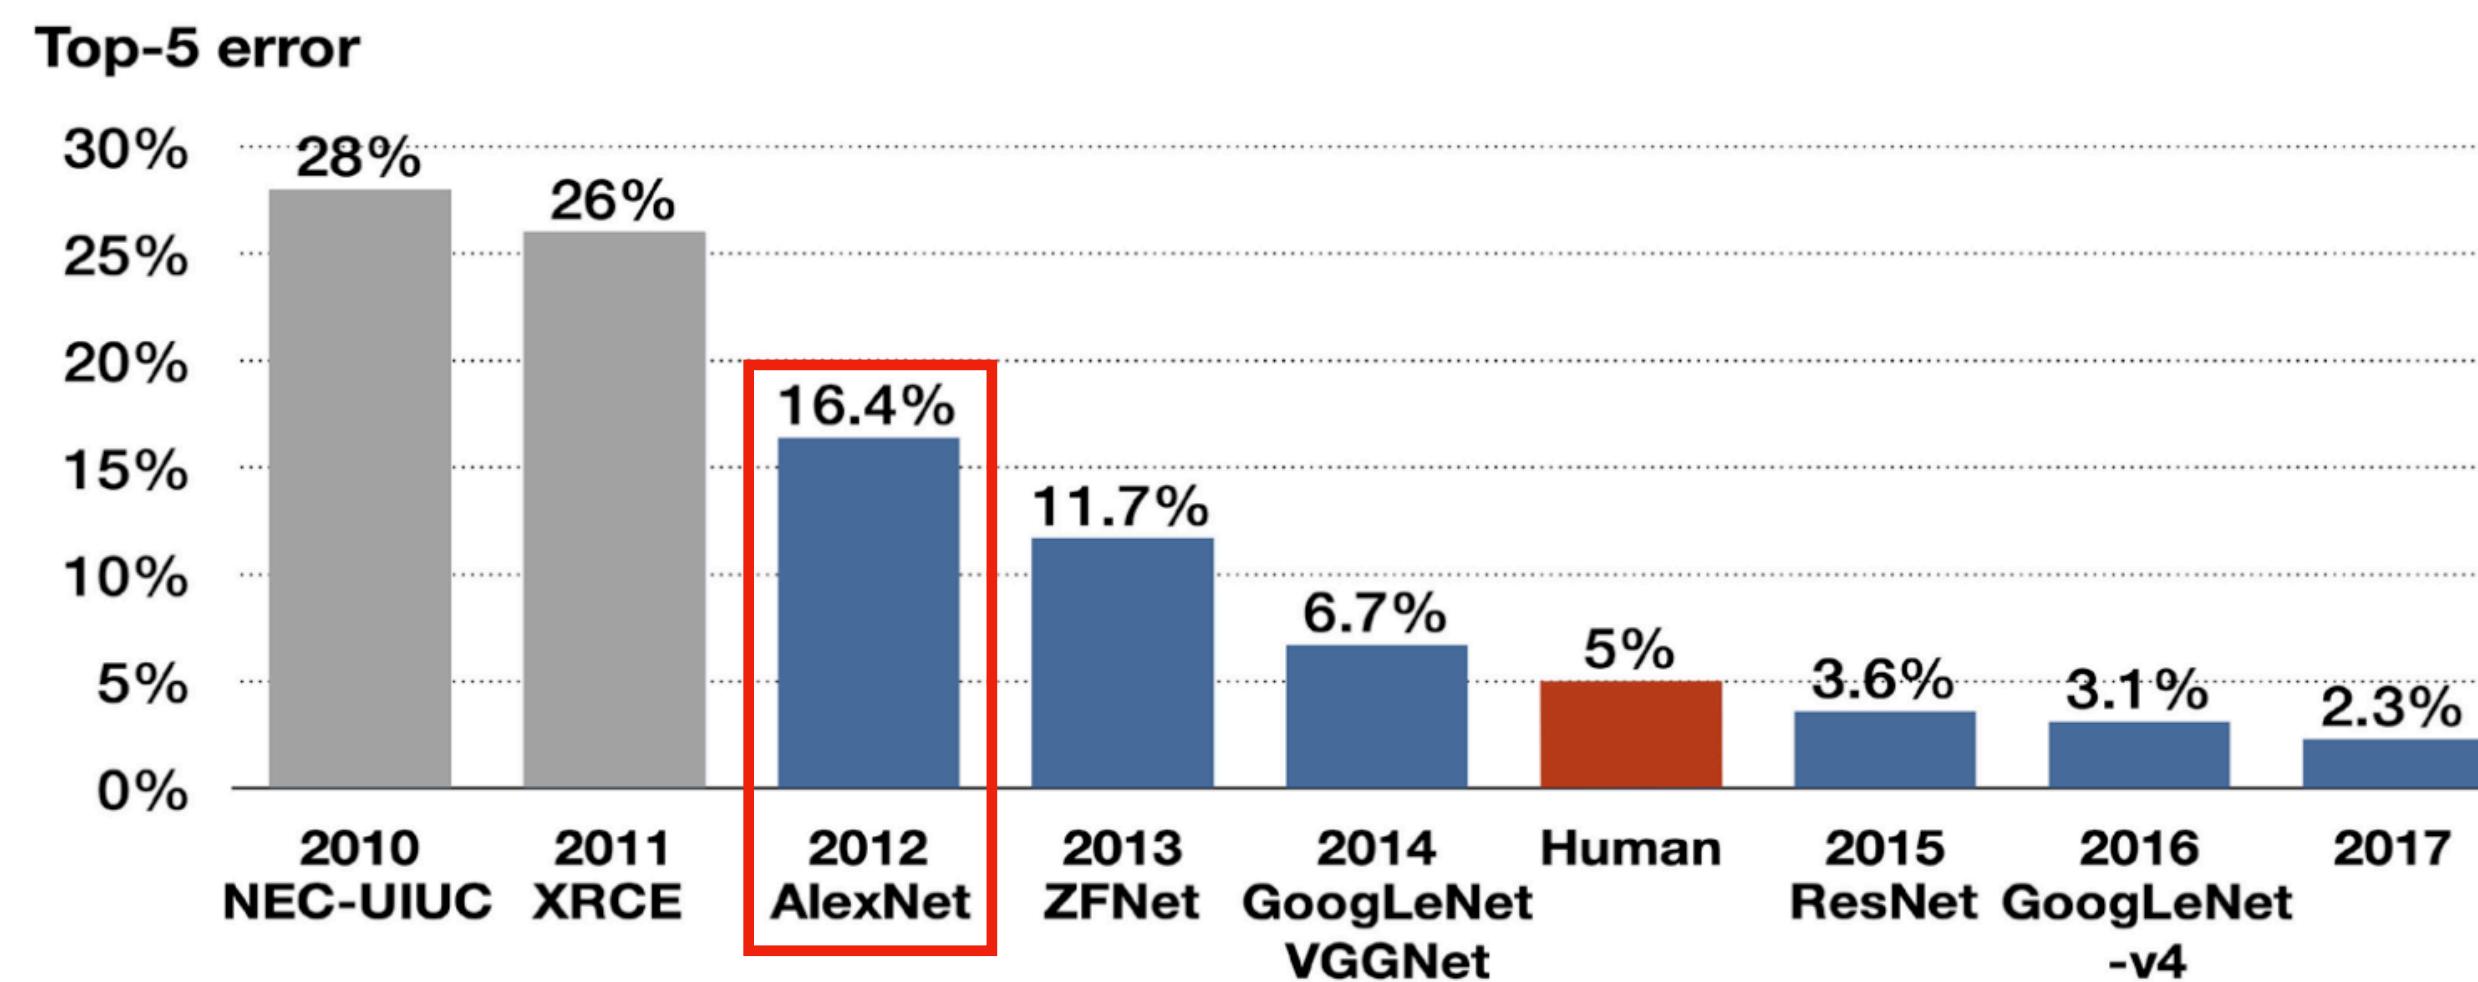
- ImageNet dataset: more than 14M images
- hand-labeled, more than 10k categories
- Since 2010: annual ImageNet challenge  
→ classification with 1000 categories
- Metric: true label must be within top 5 of predicted labels



By Gkrusze - Own work, CC BY-SA 4.0,  
<https://commons.wikimedia.org/w/index.php?curid=69750373>

# ImageNet challenge

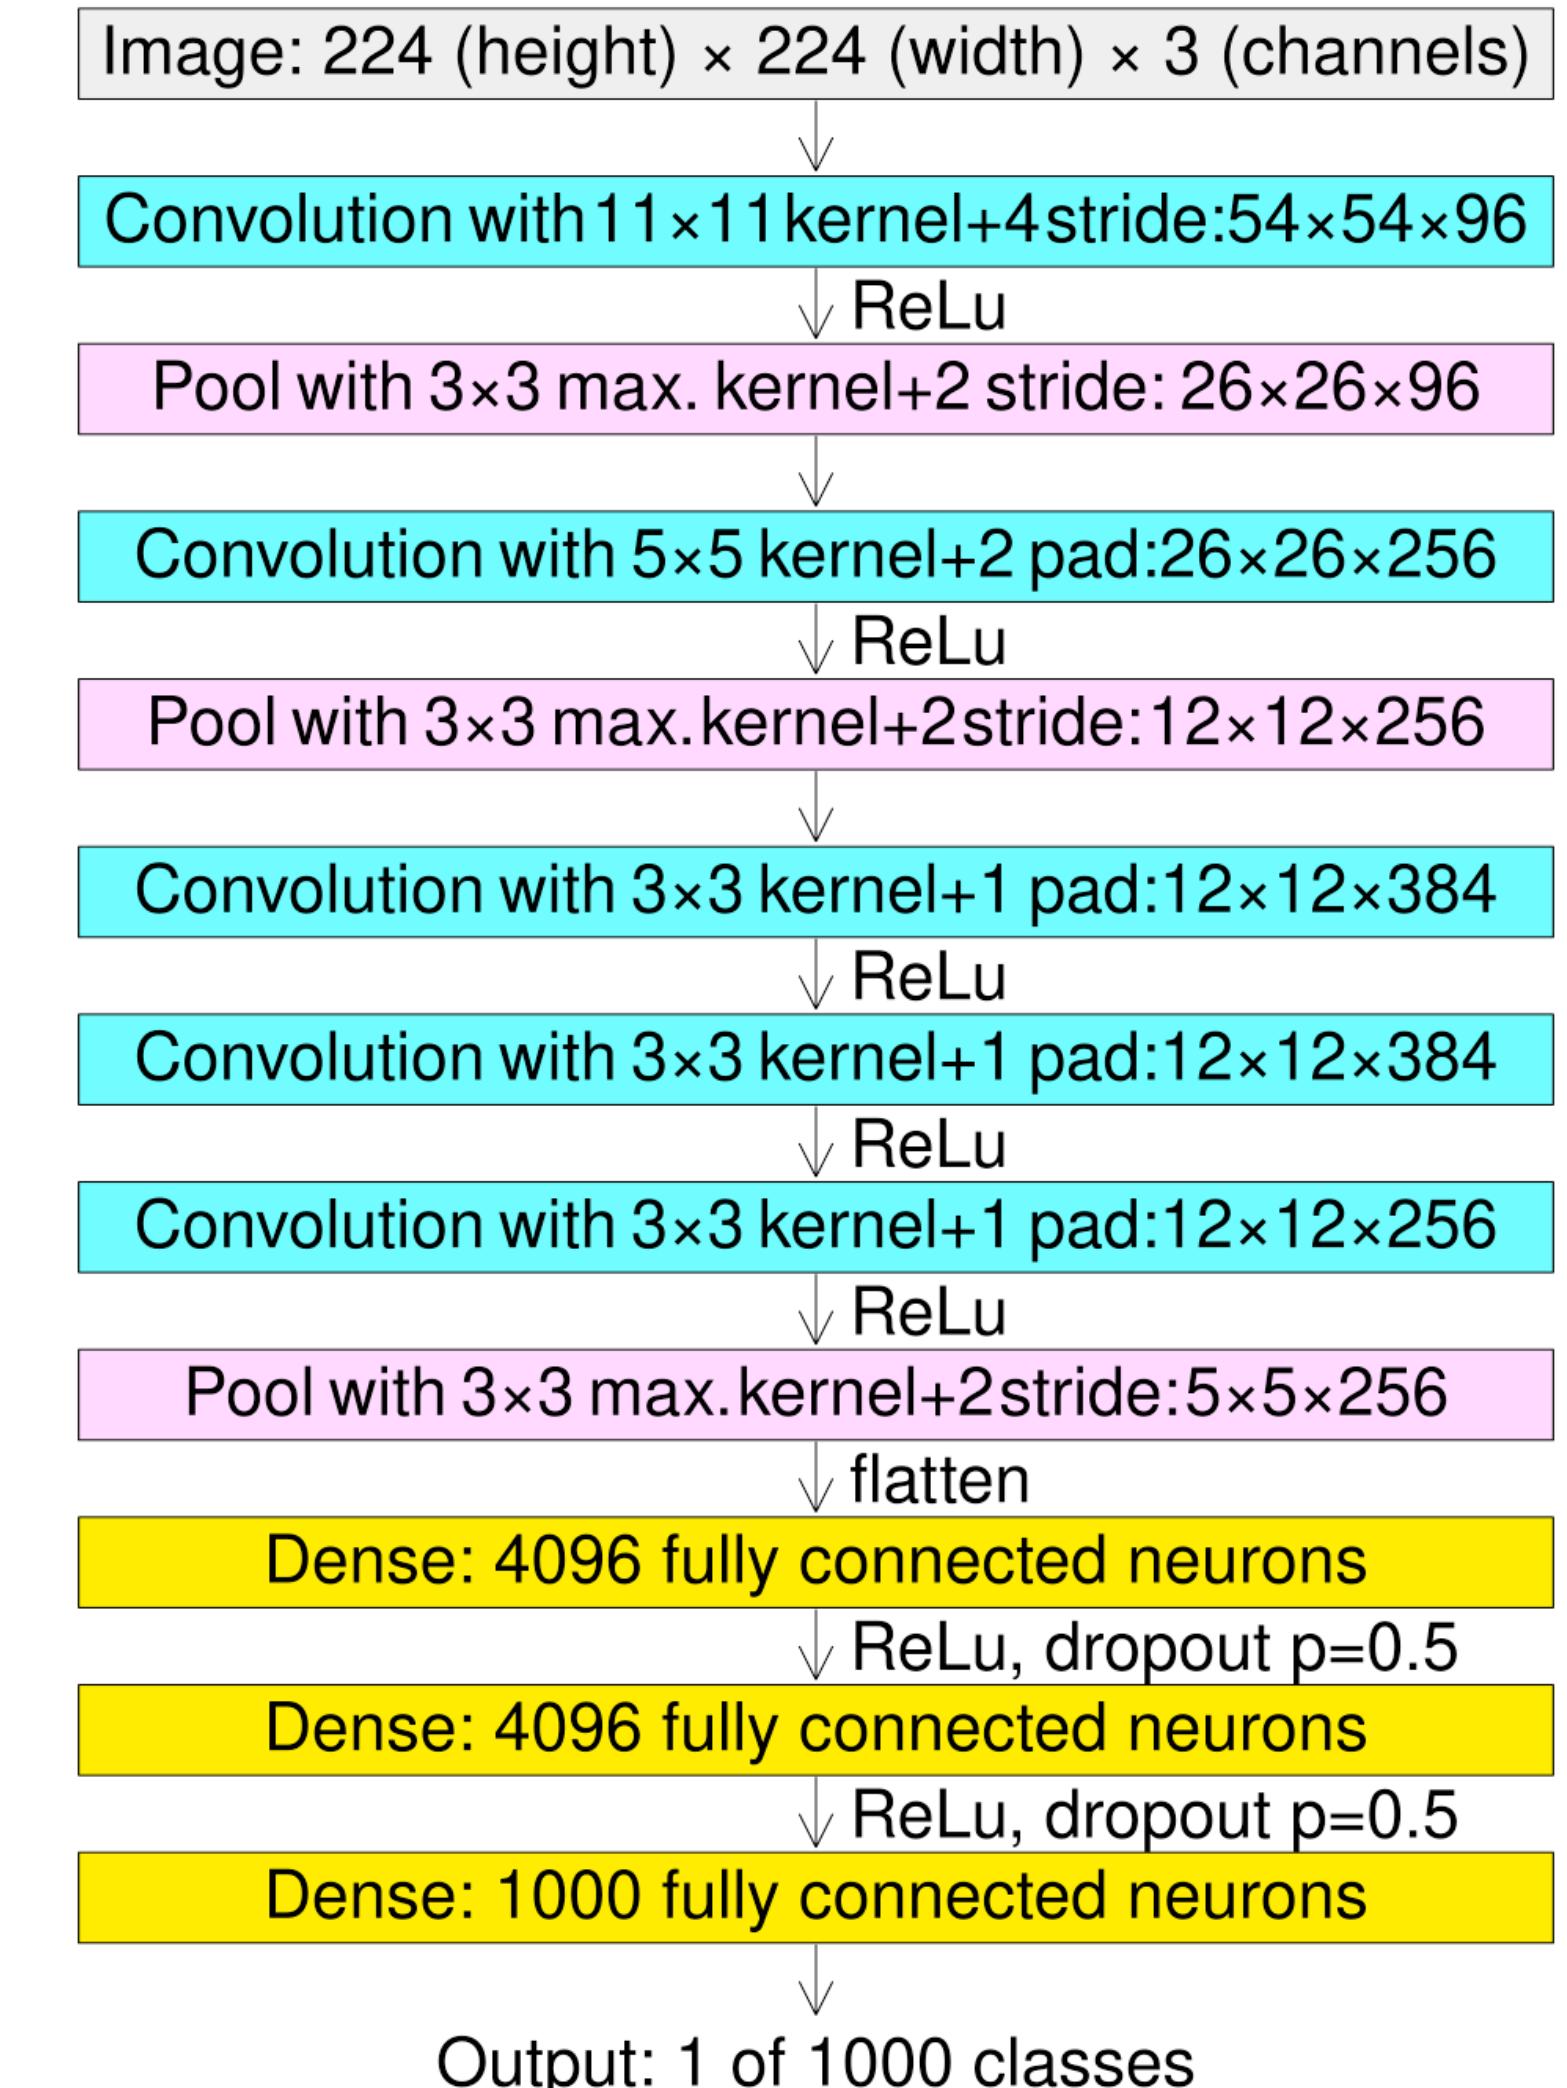
- ImageNet dataset: more than 14M images
- hand-labeled, more than 10k categories
- Since 2010: annual ImageNet challenge  
→ classification with 1000 categories
- Metric: true label must be within top 5 of predicted labels



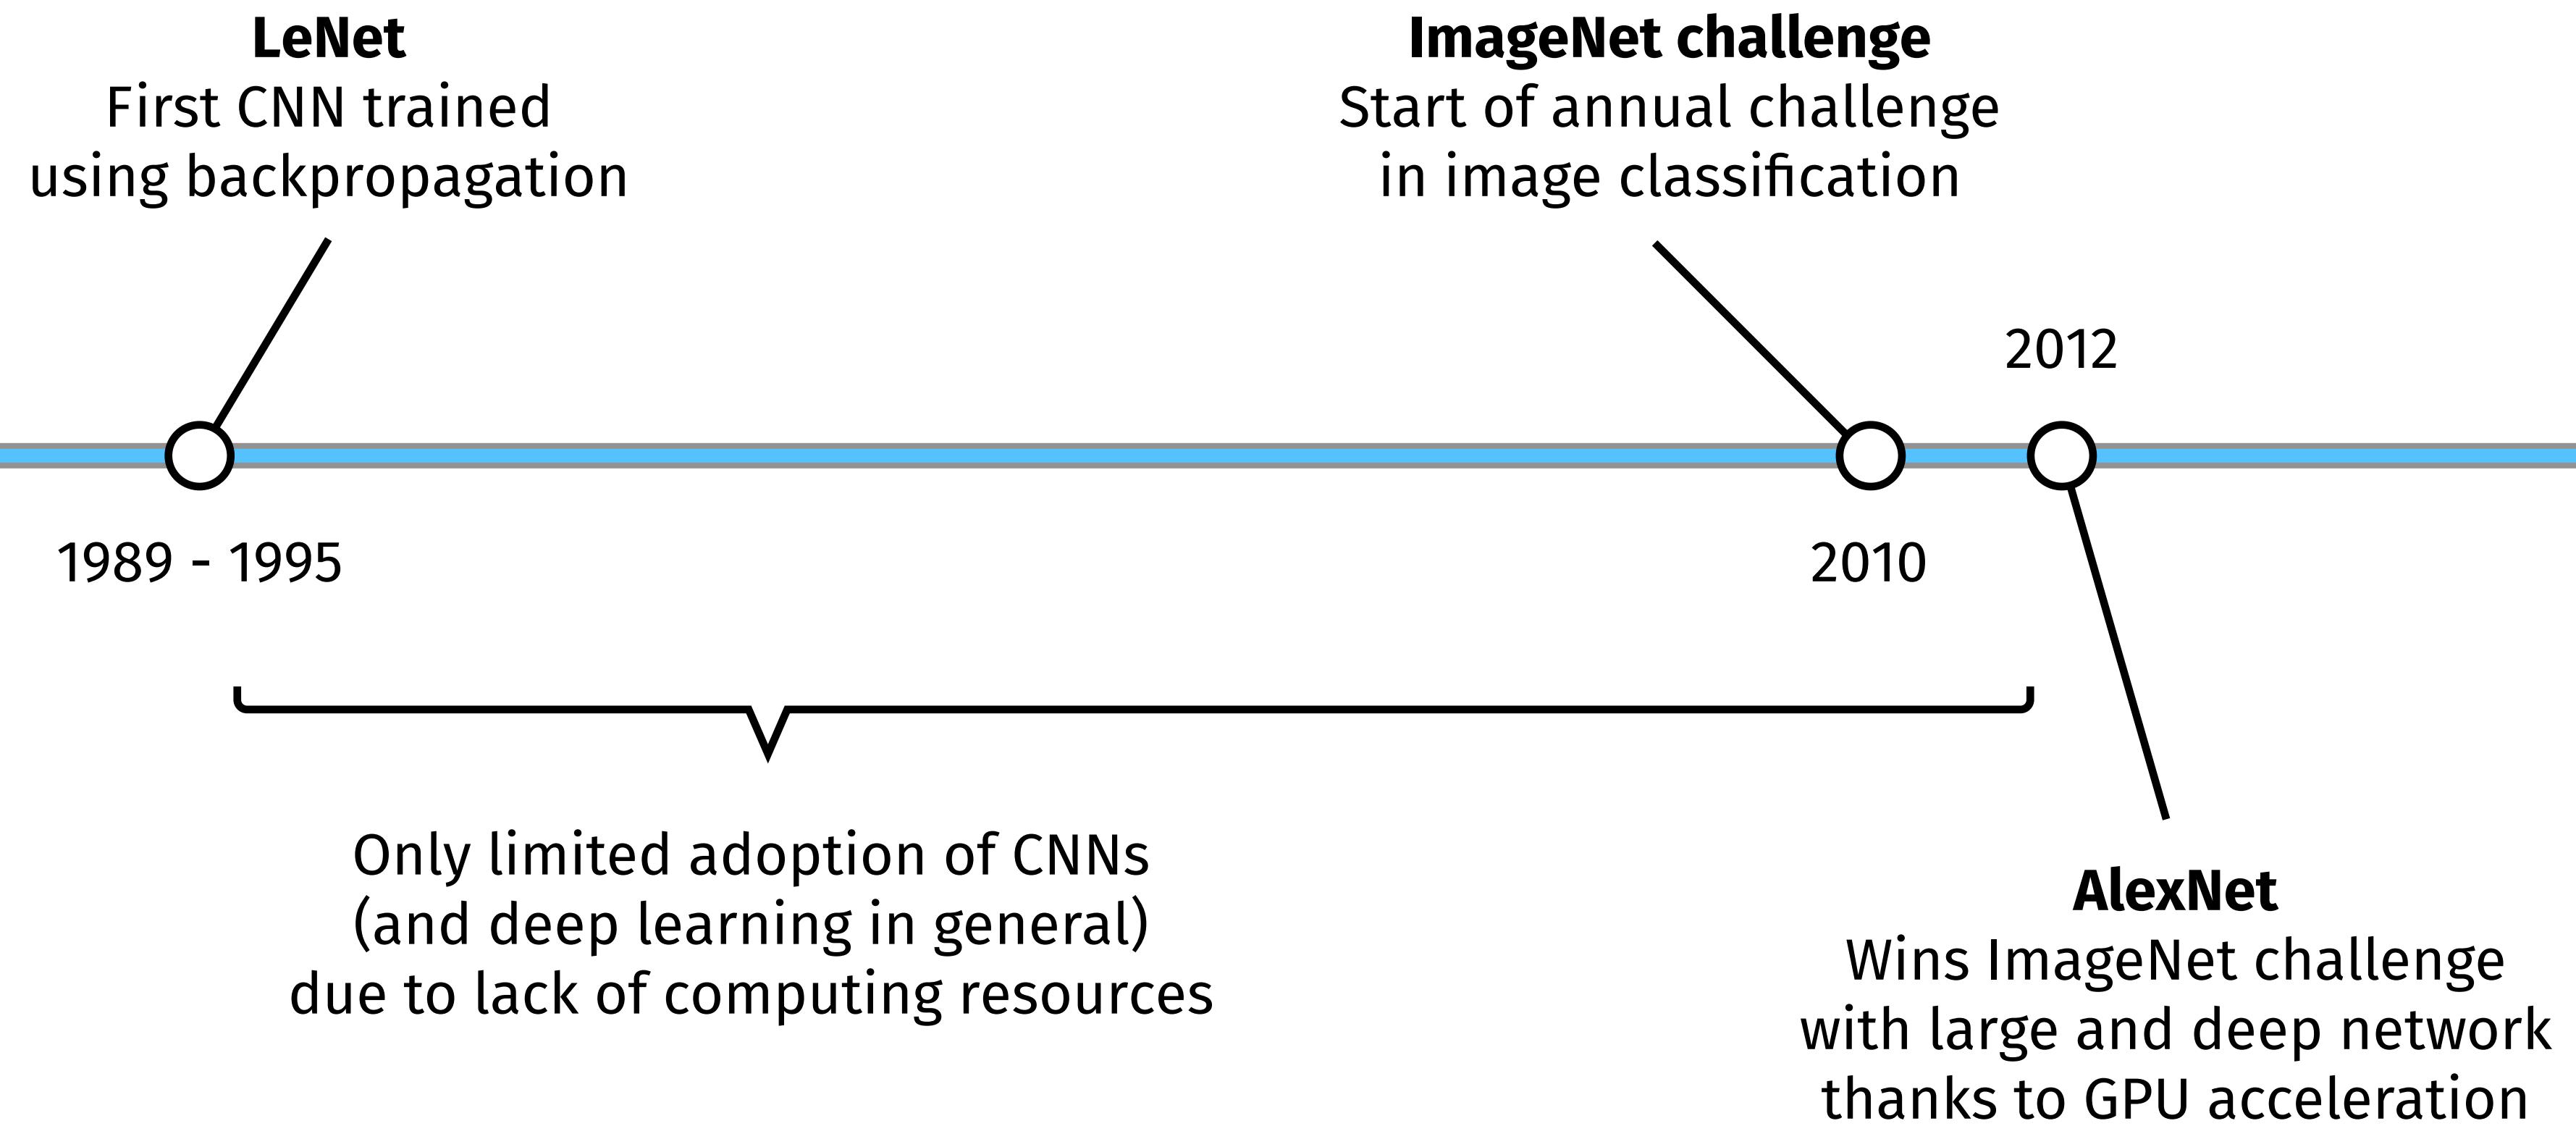
By Gkrusze - Own work, CC BY-SA 4.0,  
<https://commons.wikimedia.org/w/index.php?curid=69750373>

# AlexNet

- Developed by Alex Krizhevsky et al
- Similar to LeNet structurally, but much wider and deeper
- Among the **first to use GPUs for NNs**  
→ **much larger network possible**
- Won 2012 ImageNet challenge with much lower error rate than competitors
- Highly influential, often claimed to have started today's deep learning boom

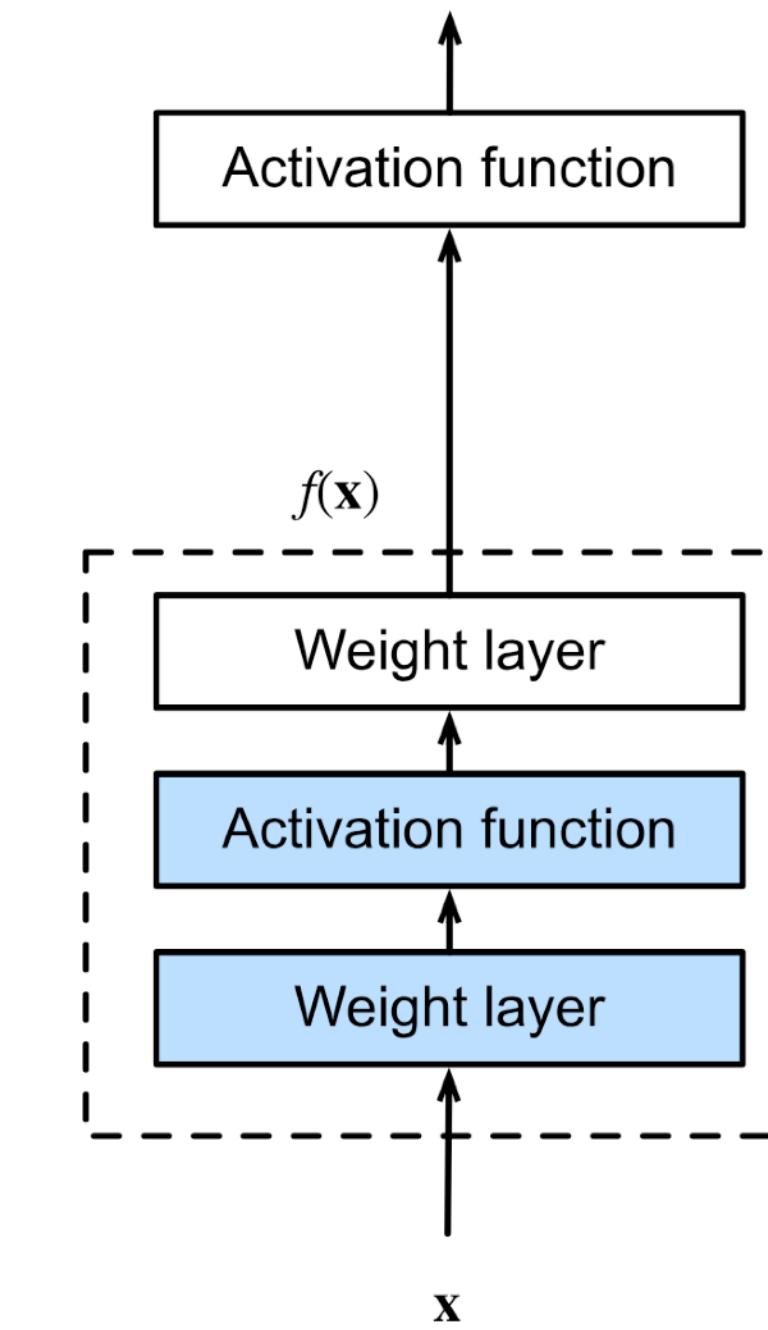


# A short history of CNNs



# ResNet

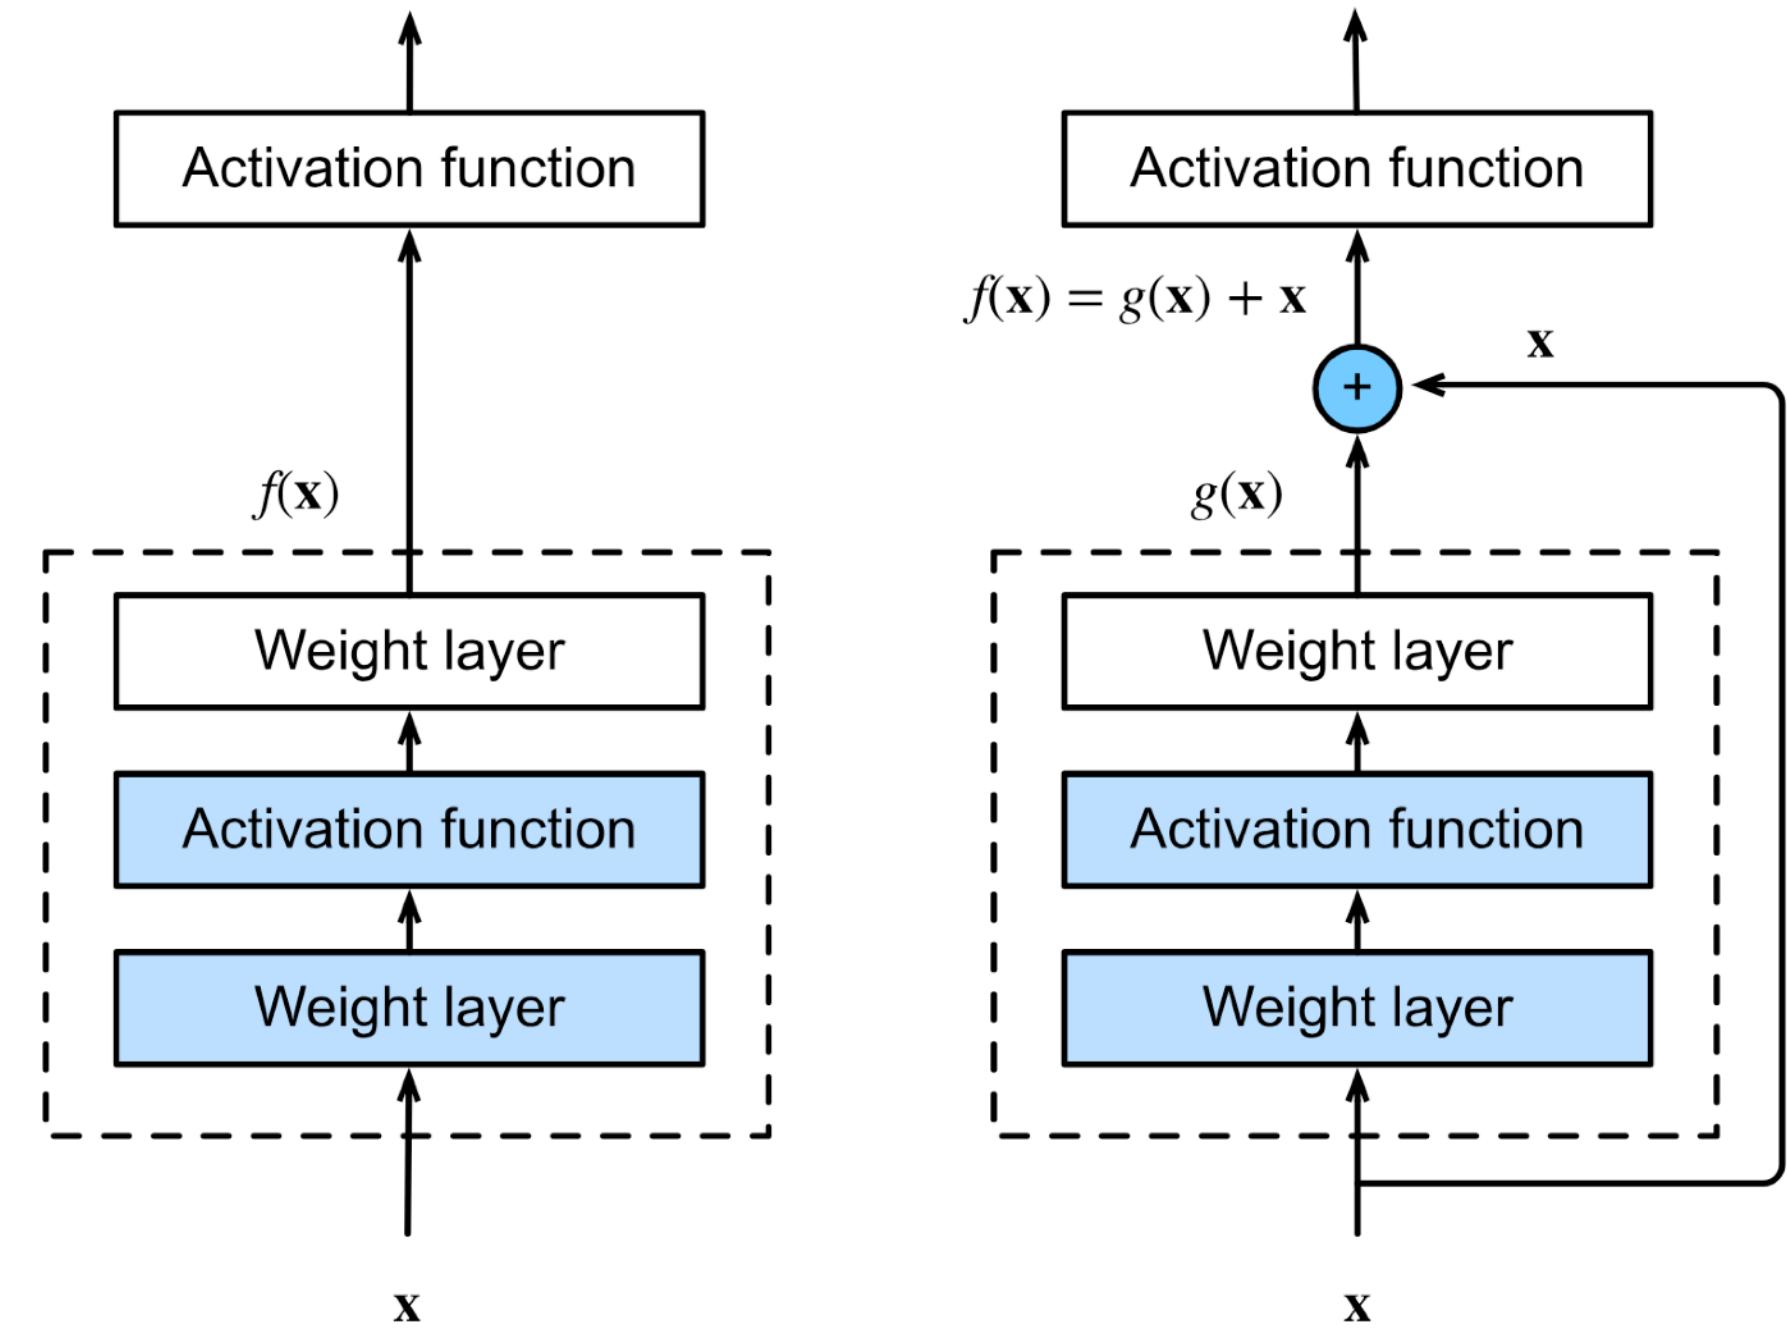
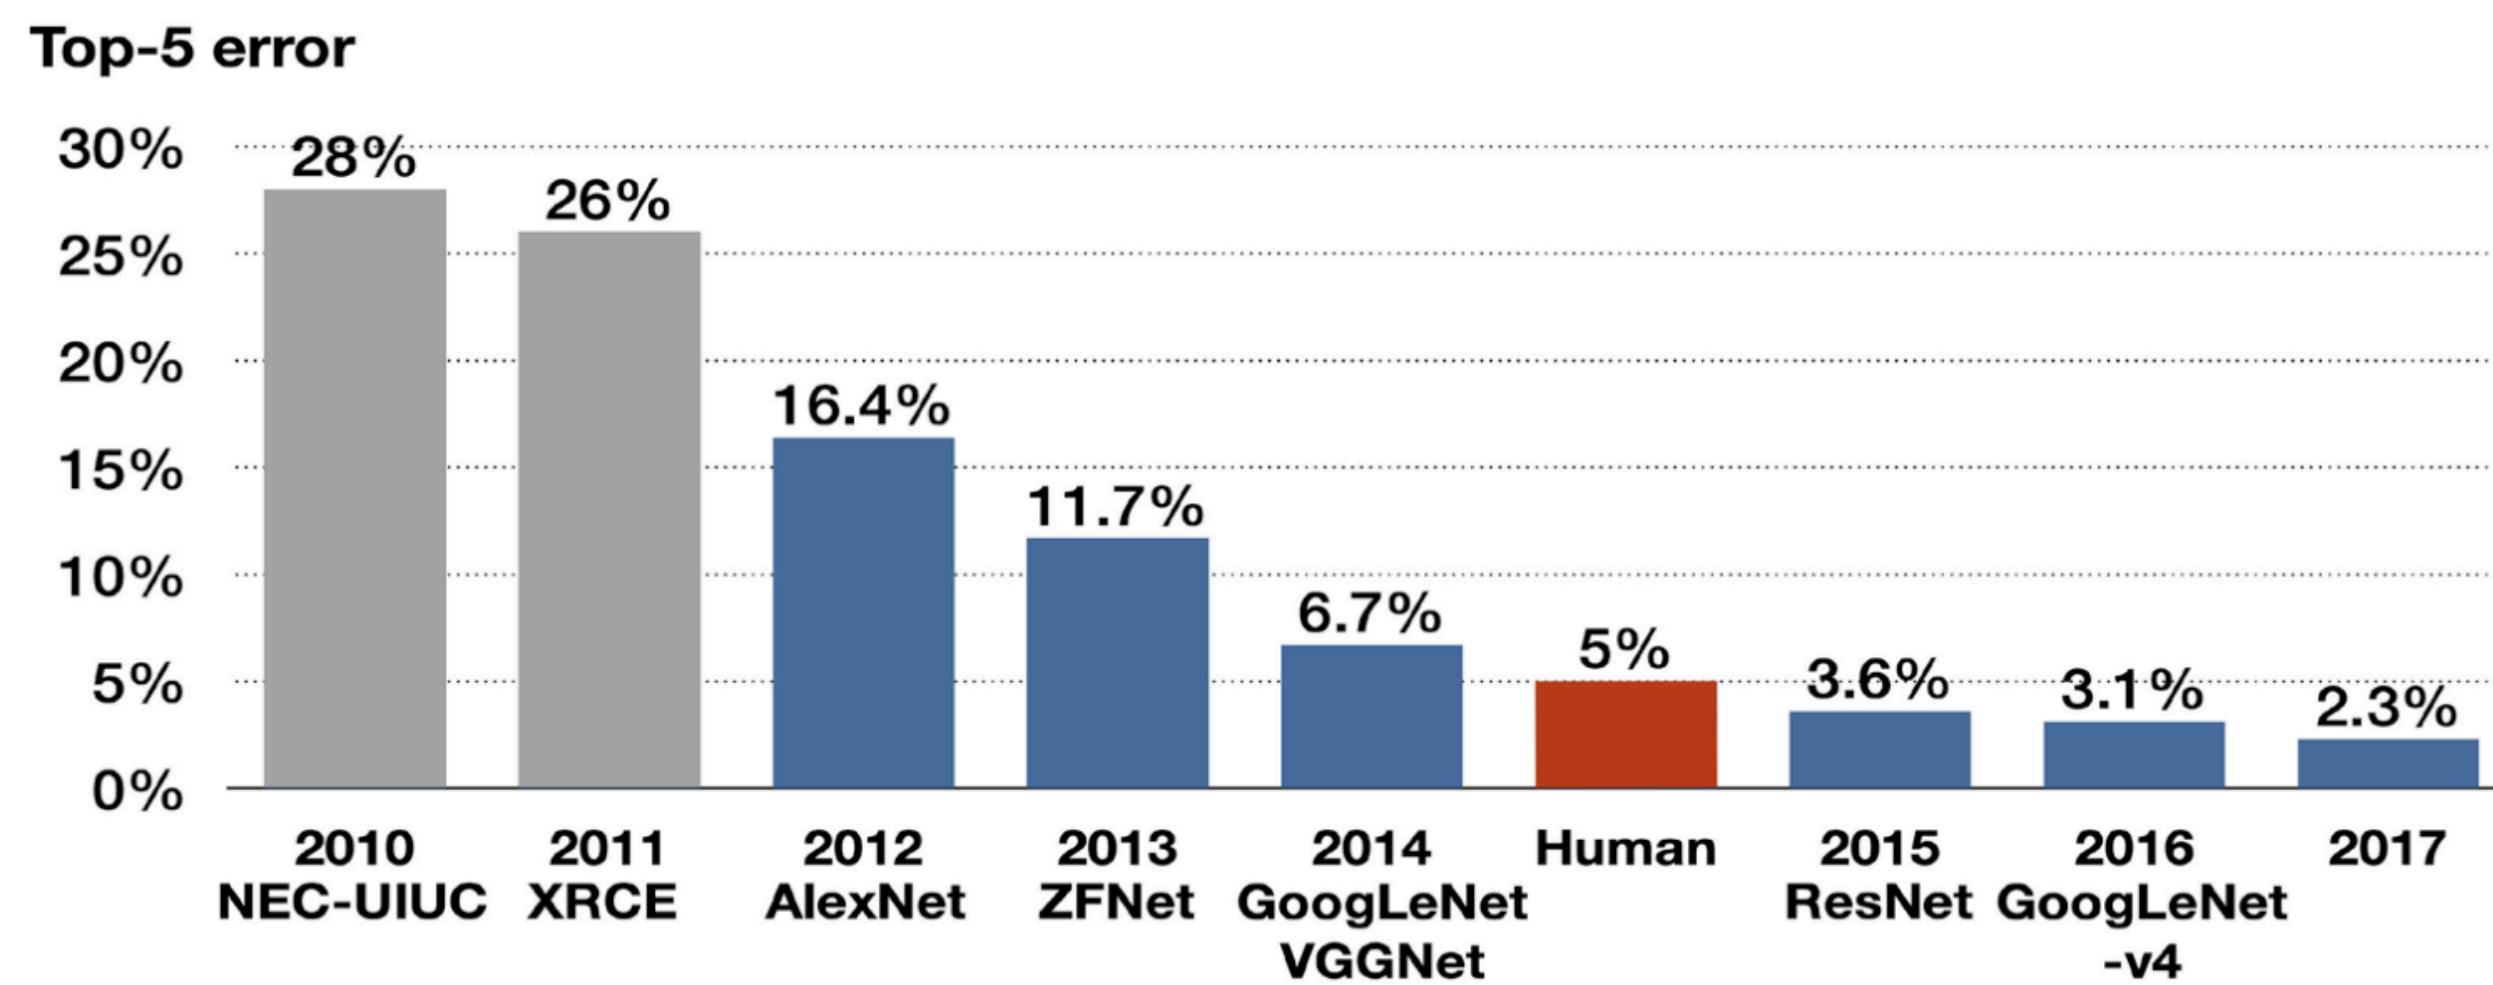
- Standard CNN: **depth limited** due to **vanishing gradients problem**
- The issue:
  - negative weights
  - negative outputs
  - ReLU activation yields 0
  - gradient is 0
  - optimizer does not adapt weights
- More likely to happen for deep networks



<https://github.com/d2l-ai/d2l-en>

# ResNet

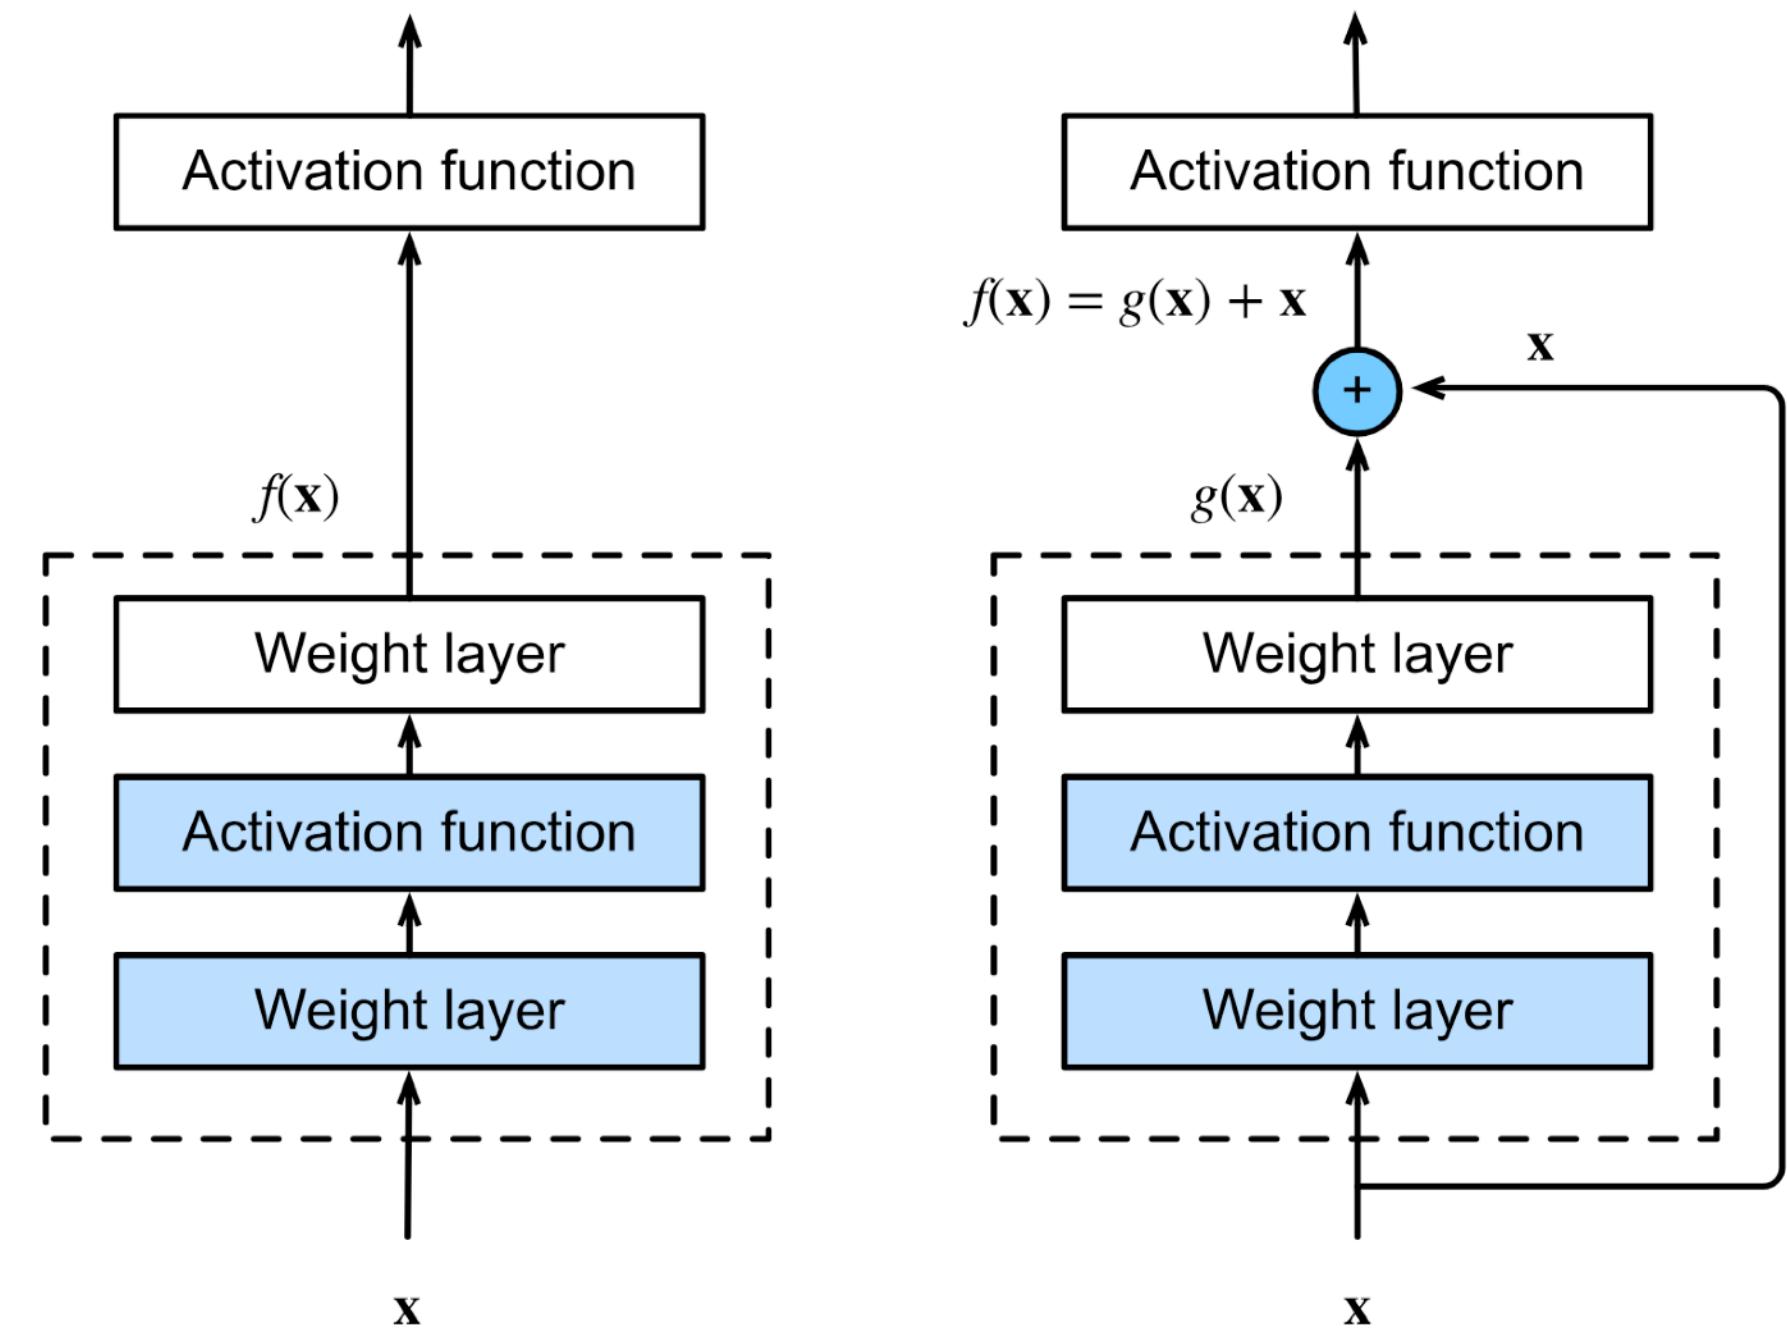
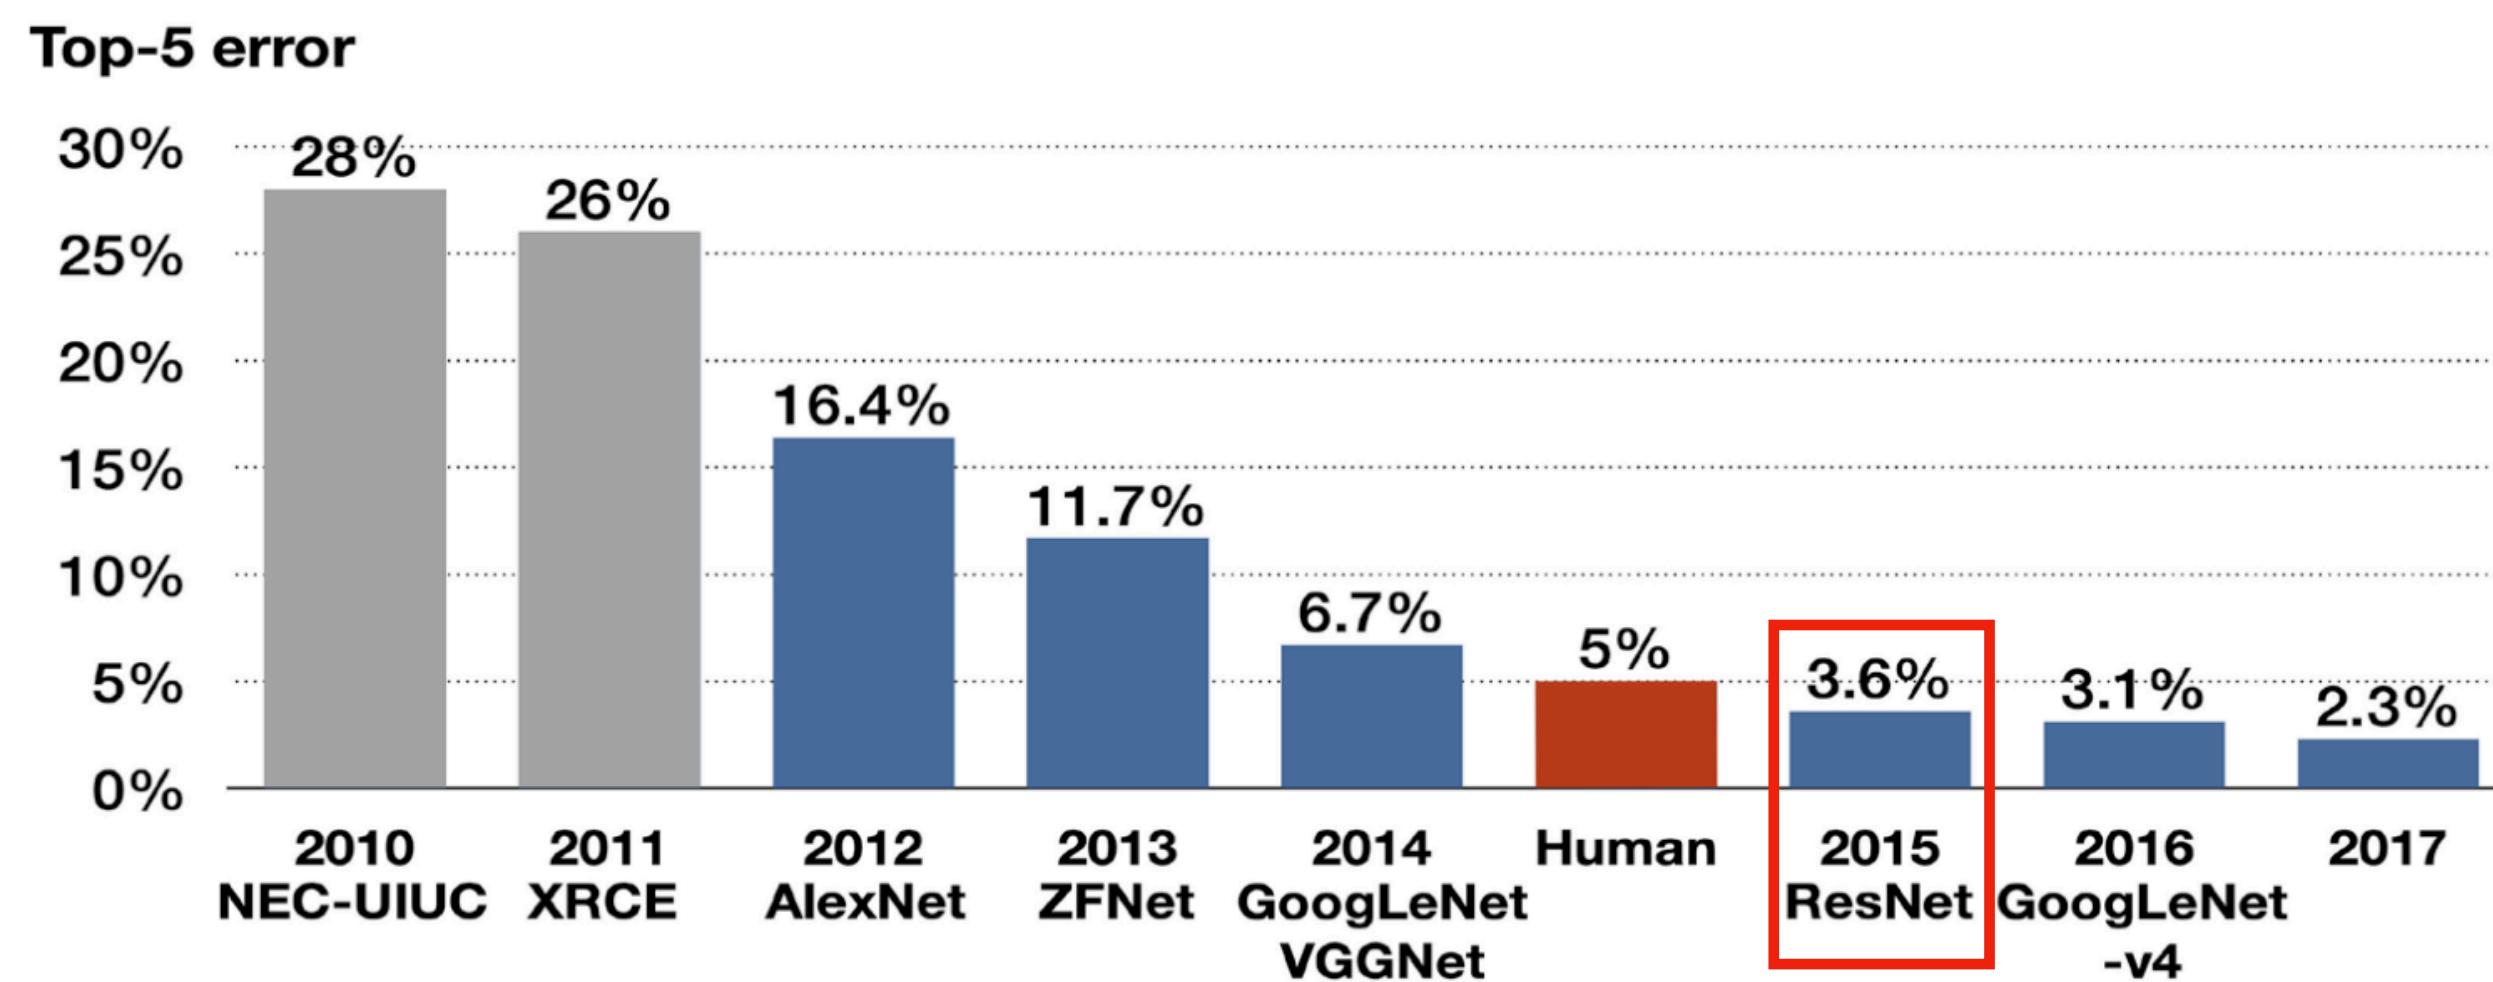
- Solution: add residual/skip connections
- Convolutions learn modification to input
- Gradients can propagate easily
  - deeper networks possible
  - can build more expressive models
- Successful in 2015 ImageNet challenge



<https://github.com/d2l-ai/d2l-en>

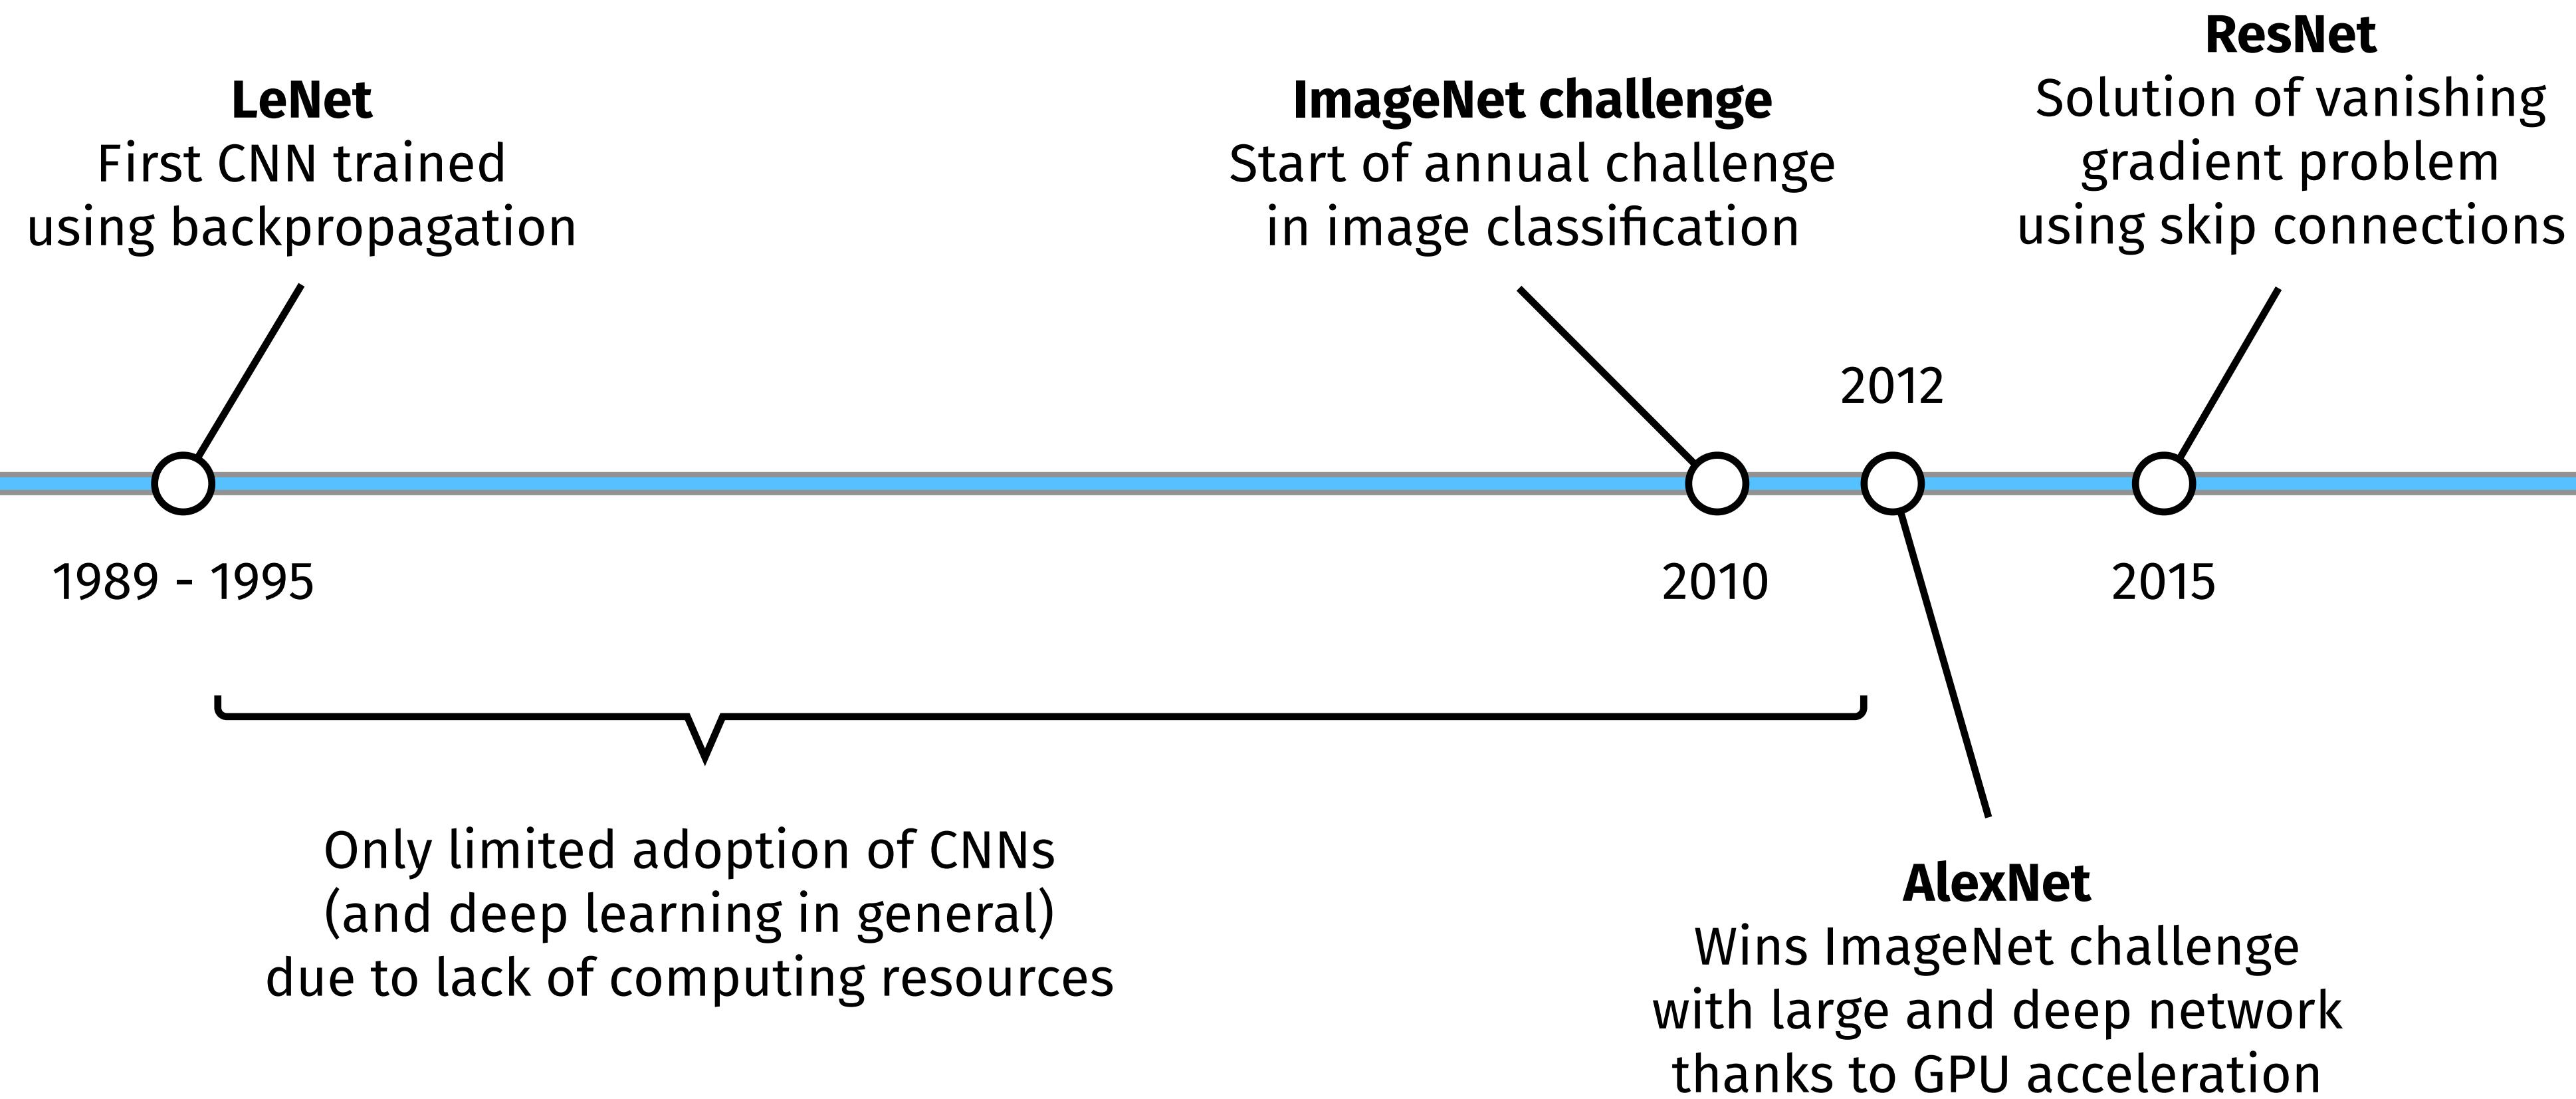
# ResNet

- Solution: add residual/skip connections
- Convolutions learn modification to input
- Gradients can propagate easily
  - deeper networks possible
  - can build more expressive models
- Successful in 2015 ImageNet challenge

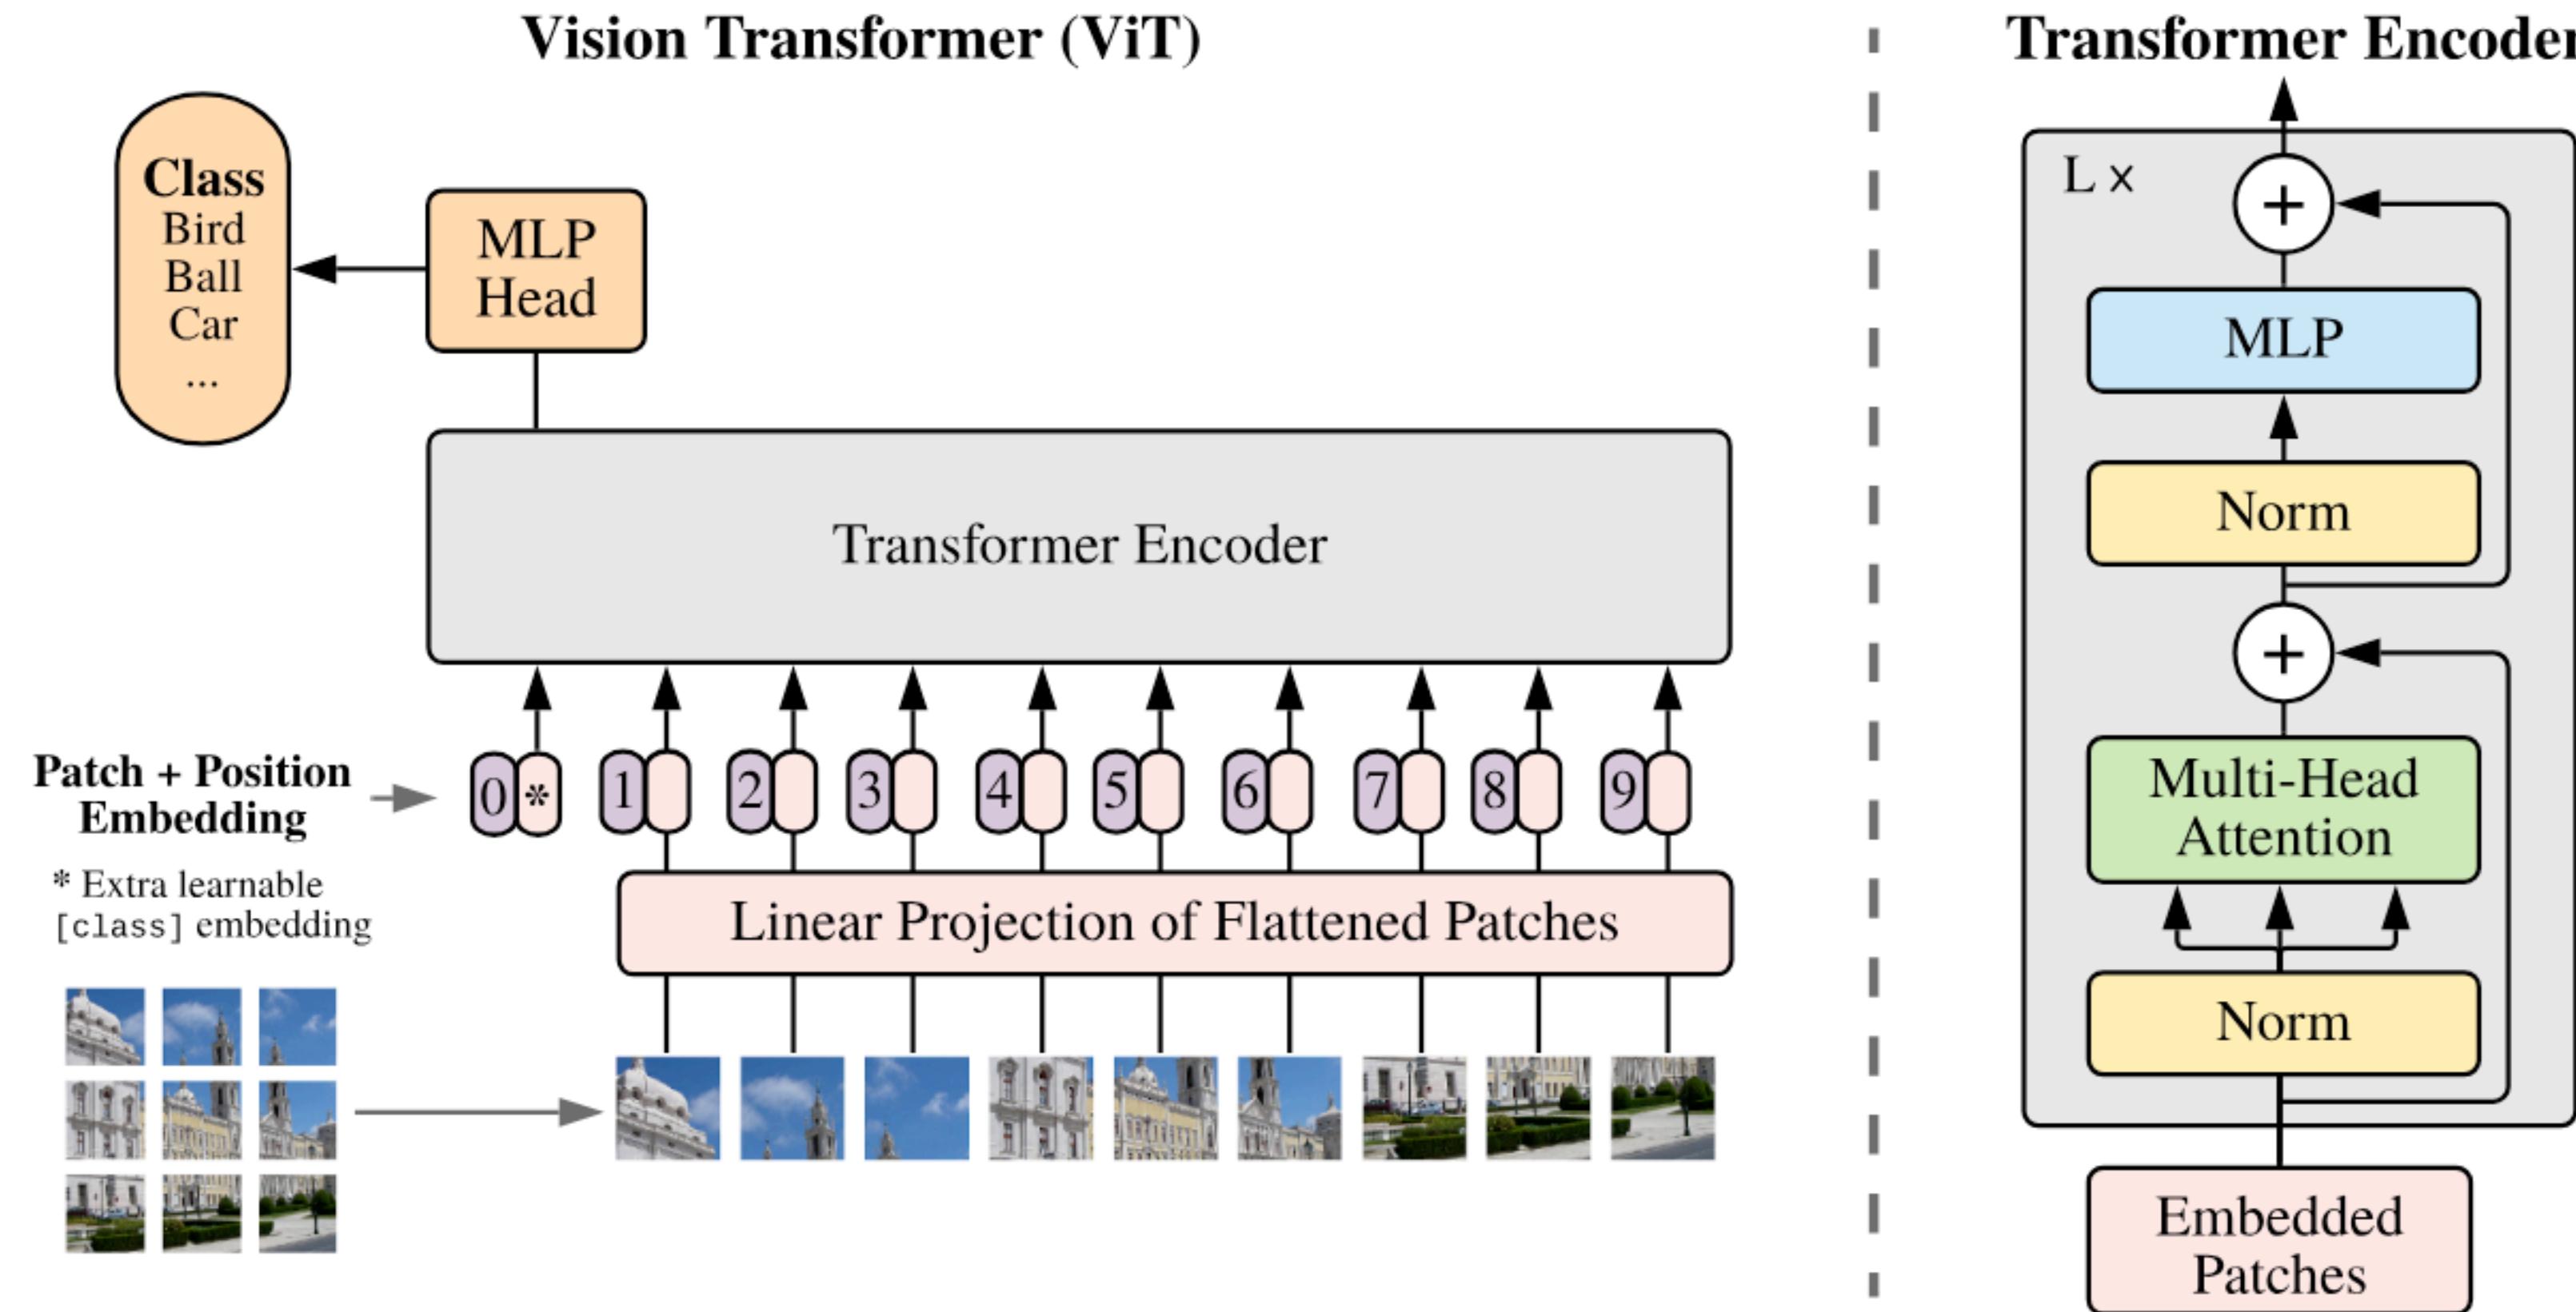


<https://github.com/d2l-ai/d2l-en>

# A short history of CNNs

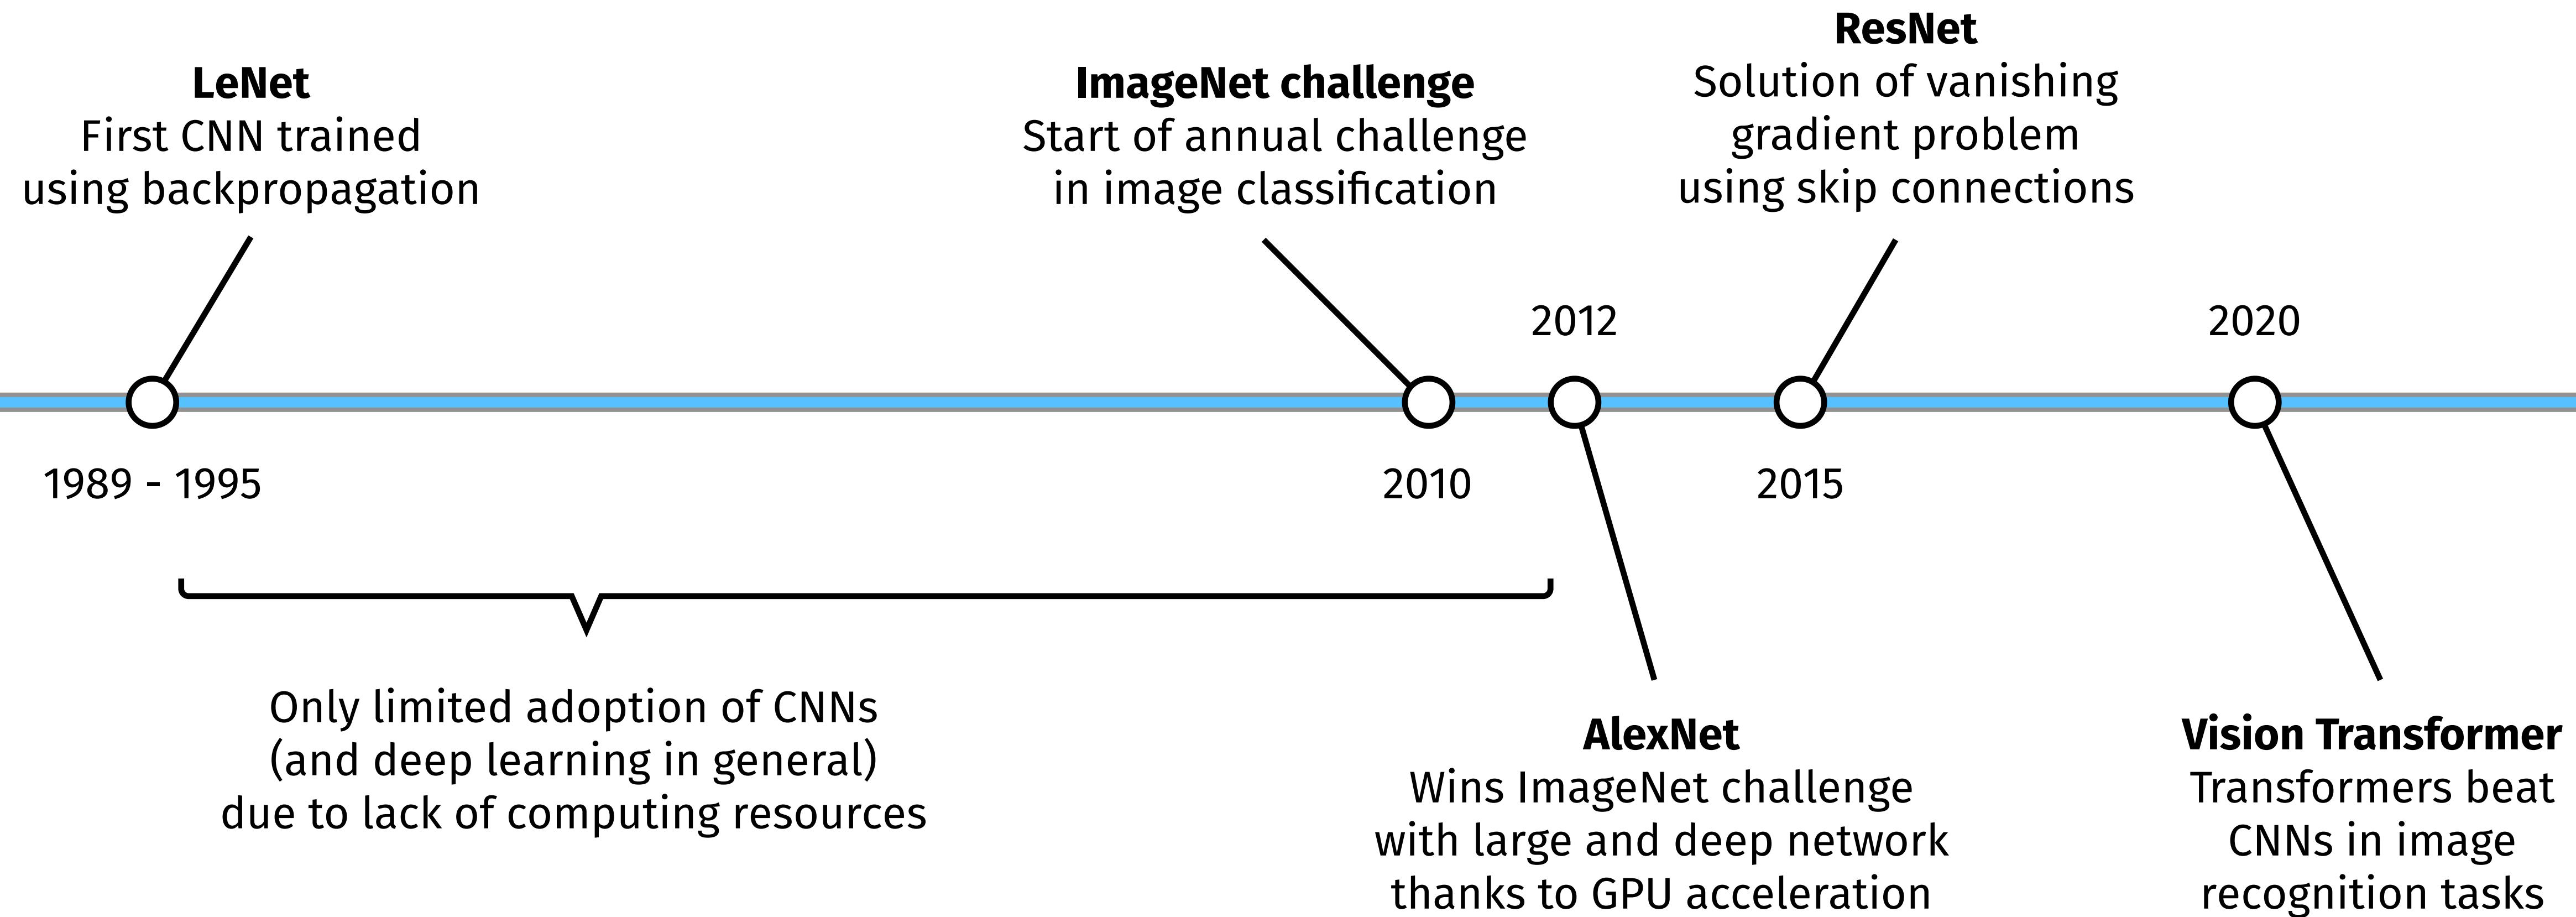


# Vision transformers

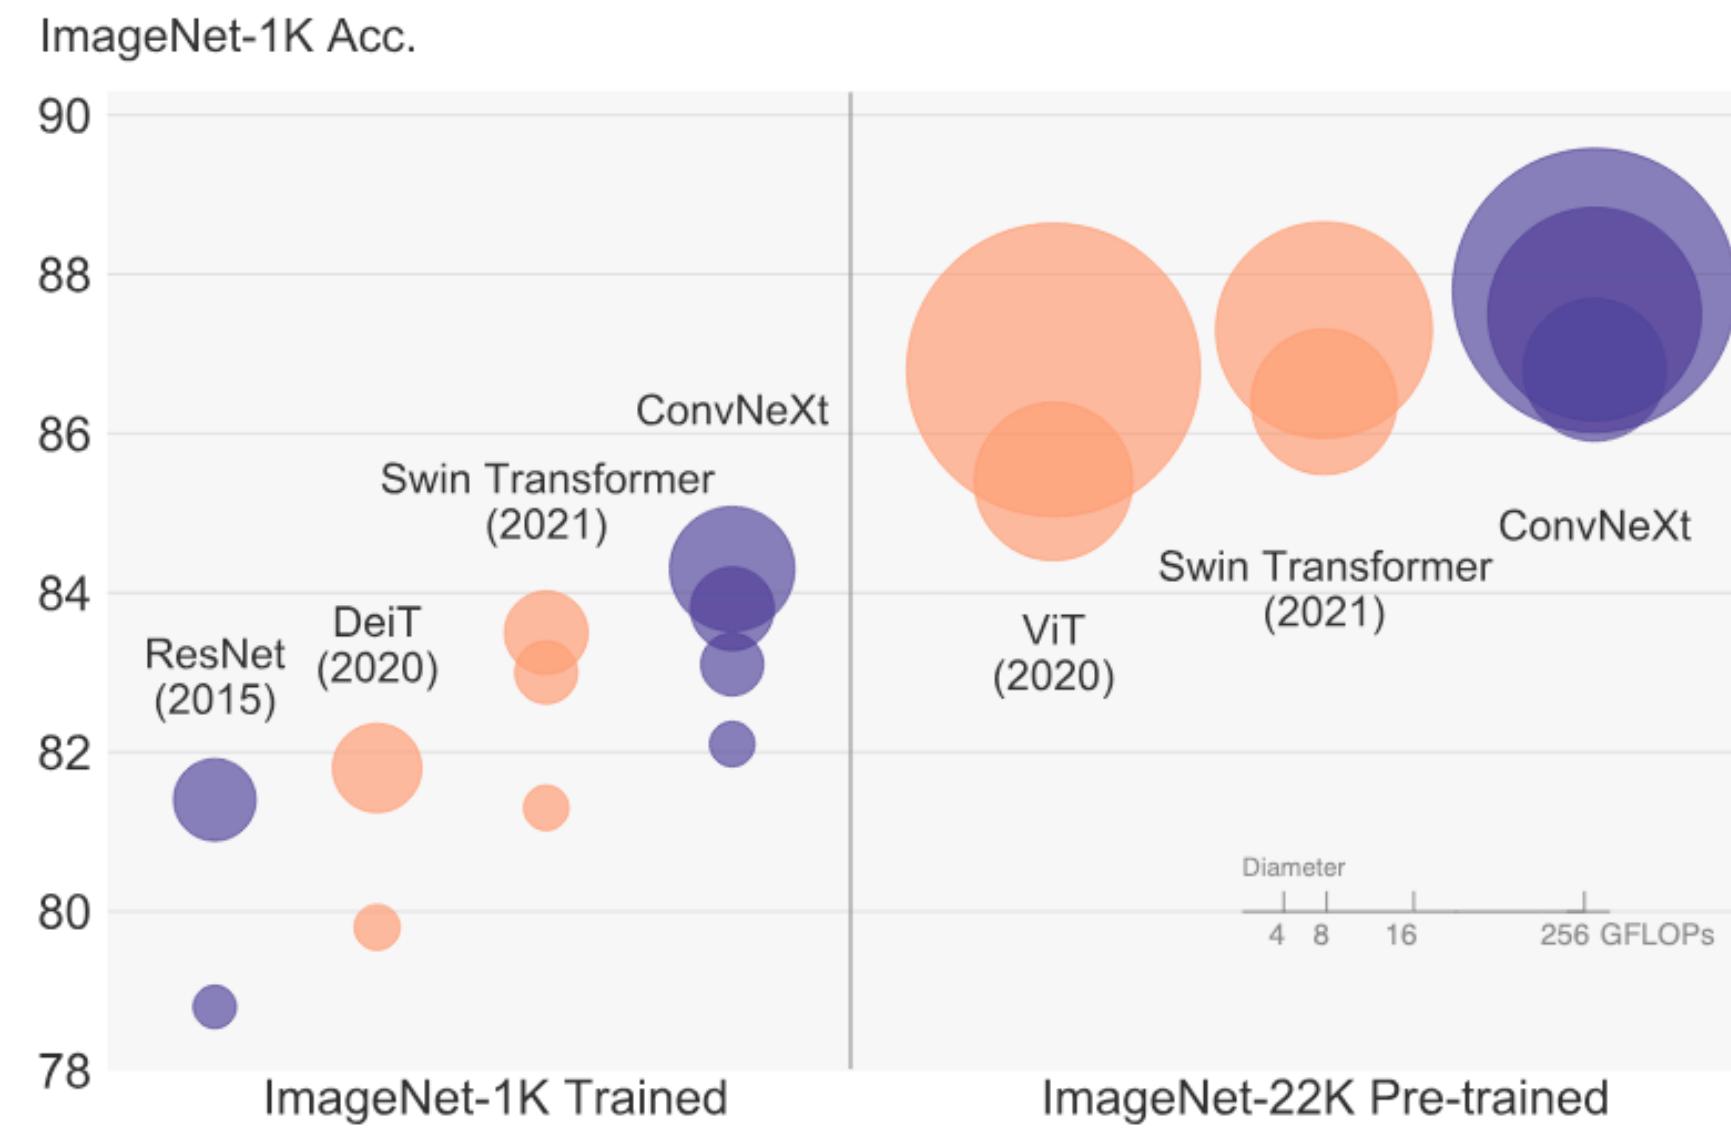


- Cut image into 16x16 patches, use as input sequence to transformer
- like CNNs: localization + weight sharing
- beats CNNs in common tasks at time of publication

# A short history of CNNs



# ConvNext

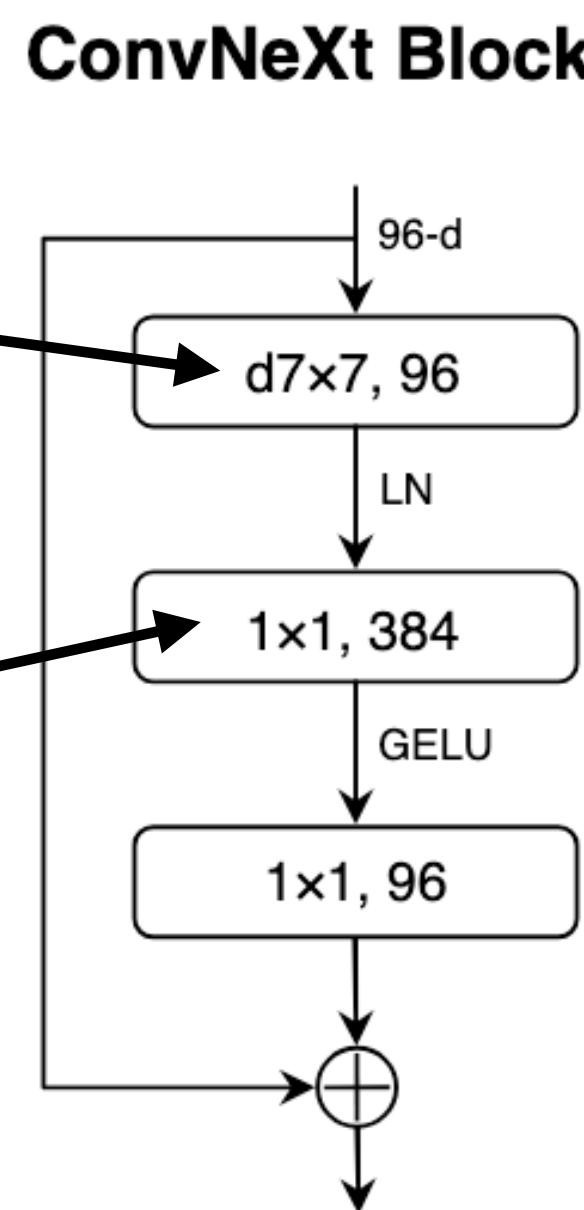


- Modernized ResNet-like architecture
  - Pure convolutional network → simple and efficient
  - Tuned taking inspiration from transformer-based models
  - state of the art performance

# **depth-wise convolutions**

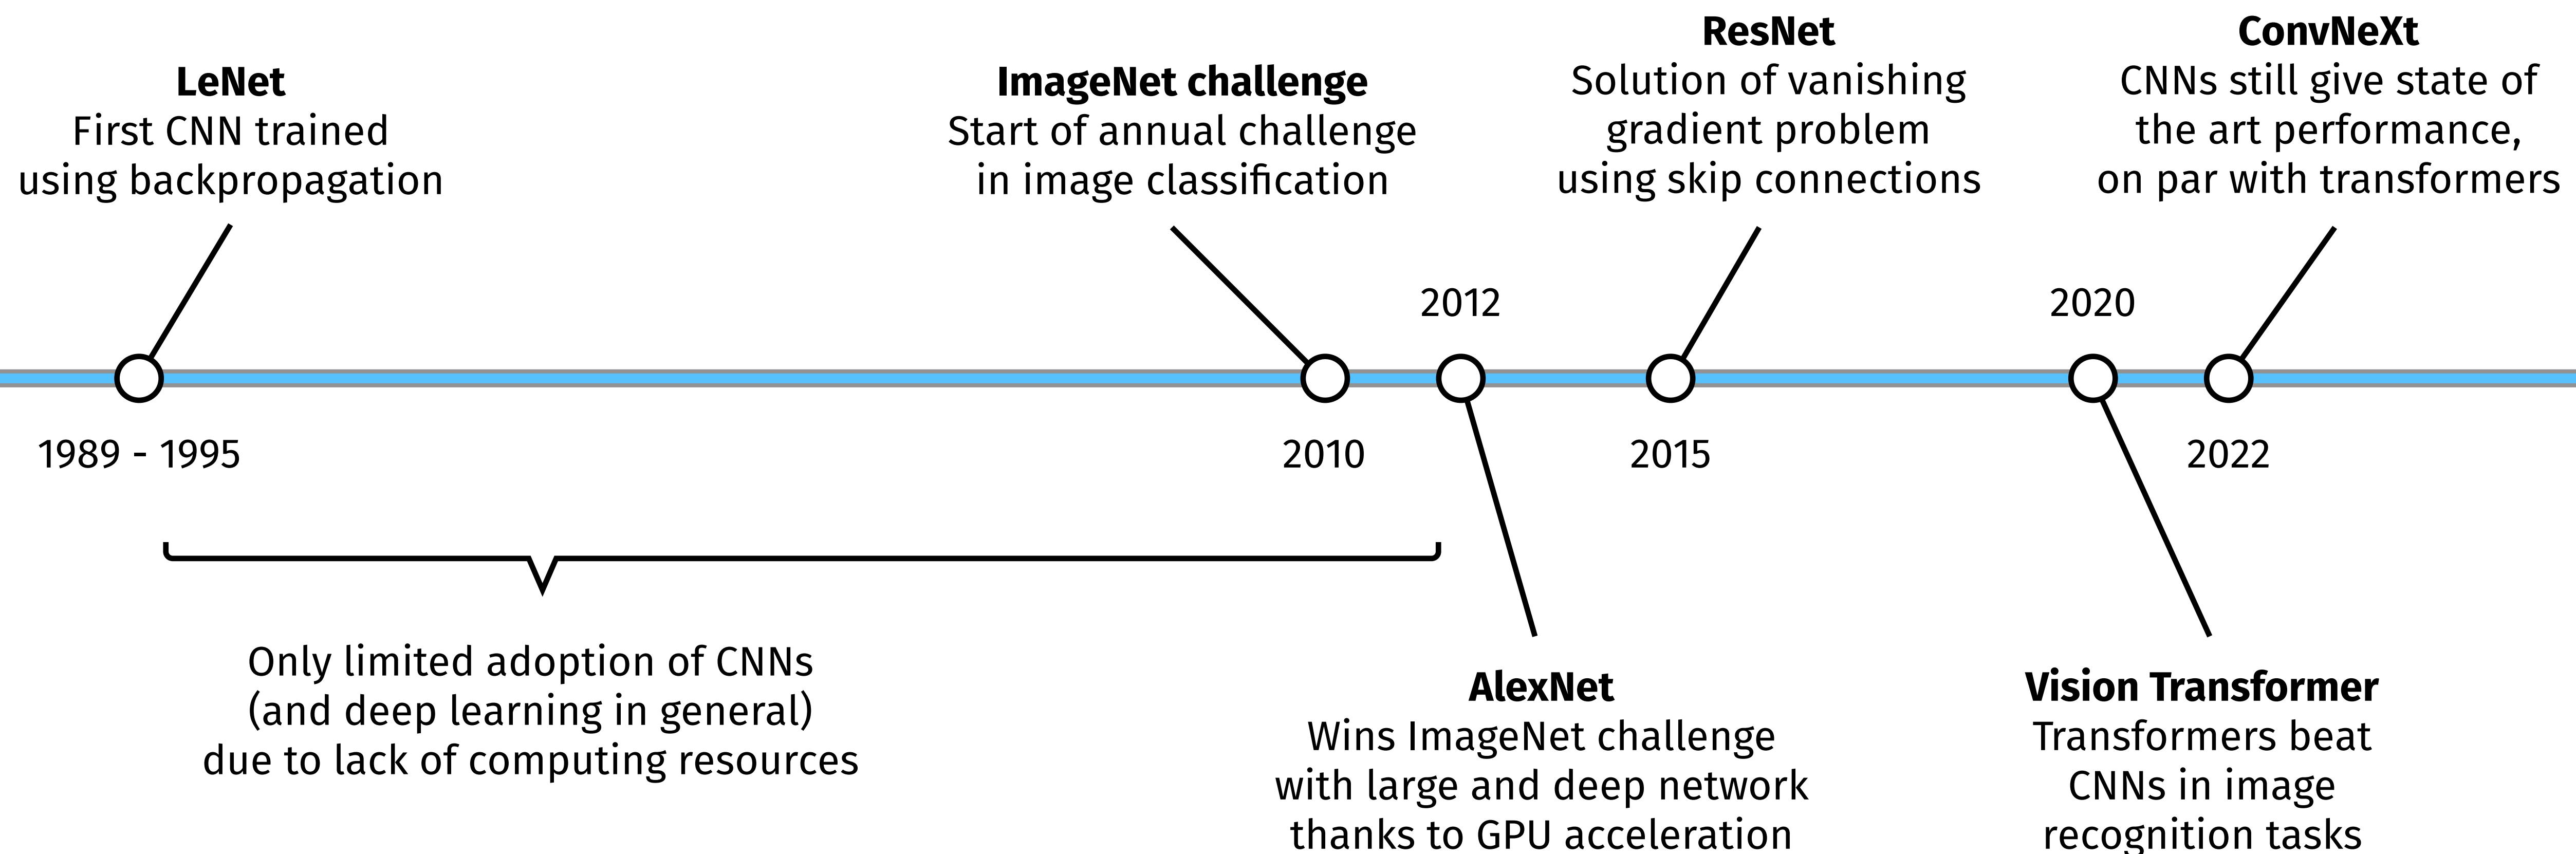
only one kernel per channel

channels interact via  
1x1 convolutions



	output size	• ConvNeXt-T
stem	$56 \times 56$	$4 \times 4, 96$ , stride 4
res2	$56 \times 56$	$\begin{bmatrix} d7 \times 7, 96 \\ 1 \times 1, 384 \\ 1 \times 1, 96 \end{bmatrix} \times 3$
res3	$28 \times 28$	$\begin{bmatrix} d7 \times 7, 192 \\ 1 \times 1, 768 \\ 1 \times 1, 192 \end{bmatrix} \times 3$
res4	$14 \times 14$	$\begin{bmatrix} d7 \times 7, 384 \\ 1 \times 1, 1536 \\ 1 \times 1, 384 \end{bmatrix} \times 9$
res5	$7 \times 7$	$\begin{bmatrix} d7 \times 7, 768 \\ 1 \times 1, 3072 \\ 1 \times 1, 768 \end{bmatrix} \times 3$
FLOPs		$4.5 \times 10^9$
# params.		$28.6 \times 10^6$

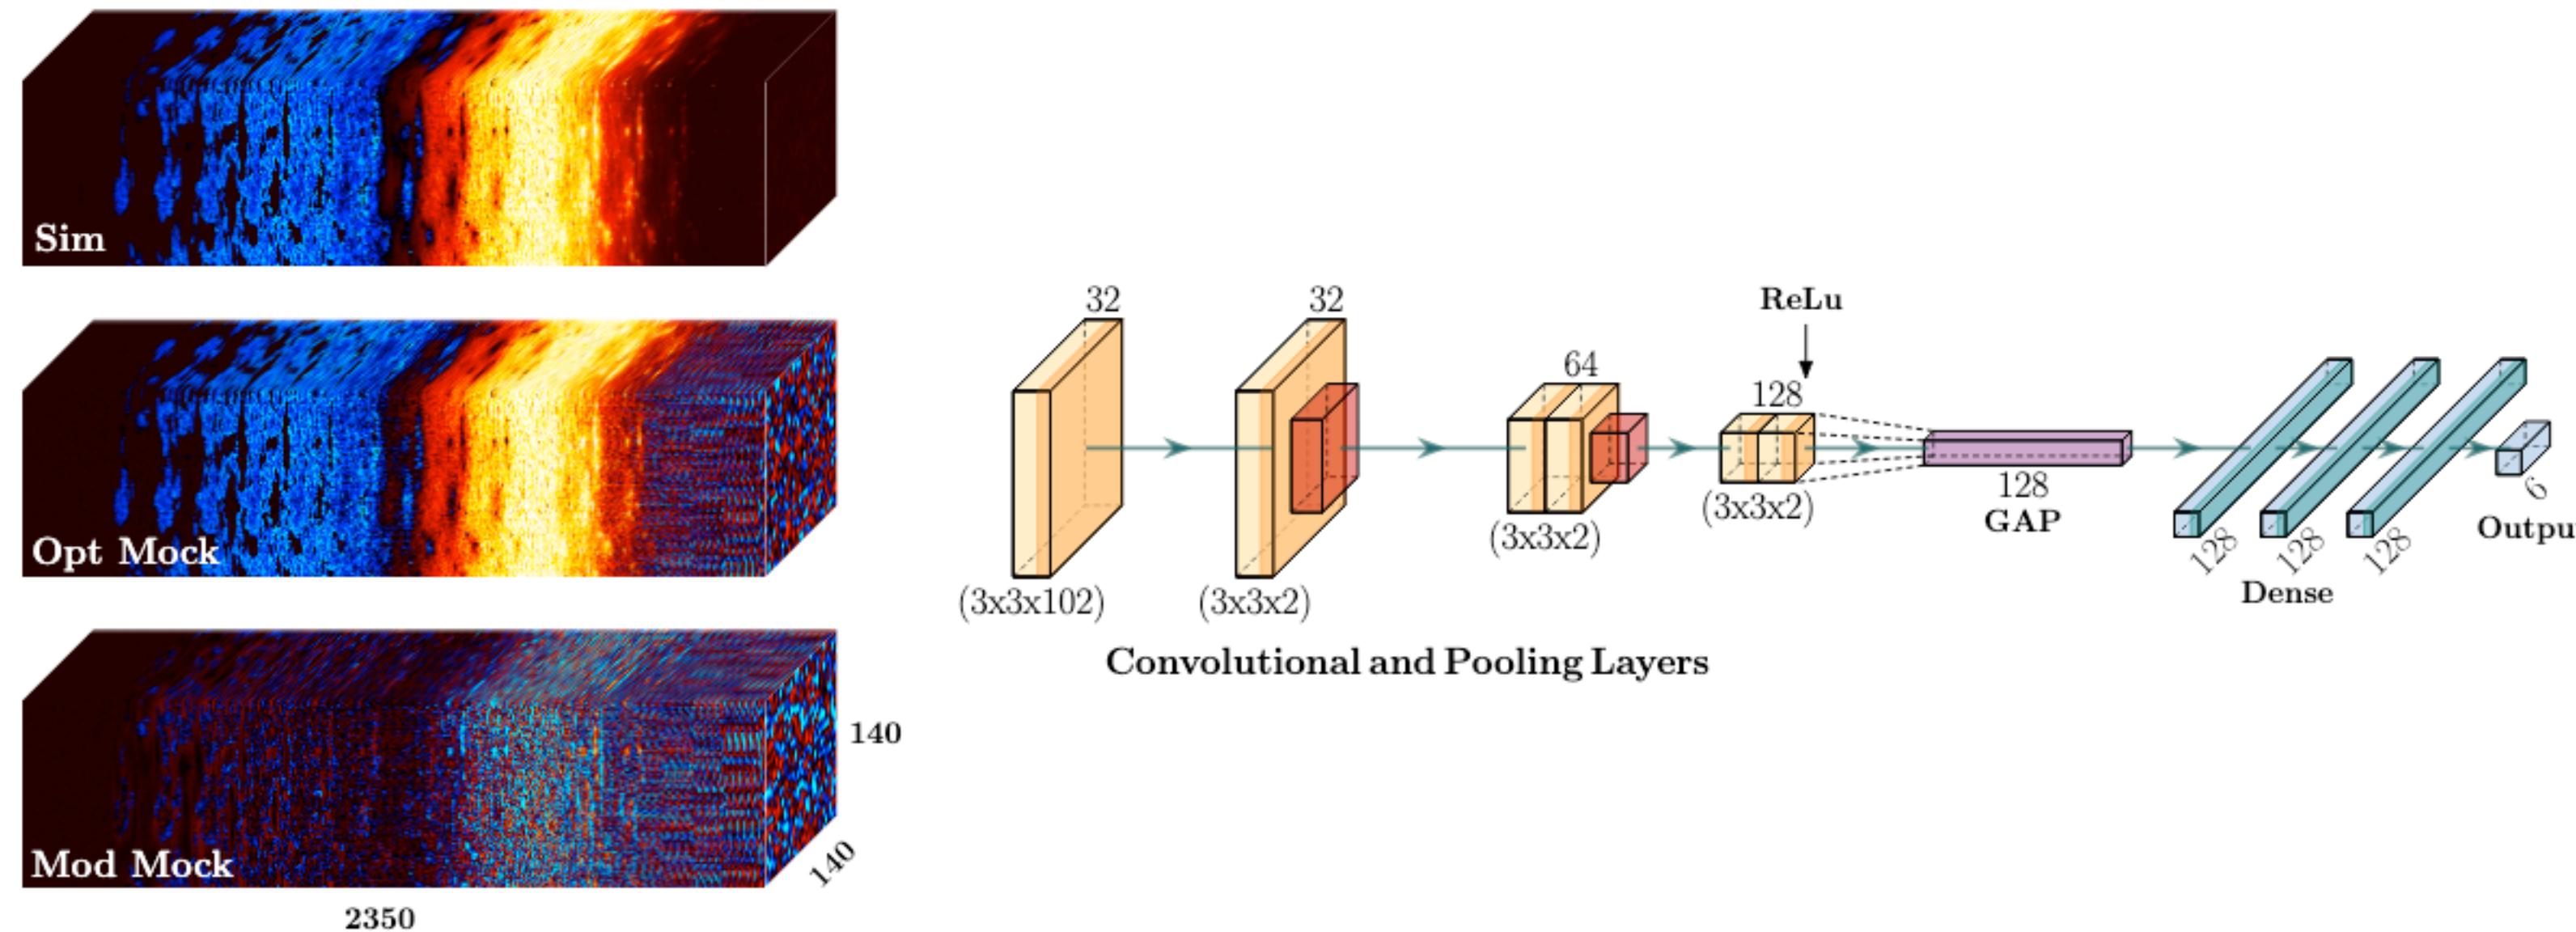
# A short history of CNNs



# Outline

- ① Challenges of image processing
- ② Introduction to convolutions
- ③ Convolutional layers
- ④ CNN building blocks
- ⑤ A short history of CNNs
- ⑥ Outlook

# Different dimensions

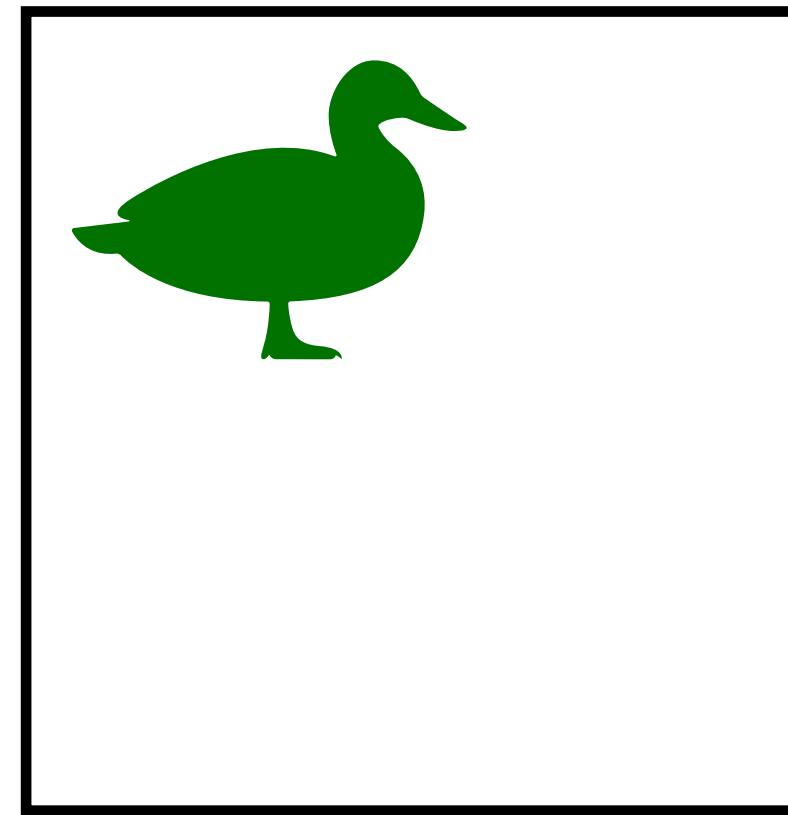


Application in astrophysics: 21cm light-cones [[arXiv:2201.07587](https://arxiv.org/abs/2201.07587)]

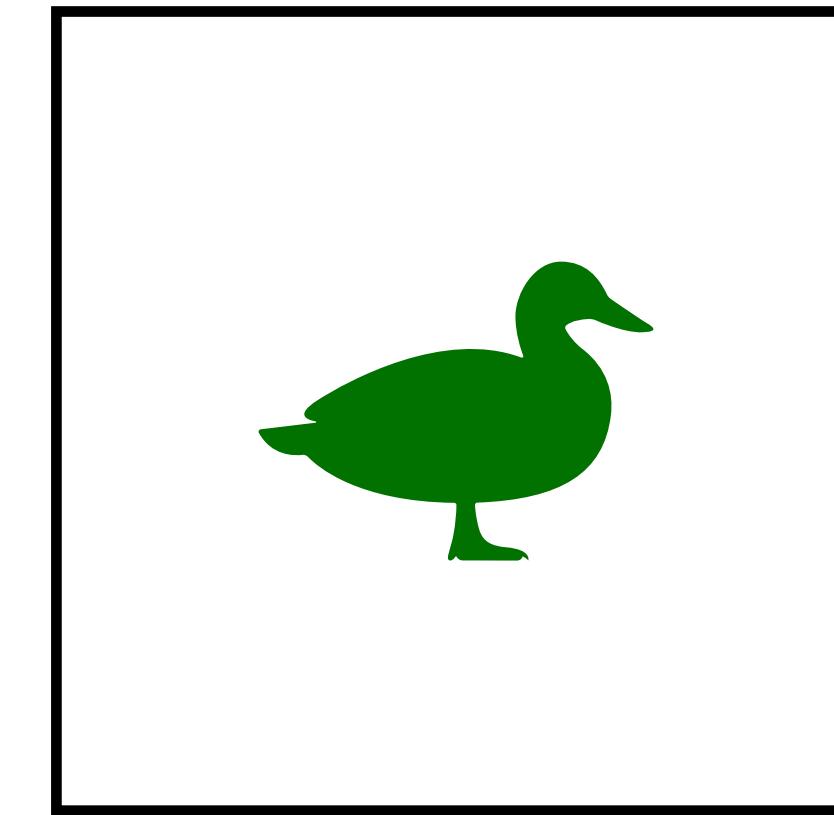
- CNNs are not restricted to two dimensions
- one-dimensional data: time series, ...
- three-dimensional data:
  - CT and MRT scans in medicine
  - observations from radio telescopes

# Invariance and equivariance

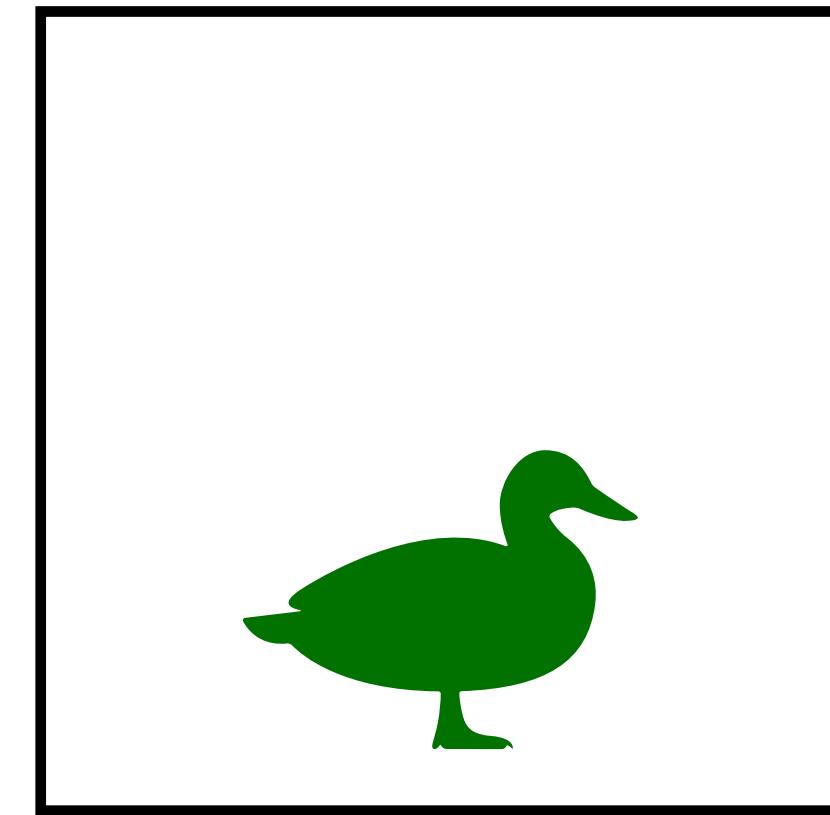
A duck



Also a duck



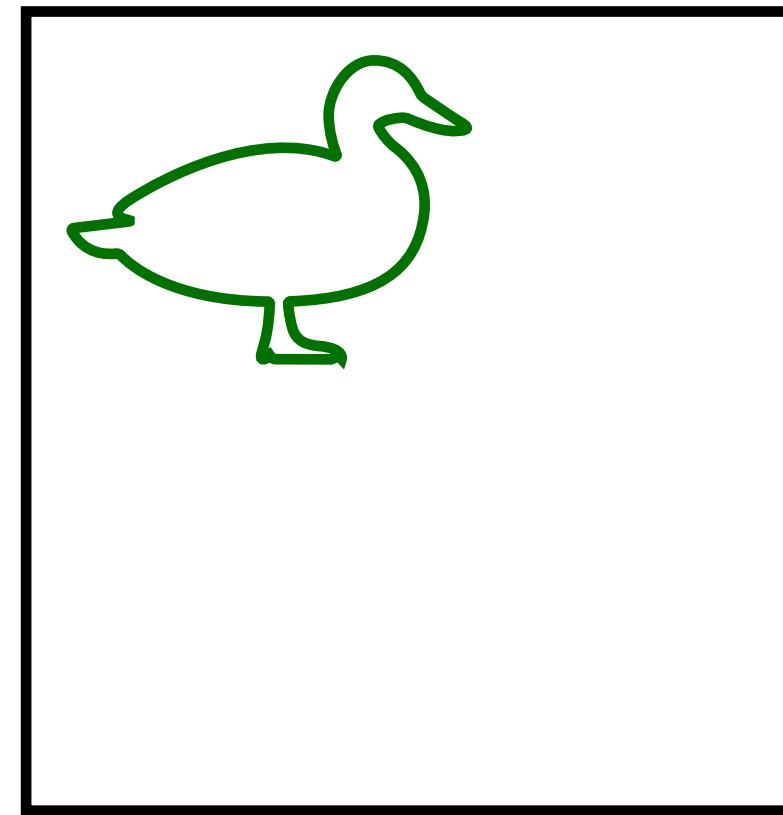
Another duck



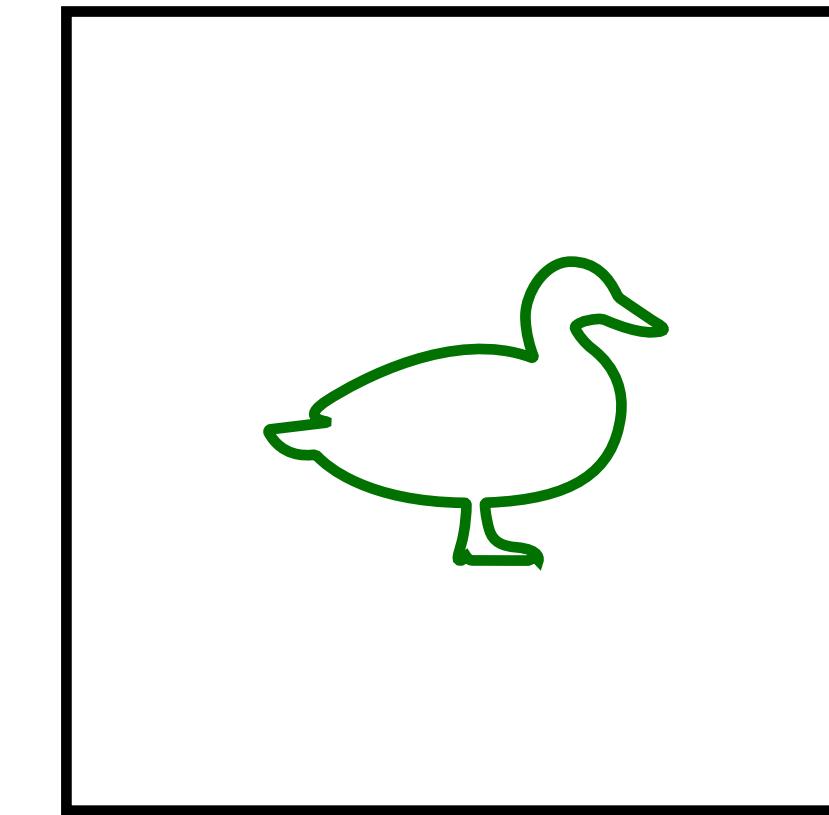
- Main motivation for introducing convolutions:  
want classifier result to be **permutation invariant**

# Invariance and equivariance

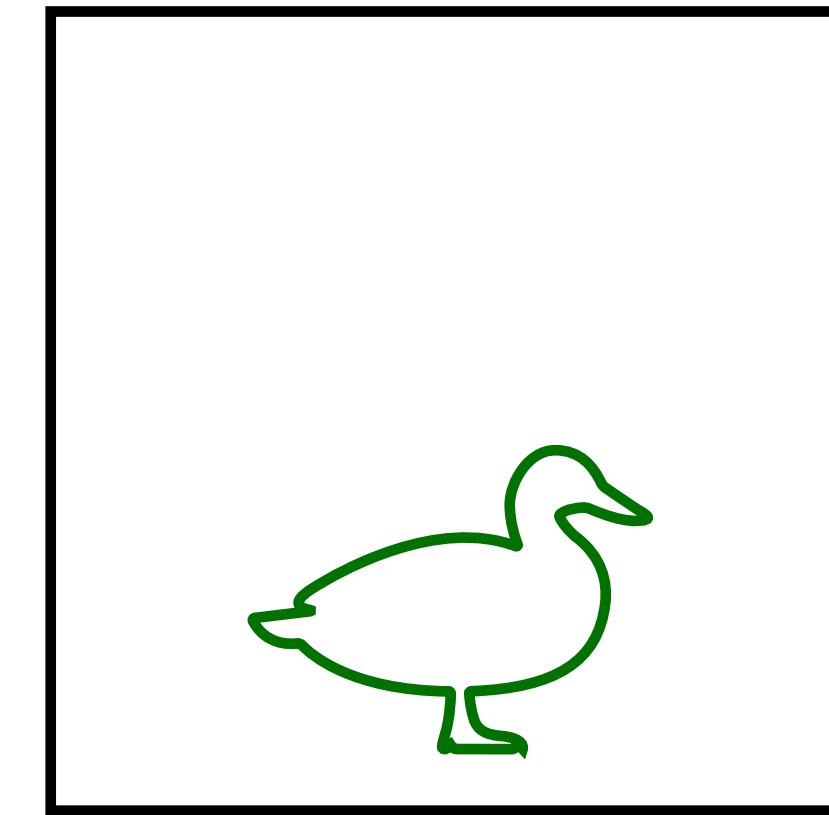
A duck



Also a duck



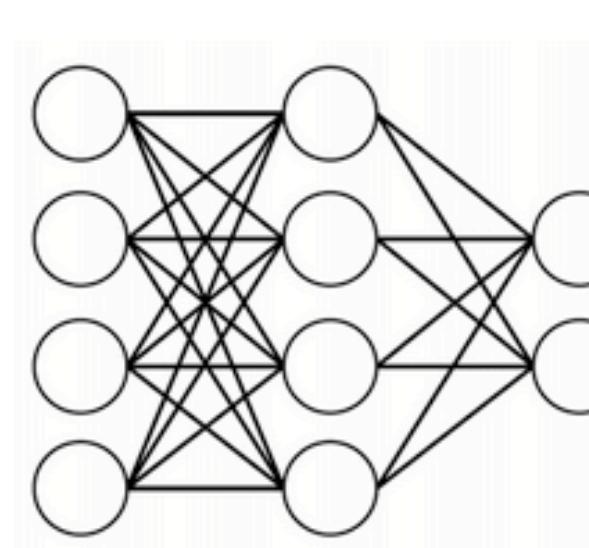
Another duck



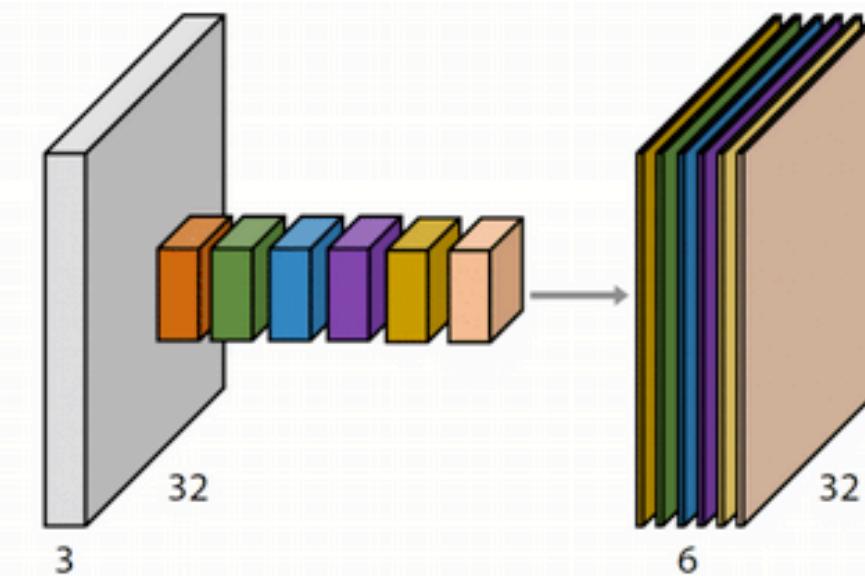
- Main motivation for introducing convolutions:  
want classifier result to be **permutation invariant**
- Solution: CNN layer is **permutation equivariant**  
→ transformation acts the same at different positions
- Symmetry of our task leads to specific architecture

# Geometric deep learning

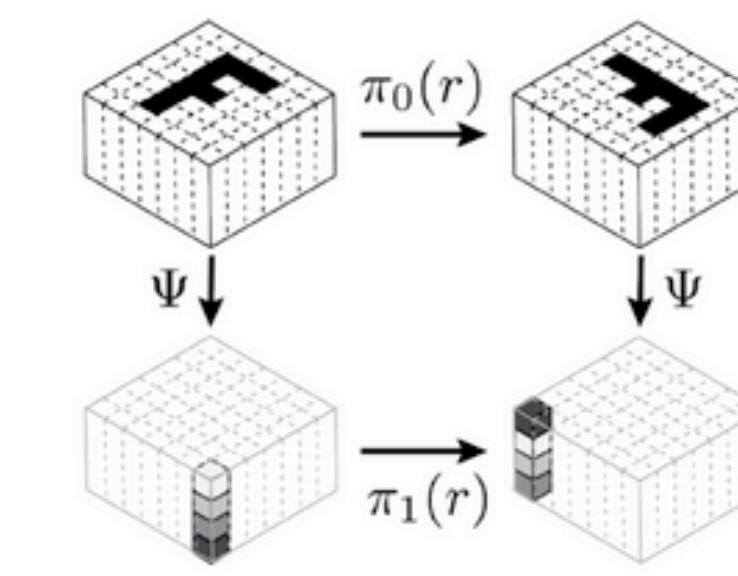
**Many network architectures follow from a symmetry!**



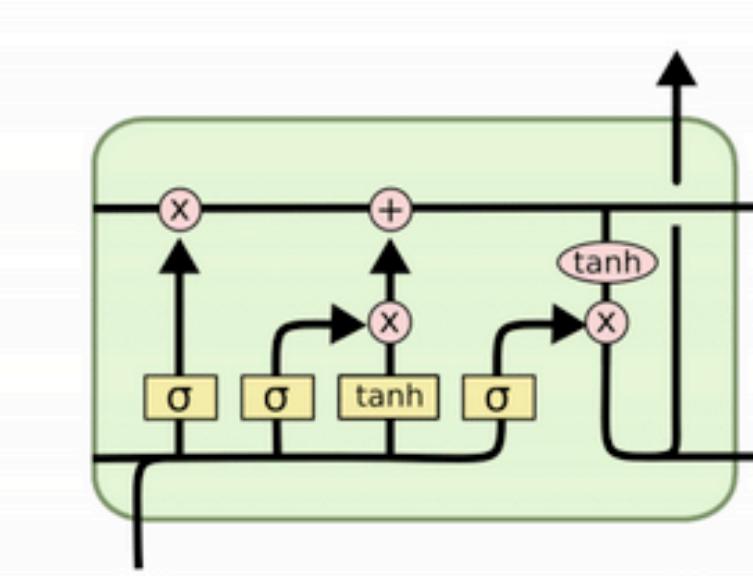
**Perceptrons**  
Function regularity



**CNNs**  
Translation



**Group-CNNs**  
Translation+Rotation,  
Global groups



**LSTMs**  
Time warping



**DeepSets / Transformers**  
Permutation



**GNNs**  
Permutation



**Intrinsic CNNs**  
Isometry / Gauge choice

Book and lectures on the topic:  
<https://geometricdeeplearning.com/>

# Exercises



Now it's time for a coffee,



and then some exercises!

You can find the material here:

<https://github.com/theoheimel/erum-cnn-2025>